

Factoid: The 5G Core Network introduced in 3GPP Release 15 adopts a Service-Based Architecture (SBA), where network functions expose services via standardized APIs (primarily HTTP/2 + JSON).

- **5G Core Network** *is introduced in 3GPP Release 15*
- **5G Core Network** *adopts architecture SBA*
- **Network Function** *exposes Service*
- **Service** *is available via API (HTTP/2 + JSON)*
- **API** *uses protocol HTTP/2*
- **API** *uses format JSON*

Factoid: SBA enables dynamic service discovery and direct NF-to-NF communication without requiring intermediaries or brokers.

- **Service-Based Architecture (SBA)** *enables dynamic service discovery*
- **Service-Based Architecture (SBA)** *enables direct NF-to-NF communication*
- **Direct NF-to-NF communication** *does not require intermediaries or brokers*

Factoid: The SBA model decouples the service consumer from the service provider through the use of a Network Repository Function (NRF).

- **SBA model** *decouples service consumer from service provider*
- **SBA model** *achieves decoupling via Network Repository Function (NRF)*
- **Network Repository Function (NRF)** *mediates interaction between service consumer and service provider*

Factoid: SBA supports horizontal scalability and modularity by treating every function (e.g., AMF, SMF, PCF) as an independent service.

- **Service-Based Architecture (SBA)** *supports horizontal scalability*
- **Service-Based Architecture (SBA)** *supports modularity*
- **Service-Based Architecture (SBA)** *treats AMF as an independent service*
- **Service-Based Architecture (SBA)** *treats SMF as an independent service*
- **Service-Based Architecture (SBA)** *treats PCF as an independent service*

Factoid: 5G core and RAN functions are designed as cloud-native components that follow the principles of containerization, statelessness, and modularity.

- **5G core functions** *are designed as cloud-native components*
- **RAN functions** *are designed as cloud-native components*
- **Cloud-native components** *follow the principle of containerization*
- **Cloud-native components** *follow the principle of statelessness*
- **Cloud-native components** *follow the principle of modularity*

Factoid: Network Functions (NFs) are implemented as microservices, enabling CI/CD deployment, independent updates, and fine-grained scaling.

- **Network Functions (NFs)** *are implemented as microservices*

- **Microservices** *enable* CI/CD deployment
- **Microservices** *enable* independent updates
- **Microservices** *enable* fine-grained scaling

Factoid: Each NF runs in isolated execution environments, typically containers orchestrated via Kubernetes or other NFVO frameworks.

- **Network Function (NF)** *runs in isolated execution environments*
- **Isolated execution environments** *are typically realized as containers*
- **Containers** *are orchestrated via Kubernetes*
- **Containers** *may be orchestrated via NFVO frameworks*

Factoid: The use of microservices and container platforms supports vendor-agnostic and hardware-agnostic deployments in hybrid cloud environments.

- **Microservices** *support* vendor-agnostic deployments
- **Microservices** *support* hardware-agnostic deployments
- **Container platforms** *support* vendor-agnostic deployments
- **Container platforms** *support* hardware-agnostic deployments
- **Vendor-agnostic deployments** *occur in hybrid cloud environments*
- **Hardware-agnostic deployments** *occur in hybrid cloud environments*

Factoid: These cloud-native features allow for slicing, multi-tenancy, and elasticity in managing network resources.

- **Cloud-native features** *allow* slicing
- **Cloud-native features** *allow* multi-tenancy
- **Cloud-native features** *allow* elasticity
- **Slicing** *is used in managing network resources*
- **Multi-tenancy** *is used in managing network resources*
- **Elasticity** *is used in managing network resources*

Factoid: 5G continues the Control and User Plane Separation (CUPS) principle introduced in 3GPP Rel-14, allowing independent placement and scaling of control-plane (AMF, SMF) and user-plane (UPF) functions.

- **5G** *continues* Control and User Plane Separation (CUPS) principle
- **Control and User Plane Separation (CUPS) principle** *was introduced in 3GPP Release 14*
- **CUPS principle** *allows* independent placement of control-plane functions
- **CUPS principle** *allows* independent scaling of control-plane functions
- **CUPS principle** *allows* independent placement of user-plane functions
- **CUPS principle** *allows* independent scaling of user-plane functions
- **Control-plane functions** *include* AMF
- **Control-plane functions** *include* SMF
- **User-plane functions** *include* UPF

Factoid: CUPS enables deploying UPF close to the network edge for latency-sensitive services, while control-plane functions remain centralized.

- **CUPS enables deploying UPF close to the network edge**
- **UPF close to the network edge supports latency-sensitive services**
- **Control-plane functions remain centralized**

Factoid: Separation of control and user planes enhances fault isolation and resource utilization across the network.

- **Separation of control and user planes (CUPS) enhances fault isolation**
- **Separation of control and user planes (CUPS) enhances resource utilization**
- **Resource utilization occurs in the network**

Factoid: CUPS is a prerequisite for enabling flexible and performant network slicing and multi-access edge computing (MEC).

- **CUPS is a prerequisite for flexible network slicing**
- **CUPS is a prerequisite for performant network slicing**
- **CUPS is a prerequisite for multi-access edge computing (MEC)**

Factoid: In 5G SBA, core functions like AMF, SMF, AUSF, and UDM can operate statelessly by offloading UE context to an external Unstructured Data Storage Function (UDSF).

- **5G SBA enables stateless operation of AMF**
- **5G SBA enables stateless operation of SMF**
- **5G SBA enables stateless operation of AUSF**
- **5G SBA enables stateless operation of UDM**
- **AMF offloads UE context to Unstructured Data Storage Function (UDSF)**
- **SMF offloads UE context to Unstructured Data Storage Function (UDSF)**
- **AUSF offloads UE context to Unstructured Data Storage Function (UDSF)**
- **UDM offloads UE context to Unstructured Data Storage Function (UDSF)**
- **Offloading UE context to UDSF enables stateless operation**

Factoid: Stateless NFs gain elasticity and failure resilience, enabling container-based deployments and fast restarts without local state loss.

- **Stateless NFs gain elasticity**
- **Stateless NFs gain failure resilience**
- **Stateless NFs enable container-based deployments**
- **Stateless NFs enable fast restarts**
- **Fast restarts occur without local state loss**

Factoid: UDSF acts as a centralized, non-relational key-value store, enabling fast read/write of UE-related state.

- **UDSF acts as centralized non-relational key-value store**
- **UDSF enables fast read/write of UE-related state**
- **Fast read/write operates on UE-related state**

Factoid: Stateful NFs are tightly coupled with internal UE context, causing failure risks and scaling limitations.

- **Stateful NFs** *are tightly coupled with internal UE context*
- **Stateful NFs tightly coupled with internal UE context** *incur failure risks*
- **Stateful NFs tightly coupled with internal UE context** *incur scaling limitations*

Factoid: The *piggyback-based approach* retrieves UE context once per procedure and embeds it in all NF-to-NF HTTP calls, reducing repeated fetches.

- **Piggyback-based approach** *retrieves UE context once per procedure*
- **Piggyback-based approach** *embeds UE context in NF-to-NF HTTP calls*
- **Embedding UE context in NF-to-NF HTTP calls** *reduces repeated fetches*

Factoid: Using piggybacking reduces procedure latency by ~44% for registration and up to ~70% for deregistration procedures.

- **Piggybacking** *reduces registration procedure latency by $\approx 44\%$*
- **Piggybacking** *reduces deregistration procedure latency by up to $\approx 70\%$*

Factoid: The *proactive-push approach* lets the AMF preemptively instruct UDSF to push context to downstream NFs via new SBI POST endpoints

- **Proactive-push approach** *enables AMF to preemptively instruct UDSF*
- **AMF** *preemptively instructs UDSF to push context*
- **UDSF** *pushes context to downstream NFs*
- **Context** *is pushed via new SBI POST endpoints*

```
ex:AMF a rdfs:Class .
ex:UDSF a rdfs:Class .
ex:NF a rdfs:Class .
ex:ContextPush a rdfs:Class .
ex:SBIPostEndpoint a rdfs:Class .
ex:proactivePushApproach a ex:Approach ;
  rdfs:label "Proactive-push approach" .
```

```
ex:preemptivelyInstructs a rdf:Property ;
  rdfs:domain ex:AMF ;
  rdfs:range ex:UDSF .
```

```
ex:pushesContextTo a rdf:Property ;
  rdfs:domain ex:UDSF ;
  rdfs:range ex:NF .
```

```
ex:usesSBIEndpoint a rdf:Property ;
  rdfs:domain ex:ContextPush ;
  rdfs:range ex:SBIPostEndpoint .
```

```
ex:proactivePushApproach ex:enables ex:AMFpreemptivelyInstructingUDSF .
```

```
ex:AMFpreemptivelyInstructingUDSF
  a ex:ContextPush ;
```

ex:performedBy ex:AMF ;
ex:target ex:UDSF ;
rdfs:comment "AMF preemptively instructs UDSF to push context." .

ex:AMF ex:preemptivelyInstructs ex:UDSF .
ex:UDSF ex:pushesContextTo ex:NF .
ex:**ContextPush** ex:usesSBIendpoint ex:SBIPostEndpoint .

Factoid: Proactive-push improves asynchronous procedure completion times (like PDU session setup) by 13–22%.

- **Proactive-push** *improves asynchronous procedure completion times by \approx 13–22 %*
- **Asynchronous procedure completion times** *include PDU session setup*

Factoid: Piggybacking is optimal for synchronous flows; proactive-push works best for asynchronous flows.

- **Piggybacking** *is optimal for synchronous flows*
- **Proactive-push** *works best for asynchronous flows*

Factoid: A hybrid of piggyback and proactive-push yields best end-to-end control procedure performance.

- **Hybrid approach** *combines piggyback and proactive-push*
- **Hybrid approach** *yields best end-to-end control procedure performance*

Factoid: These approaches require no major changes to core 5G architecture but depend on UDSF and SBI endpoint extension.

- **These approaches** *require no major changes to core 5G architecture*
- **These approaches** *depend on UDSF*
- **These approaches** *depend on SBI endpoint extension*

Factoid: Prototype deployments on Kubernetes with Free5GC confirmed stateless NFs with minimal CPU/memory overhead.

- **Prototype deployments** *were conducted on Kubernetes*
- **Prototype deployments** *used Free5GC*
- **Prototype deployments** *confirmed stateless NFs*
- **Stateless NFs** *incur minimal CPU/memory overhead*

Factoid: Procedure-aware optimizations are backward compatible and non-intrusive to existing SBI-based service flows.

- **Procedure-aware optimizations** *are backward compatible with existing SBI-based service flows*
- **Procedure-aware optimizations** *are non-intrusive to existing SBI-based service flows*

Factoid: NG-RAN is split into gNB-DU for real-time processing and gNB-CU for centralized control and core interfacing.

- **NG-RAN is split into gNB-DU**
- **NG-RAN is split into gNB-CU**
- **gNB-DU performs real-time processing**
- **gNB-CU performs centralized control**
- **gNB-CU performs core interfacing**

Factoid: gNB-CU is functionally divided into CU-CP (control plane) and CU-UP (user plane), mirroring CUPS design.

- **gNB-CU is functionally divided into CU-CP**
- **gNB-CU is functionally divided into CU-UP**
- **gNB-CU functional division mirrors CUPS design**

Factoid: CU-CP connects to AMF over NG-AP/SCTP (NG-C link); CU-UP transfers data to UPF via NG-U using GTP-U.

- **CU-CP connects to AMF**
- **CU-CP-AMF connection uses NG-AP**
- **CU-CP-AMF connection uses SCTP**
- **CU-CP-AMF connection constitutes NG-C link**
- **CU-UP transfers data to UPF**
- **Data transfer between CU-UP and UPF occurs via NG-U**
- **Data transfer between CU-UP and UPF uses GTP-U**

Factoid: E1 interface supports CU-CP → CU-UP communication via E1AP over SCTP for bearer/context management.

- **E1 interface supports CU-CP ↔ CU-UP communication**
- **CU-CP ↔ CU-UP communication uses E1AP**
- **E1AP runs over SCTP**
- **CU-CP ↔ CU-UP communication is used for bearer/context management**

Factoid: F1-C and F1-U split handles control and user-plane data between CU and DU using SCTP.

- **F1-C split handles control-plane data between CU and DU**
- **F1-U split handles user-plane data between CU and DU**
- **F1-C split uses SCTP**
- **F1-U split uses SCTP**

Factoid: CUPS architecture enables horizontal scaling and placement flexibility across both RAN and core.

- **CUPS architecture enables horizontal scaling**
- **CUPS architecture enables placement flexibility**
- **Horizontal scaling applies across RAN**
- **Horizontal scaling applies across core**

- **Placement flexibility** *applies across RAN*
- **Placement flexibility** *applies across core*

Factoid: 5G NR uses scalable OFDM numerology, ultra-lean design, and beamforming for diverse service support (eMBB, mMTC, URLLC).

- **5G NR** *uses scalable OFDM numerology*
- **5G NR** *uses ultra-lean design*
- **5G NR** *uses beamforming*
- **Scalable OFDM numerology** *supports diverse service types*
- **Ultra-lean design** *supports diverse service types*
- **Beamforming** *supports diverse service types*
- **Diverse service types** *include eMBB*
- **Diverse service types** *include mMTC*
- **Diverse service types** *include URLLC*

Factoid: Separating control and data functions allows independent edge deployment—CU-UP at edge, CU-CP centralized—for low latency.

- **Separation of control and data functions** *allows independent edge deployment*
- **Separation of control and data functions** *allows deploying CU-UP at the edge*
- **Separation of control and data functions** *allows keeping CU-CP centralized*
- **CU-UP at the edge** *reduces latency*
- **Independent edge deployment** *enables low latency*

Factoid: Functional splits enhance vendor diversity, network slicing, and multi-tenant orchestration in cloud-native environments.

- **Functional splits** *enhance vendor diversity*
- **Functional splits** *enhance network slicing*
- **Functional splits** *enhance multi-tenant orchestration*
- **Multi-tenant orchestration** *occurs in cloud-native environments*

Factoid: SBA control plane is more vulnerable due to REST over HTTP/2 interfaces between NFs, which enable new attack vectors like header manipulation and horizontal DDoS.

- **SBA control plane** *is more vulnerable due to REST over HTTP/2 interfaces*
- **REST over HTTP/2 interfaces** *connect network functions (NFs)*
- **REST over HTTP/2 interfaces between NFs** *enable new attack vectors*
- **New attack vectors** *include header manipulation*
- **New attack vectors** *include horizontal DDoS*

Factoid: NRF centralizes service registry and discovery; if compromised, it can route traffic to malicious or rogue NFs.

- **NRF** *centralizes service registry*
- **NRF** *centralizes service discovery*
- **Compromised NRF** *can route traffic to malicious NFs*

- **Compromised NRF** *can route traffic to rogue NFs*

Factoid: End-to-end mutual TLS (mTLS) is required to secure NF-to-NF communication in SBA control plane.

- **End-to-end mutual TLS (mTLS)** *is required to secure NF-to-NF communication*
- **NF-to-NF communication** *occurs in SBA control plane*

Factoid: API gateways or service proxies can detect anomalous patterns, rate-limit flows, and enforce role-based access to NF services.

- **API gateways** *detect anomalous patterns*
- **Service proxies** *detect anomalous patterns*
- **API gateways** *rate-limit flows*
- **Service proxies** *rate-limit flows*
- **API gateways** *enforce role-based access to NF services*
- **Service proxies** *enforce role-based access to NF services*

Factoid: Deploying PKI with PLMN-bound CAs allows strong NF identity verification and integrity within multitenant operator domains.

- **PKI with PLMN-bound CAs** *allows strong NF identity verification*
- **PKI with PLMN-bound CAs** *ensures integrity within multitenant operator domains*
- **Strong NF identity verification** *applies within multitenant operator domains*
- **Integrity** *applies within multitenant operator domains*

Factoid: NF deployments must include micro-segmentation and zero-trust policies to mitigate lateral threats in disaggregated environments.

- **NF deployments** *must include micro-segmentation*
- **NF deployments** *must include zero-trust policies*
- **Micro-segmentation** *mitigates lateral threats*
- **Zero-trust policies** *mitigate lateral threats*
- **Lateral threats** *occur in disaggregated environments*

Factoid: Control plane security must encompass confidentiality, integrity, authentication, segmentation, and runtime monitoring across all NF interfaces.

- **Control plane security** *must encompass confidentiality*
- **Control plane security** *must encompass integrity*
- **Control plane security** *must encompass authentication*
- **Control plane security** *must encompass segmentation*
- **Control plane security** *must encompass runtime monitoring*
- **Confidentiality** *must apply across all NF interfaces*
- **Integrity** *must apply across all NF interfaces*
- **Authentication** *must apply across all NF interfaces*
- **Segmentation** *must apply across all NF interfaces*
- **Runtime monitoring** *must apply across all NF interfaces*

Factoid: AMF manages UE registration, mobility, and NAS signaling.

- **AMF manages UE registration**
- **AMF manages mobility**
- **AMF manages NAS signaling**

Factoid: AMF authenticates UE using NAS and passes session control to SMF.

- **AMF authenticates UE using NAS**
- **AMF passes session control to SMF**

Factoid: AMF is control-plane only; it does not handle user traffic.

- **AMF is control-plane only**
- **AMF does not handle user traffic**

Factoid: AMF can redirect UEs to appropriate SMF/UPF combinations based on policies.

- **AMF redirects UEs to appropriate SMF/UPF combinations**
- **AMF selects SMF/UPF combinations based on policies**

Factoid: AMF is dependent on UDM (for subscription), AUSF (for auth), SMF (for session setup).

- **AMF depends on UDM** for subscription
- **AMF depends on AUSF** for authentication
- **AMF depends on SMF** for session setup

Factoid: AMF uses such interfaces as: N1 (UE-AMF) NAS over IP, N2 (AMF-gNB) NGAP over SCTP, N11 (AMF-SMF) SBI (HTTP/2+JSON).

- **AMF uses N1 interface**
- **N1 interface connects UE and AMF**
- **N1 interface carries NAS over IP**
- **AMF uses N2 interface**
- **N2 interface connects AMF and gNB**
- **N2 interface carries NGAP over SCTP**
- **AMF uses N11 interface**
- **N11 interface connects AMF and SMF**
- **N11 interface implements SBI over HTTP/2 + JSON**

Factoid: SMF handles PDU session lifecycle: setup, modification, and release.

- **SMF handles PDU session lifecycle**
- **SMF handles PDU session setup**
- **SMF handles PDU session modification**
- **SMF handles PDU session release**

Factoid: SMF allocates IP addresses and installs traffic steering rules in UPFs.

- **SMF allocates IP addresses**
- **SMF installs traffic steering rules**
- **Traffic steering rules are installed in UPFs**

Factoid: SMF enforces QoS via PCF policies.

- **SMF enforces QoS**
- **SMF uses PCF policies to enforce QoS**

Factoid: SMF connects to UPFs via PFCP, a protocol designed for fast, stateless user-plane control.

- **SMF connects to UPFs**
- **SMF–UPF connection uses PFCP**
- **PFCP is designed for fast user-plane control**
- **PFCP is designed for stateless user-plane control**

Factoid: SMF uses the following interfaces: N11 (SMF-AMF), N4 (SMF-UPF) PFCP over UDP, and N10 (SMF-UDM).

- **SMF uses N11 interface**
- **N11 interface connects SMF and AMF**
- **SMF uses N4 interface**
- **N4 interface connects SMF and UPF**
- **N4 interface uses PFCP**
- **PFCP runs over UDP**
- **SMF uses N10 interface**
- **N10 interface connects SMF and UDM**

Factoid: SMF is dependent on AMF, UPF, and PCF.

- **SMF depends on AMF**
- **SMF depends on UPF**
- **SMF depends on PCF**

Factoid: SMF exchanges data such as session context, QoS rules, and IP address allocation.

- **SMF exchanges session context**
- **SMF exchanges QoS rules**
- **SMF exchanges IP address allocation**

Factoid: SMF uses protocols including HTTP/2 and PFCP.

- **SMF uses HTTP/2**
- **SMF uses PFCP**

Factoid: UPF handles user traffic forwarding and QoS enforcement.

- **UPF handles user traffic forwarding**

- **UPF handles QoS enforcement**

Factoid: UPF forwards traffic between RAN and external data networks.

- **UPF forwards traffic between RAN and external data networks**
- **UPF connects RAN to external data networks**

Factoid: UPF supports data buffering, downlink packet marking, and traffic shaping.

- **UPF supports data buffering**
- **UPF supports downlink packet marking**
- **UPF supports traffic shaping**

Factoid: Multiple UPFs can be deployed regionally to support MEC and network slicing.

- **Multiple UPFs can be deployed regionally**
- **Multiple UPFs support multi-access edge computing (MEC)**
- **Multiple UPFs support network slicing**

Factoid: UPF uses the following interfaces: N3 (UPF-gNB) GTP-U, N4 (UPF-SMF) PFCP, N6 (UPF-DN) IP forwarding.

- **UPF uses N3 interface**
- **N3 interface connects UPF and gNB**
- **N3 interface carries GTP-U**
- **UPF uses N4 interface**
- **N4 interface connects UPF and SMF**
- **N4 interface uses PFCP**
- **UPF uses N6 interface**
- **N6 interface connects UPF and data network (DN)**
- **N6 interface provides IP forwarding**

Factoid: UPF is dependent on SMF.

- **UPF depends on SMF**

Factoid: UPF manages data such as PDU sessions and QoS enforcement rules.

- **UPF manages PDU sessions**
- **UPF manages QoS enforcement rules**

Factoid: UPF uses protocols including PFCP and GTP-U.

- **UPF uses PFCP**
- **UPF uses GTP-U**

Factoid: AUSF handles UE authentication using 5G AKA or EAP-AKA'.

- **AUSF handles UE authentication**
- **UE authentication uses 5G AKA**

- **UE authentication** *uses EAP-AKA'*

Factoid: AUSF performs challenge–response procedures with UEs during registration.

- **AUSF** *performs challenge–response procedures*
- **Challenge–response procedures** *involve UEs*
- **Challenge–response procedures** *occur during registration*

Factoid: AUSF is stateless by design and uses UDM to obtain authentication data (e.g., K, OPc, SQN).

- **AUSF** *is stateless by design*
- **AUSF** *uses UDM*
- **AUSF** *obtains authentication data from UDM*
- **Authentication data** *includes K*
- **Authentication data** *includes OPc*
- **Authentication data** *includes SQN*

Factoid: AUSF uses the following interfaces: N12 (AMF-AUSF) and N13 (AUSF-UDM).

- **AUSF** *uses N12 interface*
- **N12 interface** *connects AMF and AUSF*
- **AUSF** *uses N13 interface*
- **N13 interface** *connects AUSF and UDM*

Factoid: AUSF is dependent on UDM for subscription credentials.

- **AUSF** *depends on UDM*
- **AUSF** *obtains subscription credentials from UDM*

Factoid: AUSF processes authentication vectors as part of its data handling.

- **AUSF** *processes authentication vectors*
- **Authentication vectors** *are part of AUSF data handling*

Factoid: UDM serves as the central database for subscriber profiles and policies.

- **UDM** *serves as central database*
- **Central database** *stores subscriber profiles*
- **Central database** *stores policies*

Factoid: UDM provides authentication data to AUSF and access policies to AMF.

- **UDM** *provides authentication data to AUSF*
- **UDM** *provides access policies to AMF*

Factoid: UDM stores SUPI, subscription profiles, and AM Policy Association.

- **UDM** *stores SUPI*
- **UDM** *stores subscription profiles*

- **UDM stores AM Policy Association**

Factoid: UDM uses the following interfaces: N8 (UDM-AMF), N10 (UDM-SMF), N13 (UDM-AUSF).

- **UDM uses N8 interface**
- **N8 interface connects UDM and AMF**
- **UDM uses N10 interface**
- **N10 interface connects UDM and SMF**
- **UDM uses N13 interface**
- **N13 interface connects UDM and AUSF**

Factoid: UDM manages data including subscription data, access profiles, and authentication keys.

- **UDM manages subscription data**
- **UDM manages access profiles**
- **UDM manages authentication keys**

Factoid: UDM may link with the HSS in interworking scenarios for backward compatibility.

- **UDM may link with HSS**
- **UDM-HSS linking occurs in interworking scenarios**
- **UDM-HSS linking provides backward compatibility**

Factoid: PCF provides policy decisions such as QoS and charging to other network functions.

- **PCF provides policy decisions**
- **PCF provides policy decisions to other network functions**
- **Policy decisions include QoS**
- **Policy decisions include charging**

Factoid: PCF enforces subscriber-specific policies via SMF and AMF.

- **PCF enforces subscriber-specific policies**
- **PCF enforces *subscriber-specific policies via SMF**
- **PCF enforces *subscriber-specific policies via AMF**

Factoid: PCF works with AF (Application Function) to enforce application-level QoS.

- **PCF works with AF (Application Function)**
- **PCF and AF (Application Function) enforce application-level QoS**

Factoid: PCF uses the following interfaces: N7 (PCF-SMF) and N15 (PCF-AMF).

- **PCF uses N7 interface**
- **N7 interface connects PCF and SMF**
- **PCF uses N15 interface**
- **N15 interface connects PCF and AMF**

Factoid: PCF manages data including policy rules, network slicing information, and QoS profiles.

- **PCF manages policy rules**
- **PCF manages network slicing information**
- **PCF manages QoS profiles**

Factoid: PCF may connect to UDR for policy rule storage.

- **PCF may connect to UDR**
- **UDR stores policy rules**

Factoid: NRF provides service discovery and registration for the Service-Based Architecture (SBA).

- **NRF provides service discovery**
- **NRF provides registration**
- **Service-Based Architecture (SBA) relies on NRF**

Factoid: NRF enables dynamic discovery of NF services using API lookup.

- **NRF enables dynamic discovery of NF services**
- **Dynamic discovery of NF services uses API lookup**

Factoid: NRF ensures service-level routing and load-aware NF selection.

- **NRF ensures service-level routing**
- **NRF ensures load-aware NF selection**

Factoid: NRF uses the Nnrf interface (NF-NRF) over HTTP/2.

- **NRF uses Nnrf interface**
- **Nnrf interface connects network functions (NFs) and NRF**
- **Nnrf interface runs over HTTP/2**

Factoid: NRF manages data such as NF profiles, service status, and availability.

- **NRF manages NF profiles**
- **NRF manages service status**
- **NRF manages availability**

Factoid: NRF is a dependency for all SBA-based network functions.

- **NRF is a dependency for all SBA-based network functions**

Factoid: UDR acts as a centralized database backend for UDM and PCF.

- **UDR acts as centralized database backend**
- **UDR serves as backend for UDM**
- **UDR serves as backend for PCF**

Factoid: UDR decouples business logic (UDM/PCF) from persistent storage.

- **UDR decouples business logic** from **persistent storage**
- **Business logic includes UDM**
- **Business logic includes PCF**

Factoid: UDR improves fault tolerance and scaling of data-centric functions.

- **UDR improves fault tolerance**
- **UDR improves scaling**
- **Fault tolerance applies to data-centric functions**
- **Scaling applies to data-centric functions**

Factoid: UDR uses the N5 interface (UDM/PCF-UDR).

- **UDR uses N5 interface**
- **N5 interface connects UDM and UDR**
- **N5 interface connects PCF and UDR**

Factoid: UDR manages data such as subscriber data, policy data, and configuration values.

- **UDR manages subscriber data**
- **UDR manages policy data**
- **UDR manages configuration values**

Factoid: PFCP is the standardized protocol for SMF-to-UPF control, operating over UDP at the N4 and Sx reference points.

- **PFCP is standardized protocol for SMF-to-UPF control**
- **PFCP operates over UDP**
- **PFCP operates at N4 reference point**
- **PFCP operates at Sx reference point**
- **SMF-to-UPF control uses PFCP**

Factoid: SMF uses PFCP to install PFDs and FARs in UPF to govern user data forwarding rules.

- **SMF uses PFCP**
- **SMF installs PFDs in UPF**
- **SMF installs FARs in UPF**
- **PFDs govern user data forwarding rules**
- **FARs govern user data forwarding rules**

Factoid: PFCP is fundamental for enabling Control-User Plane Separation (CUPS) in 5G core.

- **PFCP is fundamental for enabling Control-User Plane Separation (CUPS)**
- **Control-User Plane Separation (CUPS) is applied in 5G core**

Factoid: PFCP includes mechanisms like Keepalive, Session, QER, and BAR to manage forwarding and buffering.

- **PFCP includes Keepalive mechanism**
- **PFCP includes Session mechanism**
- **PFCP includes QER mechanism**
- **PFCP includes BAR mechanism**
- **Keepalive mechanism manages forwarding**
- **Keepalive mechanism manages buffering**
- **Session mechanism manages forwarding**
- **Session mechanism manages buffering**
- **QER mechanism manages forwarding**
- **QER mechanism manages buffering**
- **BAR mechanism manages forwarding**
- **BAR mechanism manages buffering**

Factoid: PFCP enables dynamic session lifecycle control by allowing SMF to initiate, modify, or tear down UPF forwarding behavior in real time.

- **PFCP enables dynamic session lifecycle control**
- **PFCP allows SMF to initiate UPF forwarding behavior in real time**
- **PFCP allows SMF to modify UPF forwarding behavior in real time**
- **PFCP allows SMF to tear down UPF forwarding behavior in real time**

Factoid: Through rule-based provisioning, PFCP allows SMF to granularly instruct UPF on how to classify, buffer, or prioritize specific traffic flows.

- **SMF uses PFCP;**
- **SMF instructs UPF;**
- **UPF classifies SpecificTrafficFlows.**
- **UPF buffers SpecificTrafficFlows.**
- **UPF prioritizes SpecificTrafficFlows.**

Factoid: PFCP sessions maintain stateful associations between SMF and UPF, ensuring coherent control across multiple user sessions.

- **PFCP sessions maintain stateful associations between SMF and UPF**
- **Stateful associations between SMF and UPF ensure coherent control across multiple user sessions**

Factoid: The PFCP protocol supports heartbeats and recovery procedures to detect failure conditions and maintain high-availability pathways.

- **PFCP protocol supports heartbeats**
- **PFCP protocol supports recovery procedures**
- **Heartbeats detect failure conditions**
- **Recovery procedures detect failure conditions**
- **Heartbeats maintain high-availability pathways**
- **Recovery procedures maintain high-availability pathways**

Factoid: By separating control and data responsibilities, PFCP allows distributed UPF instances to be managed centrally by fewer SMF nodes.

- **PFCP allows central management of distributed UPF instances**
- **Central management of distributed UPF instances is performed by fewer SMF nodes**

Factoid: Advanced rule types in PFCP such as QER and BAR enable precise QoS enforcement and packet buffering, tailored to per-session policy.

- **Advanced rule types in PFCP include QER**
- **Advanced rule types in PFCP include BAR**
- **Advanced rule types in PFCP enable precise QoS enforcement**
- **Advanced rule types in PFCP enable packet buffering**
- **Precise QoS enforcement is tailored to per-session policy**
- **Packet buffering is tailored to per-session policy**

Factoid: PFCP supports scalable deployment by minimizing control overhead using compact rule structures and stateless UDP transport.

- **PFCP supports scalable deployment**
- **PFCP minimizes control overhead**
- **Control overhead is minimized using compact rule structures**
- **Control overhead is minimized using stateless UDP transport**

Factoid: PFCP is specified in 3GPP TS 29.244 and relies on UDP transport for control messaging.

- **PFCP is specified in 3GPP TS 29.244**
- **PFCP relies on UDP transport**
- **UDP transport carries control messaging**

Factoid: SEPP encrypts and proxies inter-PLMN control signaling across roaming borders.

- **SEPP encrypts inter-PLMN control signaling**
- **SEPP proxies inter-PLMN control signaling**
- **Inter-PLMN control signaling crosses roaming borders**

Factoid: SEPP ensures end-to-end control-plane security via TLS/IPsec tunnels.

- **SEPP ensures end-to-end control-plane security**
- **SEPP uses TLS/IPsec tunnels**
- **TLS/IPsec tunnels provide end-to-end control-plane security**

Factoid: N3IWF acts as a gateway for non-3GPP access, such as Wi-Fi, using IKEv2 tunnels.

- **N3IWF acts as gateway for non-3GPP access**
- **Non-3GPP access includes Wi-Fi**
- **N3IWF uses IKEv2 tunnels**

Factoid: N3IWF bridges Wi-Fi clients to the AMF securely, without requiring native RAN.

- **N3IWF bridges Wi-Fi clients to AMF securely**
- **N3IWF-mediated Wi-Fi-AMF connection does not require native RAN**

Factoid: W-AGF is the Wireline Access Gateway Function used for fixed/mobile network convergence.

- **W-AGF is Wireline Access Gateway Function**
- **Wireline Access Gateway Function is used for fixed/mobile network convergence**

Factoid: W-AGF enables fixed wireless access (FWA) and fixed-line services to connect to the 5G Core (5GC).

- **W-AGF enables fixed wireless access (FWA)**
- **W-AGF enables fixed-line services**
- **Fixed wireless access (FWA) connects to 5G Core (5GC)**
- **Fixed-line services connect to 5G Core (5GC)**

Factoid: 5G core employs SBA where NFs like AMF, SMF, UPF register/discover services via NRF using HTTP/2 + REST.

- **5G core employs Service-Based Architecture (SBA)**
- **NFs (AMF, SMF, UPF) register services via NRF**
- **NFs (AMF, SMF, UPF) discover services via NRF**
- **Service registration/discovery via NRF uses HTTP/2 + REST**

Factoid: SBA provides modularity, enabling dynamic NF registration and vendor interoperability.

- **Service-Based Architecture (SBA) provides modularity**
- **Modularity enables dynamic NF registration**
- **Modularity enables vendor interoperability**

Factoid: NSA combines 5G NR radio with 4G EPC core, while true SA uses 5GC supporting slicing, URLLC, mMTC.

- **NSA (Non-Standalone architecture) combines 5G NR radio**
- **NSA (Non-Standalone architecture) combines 4G EPC core**
- **SA (Standalone architecture) uses 5GC**
- **5GC supports network slicing**
- **5GC supports URLLC**
- **5GC supports mMTC**

Factoid: Cybersecurity in 5G requires layered defenses across virtualization, SBA APIs, RAN, and network slices.

- **Cybersecurity in 5G requires layered defenses**
- **Layered defenses span virtualization**

- **Layered defenses** *span SBA APIs*
- **Layered defenses** *span RAN*
- **Layered defenses** *span network slices*

Factoid: Standard SBA NFs include AMF, SMF, UPF, UDM, PCF; also roaming entities SEPP, N3IWF, W-AGF.

- **Service-Based Architecture (SBA)** *includes AMF*
- **Service-Based Architecture (SBA)** *includes SMF*
- **Service-Based Architecture (SBA)** *includes UPF*
- **Service-Based Architecture (SBA)** *includes UDM*
- **Service-Based Architecture (SBA)** *includes PCF*
- **Service-Based Architecture (SBA)** *includes SEPP -> roaming entity*
- **Service-Based Architecture (SBA)** *includes N3IWF -> roaming entity*
- **Service-Based Architecture (SBA)** *includes W-AGF -> roaming entity*

Factoid: AMF can operate statelessly by externalizing UE context to a Unified Data Storage Function (UDSF), enhancing horizontal scalability.

- **AMF** *can operate statelessly*
- **AMF** *externalizes UE context to Unified Data Storage Function (UDSF)*
- **Externalizing UE context to UDSF** *enhances horizontal scalability*

Factoid: SMF uses the N7 interface to retrieve policy decisions from PCF, enabling differentiated handling per session.

- **SMF** *uses N7 interface*
- **SMF** *retrieves policy decisions from PCF*
- **Policy decisions** *enable differentiated per-session handling*

Factoid: UPF maintains stateful session information to perform precise flow routing and per-user QoS enforcement.

- **UPF** *maintains stateful session information*
- **Stateful session information** *enables precise flow routing*
- **Stateful session information** *enables per-user QoS enforcement*

Factoid: NRF supports load-aware NF instance selection by maintaining up-to-date service availability data.

- **NRF** *supports load-aware NF instance selection*
- **NRF** *maintains up-to-date service availability data*
- **Maintaining up-to-date service availability data** *enables load-aware NF instance selection*

Factoid: AUSF leverages HTTP/2 interfaces to securely interact with UDM for retrieving authentication keys and vectors.

- **AUSF** *leverages HTTP/2 interfaces*
- **AUSF** *securely interacts with UDM*

- **Secure interaction with UDM** *retrieves authentication keys and vectors*

Factoid: UDM is a stateful front-end to subscriber data, often backed by UDR for persistence and redundancy.

- **UDM** *is stateful front-end to subscriber data*
- **UDM** *is backed by UDR*
- **UDR** *provides persistence*
- **UDR** *provides redundancy*

Factoid: UDR acts as a persistent, central repository supporting both subscriber data (UDM) and policy data (PCF), facilitating decoupling.

- **UDR** *acts as persistent central repository*
- **UDR** *supports subscriber data for UDM*
- **UDR** *supports policy data for PCF*
- **UDR's support for subscriber and policy data** *facilitates decoupling*

Factoid: PCF is inherently stateless and offloads data persistence to UDR, allowing lightweight and scalable deployment.

- **PCF** *is inherently stateless*
- **PCF** *offloads data persistence to UDR*
- **Offloading data persistence to UDR** *allows lightweight deployment*
- **Offloading data persistence to UDR** *allows scalable deployment*

Factoid: NSSF enables network slicing by selecting appropriate slice instances based on policies and UE subscription data.

- **NSSF** *enables network slicing*
- **NSSF** *selects slice instances*
- **Slice instance selection** *is based on policies*
- **Slice instance selection** *is based on UE subscription data*

Factoid: SEPP enforces inter-PLMN security by encrypting control-plane messages using TLS/IPsec tunnels, while hiding internal NF topology.

- **SEPP** *enforces inter-PLMN security*
- **SliceSEPP** *encrypts control-plane messages*
- **SEPP** *uses TLS/IPsec tunnels*
- **SEPP** *hides internal NF topology*

Factoid: N3IWF provides secure IPsec tunnels over Wi-Fi, allowing non-3GPP devices to attach directly to the 5G core via AMF.

- **N3IWF** *provides secure IPsec tunnels*
- **Secure IPsec tunnels** *operate over Wi-Fi*
- **N3IWF** *allows non-3GPP devices to attach directly to 5G core*
- **AMF** *facilitates non-3GPP device attachment to 5G core*

Factoid: SBA enables all control-plane NFs to expose RESTful APIs over HTTP/2, facilitating microservice-based orchestration and interop.

- **Service-Based Architecture (SBA)** *enables control-plane NFs to expose RESTful APIs*
- **Control-plane NFs** *expose RESTful APIs over HTTP/2*
- **Exposing RESTful APIs over HTTP/2** *facilitates microservice-based orchestration*
- **Exposing RESTful APIs over HTTP/2** *facilitates interoperability*

Factoid: The choice between stateful and stateless NF design impacts how context is managed, with stateless NFs enabling container-based scaling.

- **Stateful NF design** *impacts context management*
- **Stateless NF design** *impacts context management*
- **Stateless NF design** *enables container-based scaling*

Factoid: Stateless core NFs like AMF and SMF require external state repositories (e.g., UDSF, UDR) to restore context during horizontal scaling or failure recovery.

- **Stateless core NFs** *require external state repositories*
- **AMF** *requires external state repositories*
- **SMF** *requires external state repositories*
- **External state repositories** *include UDSF*
- **External state repositories** *include UDR*
- **External state repositories** *restore context during horizontal scaling*
- **External state repositories** *restore context during failure recovery*

Factoid: NRF and NSSF are stateless; AMF and SMF are stateful and store UE context in UDSF.

- **NRF** *is stateless*
- **NSSF** *is stateless*
- **AMF** *is stateful*
- **AMF** *stores UE context in UDSF*
- **SMF** *is stateful*
- **SMF** *stores UE context in UDSF*

Factoid: UDSF functions as a shared key-value store holding UE session context and security credentials.

- **UDSF** *functions as shared key-value store*
- **Shared key-value store** *holds UE session context*
- **Shared key-value store** *holds security credentials*

Factoid: Piggyback-based retrieval retrieves all required UE context at procedure start, eliminating redundant external requests.

- **Piggyback-based retrieval** *retrieves all required UE context at procedure start*
- **Piggyback-based retrieval** *eliminates redundant external requests*

Factoid: Piggyback reduces synchronous procedure latency by 44–70%.

- **Piggyback *reduces synchronous procedure latency by 44–70 %***

Factoid: Proactive-push preloads downstream NF state at procedure initiation, cutting asynchronous latency by 13–22%.

- **Proactive-push *preloads downstream NF state at procedure initiation***
- **Proactive-push *cuts asynchronous latency by \approx 13–22 %***

Factoid: The hybrid pattern (piggyback + proactive-push) optimizes latency across synchronous and asynchronous procedures without extra overhead.

- **Hybrid pattern *combines piggyback and proactive-push***
- **Hybrid pattern *optimizes latency across synchronous procedures***
- **Hybrid pattern *optimizes latency across asynchronous procedures***
- **Hybrid pattern *incurs no extra overhead***

Factoid: Container-based stateless deployment enables robust horizontal scaling and failure resilience.

- **Container-based stateless deployment *enables robust horizontal scaling***
- **Container-based stateless deployment *enables failure resilience***

Factoid: Stateless architecture maintains 3GPP compliance by preserving standard SBI interfaces and requiring only new UDSF endpoints.

- **Stateless architecture *maintains 3GPP compliance***
- **Stateless architecture *preserves standard SBI interfaces***
- **Stateless architecture *requires only new UDSF endpoints***

Factoid: Stateful NFs like AMF and SMF can offload internal memory to UDSF to become externally state-managed without altering logical behavior.

- **AMF *offloads internal memory to UDSF***
- **SMF *offloads internal memory to UDSF***
- **Offloading internal memory to UDSF *enables externally state-managed operation***
- **Externally state-managed AMF and SMF *do not alter logical behavior***

Factoid: UDSF is optimized for UE-specific metadata such as SUPI, session state, and NAS keys, supporting fine-grained key-value access.

- **UDSF *is optimized for UE-specific metadata***
- **UE-specific metadata *includes SUPI***
- **UE-specific metadata *includes session state***
- **UE-specific metadata *includes NAS keys***
- **UDSF *supports fine-grained key-value access***

Factoid: Proactive state push allows AMF to transmit UE context to dependent NFs *before* they are involved in a procedure, enabling smoother control plane flow.

- **Proactive state push** *allows AMF to transmit UE context*
- **AMF transmits UE context** *to dependent NFs*
- **Transmitting UE context to dependent NFs before procedure involvement** *enables smoother control-plane flow*

Factoid: Piggyback and proactive strategies are not mutually exclusive—hybrid use improves end-to-end responsiveness under varied traffic patterns.

- **Piggyback strategies** *are not mutually exclusive with proactive strategies*
- **Hybrid use** *combines piggyback strategies and proactive strategies*
- **Hybrid use** *improves end-to-end responsiveness*
- **End-to-end responsiveness** *is improved under varied traffic patterns*

Factoid: Stateless NF deployment reduces the operational burden of state synchronization across instances during autoscaling or rolling updates.

- **Stateless NF deployment** *reduces operational burden of state synchronization across instances*
- **Operational burden of state synchronization** *occurs during autoscaling*
- **Operational burden of state synchronization** *occurs during rolling updates*

Factoid: Kubernetes-based control plane orchestration benefits from stateless NF design by enabling auto-replacement of crashed pods without state loss.

- **Kubernetes-based control plane orchestration** *benefits from stateless NF design*
- **Stateless NF design** *enables auto-replacement of crashed pods*
- **Auto-replacement of crashed pods** *occurs without state loss*

Factoid: No changes to 3GPP signaling protocols (e.g., NAS, NGAP) are required to adopt stateless control—only the UDSF interface and optional push hooks are added.

- **Stateless control** *requires no changes to 3GPP signaling protocols (NAS, NGAP)*
- **Stateless control** *adds UDSF interface*
- **Stateless control** *adds optional push hooks*

Factoid: UDSF enables centralized or distributed state storage strategies depending on network topology, aiding deployment in edge or cloud regions.

- **UDSF** *enables centralized state storage strategies*
- **UDSF** *enables distributed state storage strategies*
- **Centralized state storage strategies** *depend on network topology*
- **Distributed state storage strategies** *depend on network topology*
- **Centralized state storage strategies** *aid deployment in edge regions*
- **Centralized state storage strategies** *aid deployment in cloud regions*
- **Distributed state storage strategies** *aid deployment in edge regions*
- **Distributed state storage strategies** *aid deployment in cloud regions*

Factoid: Stateless NF instances may exhibit lower average CPU usage than stateful ones, due to wait periods during external state retrieval.

- **Stateless NF instances** *exhibit lower average CPU usage than stateful NF instances*
- **Lower average CPU usage** *is due to wait periods during external state retrieval*
- **Wait periods during external state retrieval** *occur in stateless NF instances*

Factoid: Queue buildup in AMF propagates to SMF and UPF as CPU spikes once state responses are processed.

- **Queue buildup in AMF** *propagates to SMF*
- **Queue buildup in AMF** *propagates to UPF*
- **SMF experiences CPU spikes** *once state responses are processed*
- **UPF experiences CPU spikes** *once state responses are processed*

Factoid: Transactional stateless NFs increase latency and cost by performing synchronous state fetches and JSON deserialization.

- **Transactional stateless NFs** *increase latency*
- **Transactional stateless NFs** *increase cost*
- **Transactional stateless NFs** *perform synchronous state fetches*
Transactional stateless NFs *perform JSON deserialization*

Factoid: Non-blocking stateless strategies improve latency and throughput by asynchronous state access.

- **Non-blocking stateless strategies** *improve latency*
- **Non-blocking stateless strategies** *improve throughput*
- **Non-blocking stateless strategies** *use asynchronous state access*

Factoid: Working in cloud environments, stateless NF designs increase billing costs because longer request times increase runtime charges.

- **Stateless NF designs** *increase billing costs*
- **Stateless NF designs** *operate in cloud environments*
Stateless NF designs *cause longer request times*
- **Longer request times** *increase runtime charges*
- **Increased runtime charges** *lead to higher billing costs*

Factoid: Sharing global UE context between NFs reduces unnecessary database operations, improving performance by ~33%.

- **Global UE context** *is shared between NFs*
- **Sharing global UE context between NFs** *reduces unnecessary database operations*
- **Sharing global UE context between NFs** *improves performance by $\approx 33\%$*

Factoid: Embedding user context into NF-to-NF messages cuts DB reads from $4 \times n$ to 2, reducing overhead by ~22%.

- **Embedding user context into NF-to-NF messages** *cuts* DB reads from $4 \times n$ to 2
- **Cutting DB reads from $4 \times n$ to 2** *reduces* overhead by $\approx 22\%$

Factoid: Stateless NFs allow container-based scaling but require optimized caching strategies to avoid latency and cost penalties.

- **Stateless NFs** *allow* container-based scaling
- **Stateless NFs** *require* optimized caching strategies
- **Optimized caching strategies** *avoid* latency penalties
- **Optimized caching strategies** *avoid* cost penalties

Factoid: Stateless NF designs can create a misleading CPU usage profile where underutilized compute masks backend bottlenecks in state access layers.

- **Stateless NF designs** *can create* misleading CPU usage profile
- **Stateless NF designs** *can lead to* underutilized compute
- **Underutilized compute** *masks* backend bottlenecks
- **Misleading CPU usage profile** *masks* backend bottlenecks in state access layers

Factoid: Latency spikes in stateless control chains often correlate with batched state-response arrivals, leading to temporary CPU contention across NF tiers.

- **Latency spikes** *occur in* stateless control chains
- **Latency spikes** *correlate with* batched state-response arrivals
- **Batched state-response arrivals** *lead to* temporary CPU contention
- **Temporary CPU contention** *occurs across* NF tiers

Factoid: Frequent state serialization and deserialization in transactional stateless models leads to inflated processing delay, especially under high concurrency.

- **Transactional stateless models** *perform* frequent state serialization and deserialization
- **Frequent state serialization and deserialization** *leads to* inflated processing delay
- **Inflated processing delay** *is exacerbated under* high concurrency

Factoid: Non-blocking stateless strategies leverage concurrency and pipelining to maintain throughput while masking state-fetch latency.

- **Non-blocking stateless strategies** *leverage* concurrency
- **Non-blocking stateless strategies** *leverage* pipelining
- **Non-blocking stateless strategies** *maintain* throughput
- **Non-blocking stateless strategies** *mask* state-fetch latency

Factoid: Stateless deployment models require broader horizontal scaling to achieve parity with stateful throughput, despite lower CPU use per instance.

- **Stateless deployment models** *require broader horizontal scaling*
- **Broader horizontal scaling** *achieves parity with stateful throughput*
- **Stateless deployment models** *exhibit lower CPU use per instance*

Factoid: Cloud-native stateless NFs must be capacity-planned for bursty workloads, as idle periods followed by sudden spikes are common in real-world flows.

- **Cloud-native stateless NFs** *must be capacity-planned for bursty workloads*
- **Bursty workloads** *include idle periods*
- **Bursty workloads** *include sudden spikes*
- **Idle periods** *are followed by sudden spikes*
- **Idle periods** *are common in real-world flows*
- **Sudden spikes** *are common in real-world flows*

Factoid: State-sharing optimizations across service chains can significantly cut control-plane I/O without sacrificing NF modularity.

- **State-sharing optimizations** *occur across service chains*
- **State-sharing optimizations** *significantly cut control-plane I/O*
- **State-sharing optimizations** *maintain NF modularity*

Factoid: Embedding minimal but sufficient session metadata in upstream NF messages helps reduce cumulative state-store interactions across the chain.

- **Upstream NF messages** *embed minimal session metadata*
- **Embedding minimal session metadata in upstream NF messages** *reduces cumulative state-store interactions across the chain*
- **Cumulative state-store interactions** *occur across the service chain*

Factoid: Persistent caching layers or sidecars that prefetch or retain context between NF invocations offer critical improvements in latency and cost-efficiency.

- **Persistent caching layers** *prefetch context between NF invocations*
- **Persistent caching layers** *retain context between NF invocations*
- **Sidecars** *prefetch context between NF invocations*
- **Sidecars** *retain context between NF invocations*
- **Prefetching or retaining context between NF invocations** *offers critical improvements in latency*
- **Prefetching or retaining context between NF invocations** *offers critical improvements in cost-efficiency*

Factoid : Quasi-local model uses fetch-and-cache strategy, storing UE state locally after initial procedures.

- **Quasi-local model** *uses fetch-and-cache strategy*
- **Fetch-and-cache strategy** *stores UE state locally*

- **UE state** *is stored after initial procedures*

Factoid: Decoupling control compute from storage enables fast NF instance recovery mid-session.

- **Decoupling control compute from storage** *enables fast NF instance recovery*
Fast NF instance recovery *occurs mid-session*

Factoid : Per-procedure caching optimizes storage accesses and controls DB IO per NF.

- **Per-procedure caching** *optimizes storage accesses*
- **Per-procedure caching** *controls DB IO per NF*

Factoid: Quasi-local cache supports latency-critical use cases like V2X and telesurgery.

- **Quasi-local cache** *supports latency-critical use cases*
- **Latency-critical use cases** *include V2X*
- **Latency-critical use cases** *include telesurgery*

Factoid: State metrics include volume, size, and frequency of operations; helps dimension datastore loads.

- **State metrics** *include volume*
- **State metrics** *include size*
- **State metrics** *include frequency of operations*
- **State metrics** *help dimension datastore loads*

Factoid: Quasi-local caching reduces network-wide datastore interactions during user-plane processing.

- **Quasi-local caching** *reduces network-wide datastore interactions*
- **Network-wide datastore interactions** *occur during user-plane processing*

Factoid: The quasi-local model maintains a session-scoped cache that reduces repeated external lookups during procedures like registration and handover.

- **Quasi-local model** *maintains session-scoped cache*
- **Session-scoped cache** *reduces repeated external lookups*
- **Repeated external lookups** *occur during registration procedures*
- **Repeated external lookups** *occur during handover procedures*

Factoid: Caching state at well-defined checkpoints ensures that NFs have immediate access to necessary context during bursty signaling events.

- **Caching state at well-defined checkpoints** *ensures immediate access to necessary context*
- **Immediate access to necessary context** *is provided to NFs during bursty signaling events*

Factoid: Separating compute from persistent state allows newly spawned NF instances to rejoin active sessions without full protocol reinitialization.

- **Separation of compute from persistent state** *allows newly spawned NF instances to rejoin active sessions*
- **Separation of compute from persistent state** *eliminates full protocol reinitialization*

Factoid: Mid-session recovery is enabled by fetching the most recent session snapshot from a distributed state store, avoiding session resets.

- **Mid-session recovery** *is enabled by fetching the most recent session snapshot*
- **Most recent session snapshot** *is stored in distributed state store*
- **Fetching the most recent session snapshot** *avoids session resets*

Factoid: Auto-persistence strategies allow caching to be dynamically tailored per procedure type, optimizing both performance and memory footprint.

- **Auto-persistence strategies** *allow caching to be dynamically tailored per procedure type*
- **Auto-persistence strategies** *optimize performance*
- **Auto-persistence strategies** *optimize memory footprint*

Factoid: Procedure-aware caching models help NFs selectively retain short-lived versus long-lived state, balancing responsiveness with memory efficiency.

- **Procedure-aware caching models** *help NFs selectively retain short-lived state*
- **Procedure-aware caching models** *help NFs selectively retain long-lived state*
- **Selective retention of short-lived state** *balances responsiveness*
- **Selective retention of long-lived state** *balances memory efficiency*

Factoid: The fetch-and-cache model minimizes control-plane latency variance by reducing dependency on synchronous state-store reads during critical flows.

- **Fetch-and-cache model** *minimizes control-plane latency variance*
- **Fetch-and-cache model** *reduces dependency on synchronous state-store reads*
- **Dependency on synchronous state-store reads** *occurs during critical flows*

Factoid: Tailored cache lifetimes and eviction policies allow NFs to meet reliability requirements across diverse traffic types, including eMBB and IoT.

- **Tailored cache lifetimes** *allow NFs to meet reliability requirements*
- **Eviction policies** *allow NFs to meet reliability requirements*
- **Reliability requirements** *apply across diverse traffic types*
- **Diverse traffic types** *include eMBB*
- **Diverse traffic types** *include IoT*

Factoid: N1 is the control-plane interface for NAS signaling between UE and AMF, across both gNB and access network.

- **N1 interface** *is control-plane interface*
- **N1 interface** *carries NAS signaling*
NAS signaling *occurs between UE and AMF*
- **N1 interface** *operates across gNB*
- **N1 interface** *operates across access network*

Factoid: N1 NAS messages are transparently forwarded by gNB without processing.

- **gNB** *transparently forwards N1 NAS messages*
- **gNB** *does not process N1 NAS messages*

Factoid: In 3GPP access, N1 uses RRC over Uu and NGAP over N2 between UE and AMF.

- **N1 interface** *uses RRC*
- **RRC** *runs over Uu*
- **N1 interface** *uses NGAP*
- **NGAP** *runs over N2*
- **N1 interface** *operates between UE and AMF*

Factoid: In non-3GPP access, N1 includes an IPsec tunnel (NWu) to link UE to AMF securely.

- **N1 interface** *includes IPsec tunnel (NWu)*
- **N1 interface** *uses IPsec tunnel (NWu) in non-3GPP access*
- **IPsec tunnel (NWu)** *links UE and AMF securely*
- **IPsec tunnel (NWu)** *is used in non-3GPP access*

Factoid: N1 mode denotes Standalone (SA) deployment with direct 5G Core connectivity to UE, whereas S1 mode uses NSA via 4G.

- **N1 mode** *denotes Standalone (SA) deployment*
- **Standalone (SA) deployment** *connects UE to 5G Core*
- **S1 mode** *uses Non-Standalone (NSA) deployment*
- **NSA deployment** *operates via 4G*

Factoid: N1 interface carries NAS signaling transparently via gNB, using RRC and NGAP for core-plane UE–AMF communication.

- **N1 interface** *carries NAS signaling*
- **NAS signaling** *is carried transparently via gNB*
- **N1 interface** *uses RRC for core-plane UE–AMF communication*
- **N1 interface** *uses NGAP for core-plane UE–AMF communication*

Factoid: Control-plane interfaces (N2, N11, N12, N14, N15) use HTTP/2 REST over SBA, while data-plane interfaces (N3, N6, N9) use UDP-based protocols.

- **Control-plane interfaces** *use HTTP/2 REST*
- **N2** *is a Control-plane interface*

- **N11 is a Control-plane interface**
- **N12 is a Control-plane interface**
- **N14 is a Control-plane interface**
- **N15 is a Control-plane interface**
- **HTTP/2 REST runs over Service-Based Architecture (SBA)**
- **Data-plane interfaces use UDP-based protocols**
- **N3 is a Data-plane interface**
- **N6 is a Data-plane interface**
- **N9 is a Data-plane interface**
- **N15 is a Control-plane interface**

Factoid: N4 interface (SMF–UPF) employs PFCP to manage user-plane sessions and forwarding behavior.

- **N4 interface connects SMF and UPF**
- **N4 interface employs PFCP**
- **PFCP manages user-plane sessions**
- **PFCP manages forwarding behavior**

Factoid: Inter-NF service calls (e.g. N7, N8, N10) enable SMB workload routing and allow dynamic policy and data flow management.

- **Inter-NF service calls enable SMB workload routing**
- **Inter-NF service calls allow dynamic policy management**
- **Inter-NF service calls allow dynamic data flow management**
- **N7 interface is an Inter-NF service call**
- **N8 interface is an Inter-NF service call**
- **N10 interface is an Inter-NF service call**

_____ BATCH 1 TO GEMINI _____

Factoid: Roaming-related interfaces such as N16, N27, N32 facilitate control-plane continuity across administrative domains

- **Roaming-related interfaces facilitate control-plane continuity across administrative domains**
- **N16 interface is a roaming-related interface**
- **N27 interface is a roaming-related interface**
- **N32 interface is a roaming-related interface**
- **Control-plane continuity occurs across administrative domains**

Factoid: NSSF selects network slices via N22 after UE–AMF registration for slice-specific NF association.

- **NSSF selects network slices**
- **NSSF uses N22 interface**
- **Slice-specific NF association occurs after UE–AMF registration**

Factoid: N9 supports cascading of UPFs for scaling and multi-hop data-plane routing.

- **N9 interface** *supports cascading of UPFs*
- **Cascading of UPFs** *enables scaling*
- **Cascading of UPFs** *enables multi-hop data-plane routing*

Factoid: The interface structure supports cloud-native scaling via clear separation and SLA isolation.

- **Interface structure** *supports cloud-native scaling*
- **Interface structure** *provides clear separation*
- **Interface structure** *provides SLA isolation*

Factoid: Security features across interfaces include TLS for HTTP/2 SBA and IPsec tunneling for non-3GPP and RAN access paths.

- **Security features across interfaces** *include TLS for HTTP/2 SBA*
- **Security features across interfaces** *include IPsec tunneling for non-3GPP access paths*
- **Security features across interfaces** *include IPsec tunneling for RAN access paths*

Factoid: N1 carries NAS signaling between the UE and AMF and is transparent through the gNB, forming the control link for UE registration and session setup.

- **N1** *carries NAS signaling*
- **NAS signaling** *flows between UE and AMF*
- **NAS signaling** *is transparent through gNB*
- **N1** *forms control link for UE registration and session setup*

Factoid: N2 connects gNB to AMF using NGAP over SCTP and enables the transmission of access signaling and mobility control messages.

- **N2** *connects gNB to AMF*
- **N2** *uses NGAP over SCTP*
- **N2** *enables transmission of access signaling*
- **N2** *enables transmission of mobility control messages*

Factoid: N3 transports user-plane traffic between gNB and UPF using GTP-U over UDP, supporting high-throughput PDU sessions.

- **N3** *transports user-plane traffic between gNB and UPF*
- **N3** *uses GTP-U over UDP*
- **Transporting user-plane traffic via N3 using GTP-U over UDP** *supports high-throughput PDU sessions*

Factoid: N4 allows the SMF to control UPF behavior using PFCP, setting up and modifying forwarding rules and QoS enforcement.

- **N4 allows SMF to control UPF behavior**
- **SMF uses PFCP to set up forwarding rules**
- **SMF uses PFCP to modify forwarding rules**
- **PFCP enables QoS enforcement**

Factoid: N5 enables application functions (AF) to communicate policy triggers to the PCF via HTTP/2, often for content-aware optimization.

- **N5 interface enables application functions (AF)**
- **Application functions (AF) communicate policy triggers to PCF**
- **Communication of policy triggers to PCF occurs via HTTP/2**
- **Policy triggers support content-aware optimization**

Factoid: N6 links the UPF to external data networks (DN) over IP and handles routing, NAT, and service exposure for user-plane packets.

- **N6 links UPF to external data networks (DN) over IP**
- **N6 handles routing for user-plane packets**
- **N6 handles NAT for user-plane packets**
- **N6 handles service exposure for user-plane packets**

Factoid: N7 interface is used by the SMF to retrieve policy rules and QoS profiles from the PCF using RESTful HTTP APIs.

- **N7 interface is used by SMF**
- **SMF retrieves policy rules**
- **SMF retrieves QoS profiles**
- **Retrieval of policy rules and QoS profiles occurs via N7 interface**
- **N7 interface uses RESTful HTTP APIs**

Factoid: N8 provides the AMF access to subscriber and authentication data by querying the UDM over HTTP/2.

- **N8 interface provides AMF access to subscriber data**
- **N8 interface provides AMF access to authentication data**
- **AMF access to subscriber and authentication data occurs by querying UDM**
- **Querying UDM uses HTTP/2**

Factoid: N9 enables inter-UPF communication via GTP-U for distributed user-plane routing, load balancing, or service chaining.

- **N9 interface enables inter-UPF communication**
- **N9 interface uses GTP-U**
- **Inter-UPF communication supports distributed user-plane routing**
- **Inter-UPF communication supports load balancing**
- **Inter-UPF communication supports service chaining**

Factoid: N10 allows SMF to obtain subscriber policy data from the UDM, including PDU session authorization and QoS configuration.

- **N10 interface** *allows SMF to obtain subscriber policy data*
- **SMF** *obtains subscriber policy data from UDM*
- **Subscriber policy data** *includes PDU session authorization*
- **Subscriber policy data** *includes QoS configuration*

Factoid: N11 connects the AMF and SMF, enabling control-plane interactions such as PDU session setup and mobility handling.

- **N11 interface** *connects AMF and SMF*
- **N11 interface** *enables control-plane interactions*
- **Control-plane interactions** *include PDU session setup*
- **Control-plane interactions** *include mobility handling*

Factoid: N12 allows the AMF to forward authentication requests to the AUSF during UE registration or security mode command procedures.

- **N12 interface** *allows AMF to forward authentication requests to AUSF*
- **AMF** *forwards authentication requests to AUSF*
- **Forwarding authentication requests to AUSF** *occurs during UE registration*
- **Forwarding authentication requests to AUSF** *occurs during security mode command procedures*

Factoid: N13 connects AUSF to UDM for obtaining authentication vectors like AV/EAP for 5G-AKA procedures.

- **N13 interface** *connects AUSF and UDM*
- **AUSF** *obtains authentication vectors from UDM*
- **Authentication vectors** *include AV*
- **Authentication vectors** *include EAP*
- **Authentication vectors** *are used for 5G-AKA procedures*

Factoid: N14 supports AMF-to-AMF communication for inter-region handover, UE context transfer, and mobility continuity.

- **N14** *supports AMF-to-AMF communication*
- **AMF-to-AMF communication** *enables inter-region handover*
- **AMF-to-AMF communication** *enables UE context transfer*
- **AMF-to-AMF communication** *enables mobility continuity*

Factoid: N15 is used by the PCF to supply policy rules directly to the AMF, influencing access and mobility behaviors.

- **N15 interface** *is used by PCF*
- **PCF** *supplies policy rules to AMF*
- **Supplying policy rules to AMF** *influences access behaviors*
- **Supplying policy rules to AMF** *influences mobility behaviors*

Factoid: N16 links two SMFs (typically in roaming scenarios) to coordinate session continuity and policy across PLMN boundaries.

- **N16 links two SMFs**
- **Linking two SMFs occurs in roaming scenarios**
- **N16 coordinates session continuity across PLMN boundaries**
- **N16 coordinates policy across PLMN boundaries**

Factoid: N1 is the NAS signaling interface between UE and AMF, used for registration, authentication, and mobility procedures.

- **N1 interface carries NAS signaling** between UE and AMF
- **N1 interface is used for registration procedures**
- **N1 interface is used for authentication procedures**
- **N1 interface is used for mobility procedures**

Factoid: N1 signaling is bidirectional, encapsulated in RRC at the UE side and transported via NGAP through gNB to AMF.

- **N1 signaling is bidirectional**
- **N1 signaling is encapsulated in RRC at the UE side**
- **RRC encapsulation of N1 signaling occurs at UE side**
- **N1 signaling is transported via NGAP through gNB to AMF**
- **NGAP through gNB transports N1 signaling to AMF**

Factoid: N2 connects the gNB and AMF, using NGAP over SCTP to handle UE context and handover signaling.

- **N2 connects gNB and AMF**
- **N2 uses NGAP over SCTP**
- **NGAP over SCTP handles UE context**
- **NGAP over SCTP handles handover signaling**

Factoid: N2 transmits messages such as Initial UE Message and UE Context Release between RAN and Core.

- **N2 transmits messages such as Initial UE Message**
- **N2 transmits messages such as UE Context Release**
- **Initial UE Message is transmitted between RAN and Core**
- **UE Context Release is transmitted between RAN and Core**

Factoid: N3 uses GTP-U over UDP to forward user-plane traffic between gNB and UPF after PDU session establishment.

- **N3 uses GTP-U over UDP**
- **N3 forwards user-plane traffic between gNB and UPF**
- **Forwarding user-plane traffic between gNB and UPF occurs after PDU session establishment**

Factoid: N3 is unidirectional for data flow (uplink/downlink) and not involved in control signaling.

- **N3 is unidirectional for uplink data flow**
- **N3 is unidirectional for downlink data flow**
- **N3 is not involved in control signaling**

Factoid: N4 is the control interface between SMF and UPF, enabling session creation, QoS enforcement, and traffic routing via PFCP.

- **N4 is control interface between SMF and UPF**
- **N4 enables session creation**
- **N4 enables QoS enforcement**
- **N4 enables traffic routing via PFCP**

Factoid: N4 manages forwarding rules through PFCP messages such as FAR, QER, and BAR to control UPF behavior dynamically.

- **N4 manages forwarding rules**
- **N4 uses PFCP message FAR**
- **N4 uses PFCP message QER**
- **N4 uses PFCP message BAR**
- **Managing forwarding rules through FAR, QER, and BAR enables dynamic UPF behavior control**

Factoid: 5G supports simultaneous NAS sessions over cellular and Wi-Fi by establishing multiple N1 control-plane connections.

- **5G supports simultaneous NAS sessions over cellular and Wi-Fi**
- **Simultaneous NAS sessions over cellular and Wi-Fi are established by multiple N1 control-plane connections**

Factoid: UEs authenticate to 5G core over Wi-Fi using EAP-AKA' or 5G-AKA prior to NAS signaling.

- **UEs authenticate to 5G core over Wi-Fi**
- **Authentication to 5G core over Wi-Fi uses EAP-AKA'**
- **Authentication to 5G core over Wi-Fi uses 5G-AKA**
- **Authentication using EAP-AKA' or 5G-AKA occurs prior to NAS signaling**

Factoid: IPsec tunnels (NWu/NWt) using IKEv2 and EAP-5G secure NAS traffic to N3IWF/TNGF gateways.

- **IPsec tunnels (NWu/NWt) use IKEv2**
- **IPsec tunnels (NWu/NWt) use EAP-5G**
- **IPsec tunnels (NWu/NWt) secure NAS traffic to N3IWF gateways**
- **IPsec tunnels (NWu/NWt) secure NAS traffic to TNGF gateways**

Factoid: N1 over trusted Wi-Fi uses IPsec with NULL encryption to avoid double encryption while preserving link-level security.

- **N1 over trusted Wi-Fi uses IPsec with NULL encryption**
- **IPsec with NULL encryption avoids double encryption**
- **IPsec with NULL encryption preserves link-level security**

Factoid: N2 control-plane messages over non-3GPP access use NGAP/SCTP between gateway functions and AMF.

- **N2 control-plane messages occur over non-3GPP access**
- **N2 control-plane messages use NGAP/SCTP**
- **NGAP/SCTP connects gateway functions and AMF**

Factoid: TWIF enables NAS signaling for legacy Wi-Fi devices lacking EAP-5G support by acting as a NAS proxy.

- **TWIF enables NAS signaling for legacy Wi-Fi devices lacking EAP-5G support**
- **TWIF acts as NAS proxy**

Factoid: N3 user-plane traffic between Wi-Fi gateways and UPF uses GTP-U over UDP, mirroring cellular data flow.

- **N3 user-plane traffic between Wi-Fi gateways and UPF uses GTP-U over UDP**
- **N3 user-plane traffic between Wi-Fi gateways and UPF mirrors cellular data flow**

Factoid: Service-Based Architecture (SBA) exists only between 5GC control-plane NFs via HTTP/2 REST APIs.

- **Service-Based Architecture (SBA) exists only between 5GC control-plane NFs**
- **Service-Based Architecture (SBA) operates via HTTP/2 REST APIs**

Factoid: Interfaces N1, N2, N3, N4, N6, and N9 are executed outside the SBA domain using traditional control/user-plane protocols.

- **Interfaces N1, N2, N3, N4, N6, and N9 are executed outside SBA domain**
- **Interfaces N1, N2, N3, N4, N6, and N9 use traditional control/user-plane protocols**

Factoid: AMF and SMF represent separate functional domains in control-plane, enabling independent scaling and specialization.

- **AMF represents a separate functional domain in the control-plane**
- **SMF represents a separate functional domain in the control-plane**
- **Separate functional domains in the control-plane enable independent scaling**
- **Separate functional domains in the control-plane enable specialization**

Factoid: CUPS design separates control and user gateways for flexible placement and scalability in the 5G core.

- **CUPS design separates control gateways and user gateways**
- **CUPS design enables flexible placement in the 5G core**
- **CUPS design enables scalability in the 5G core**

Factoid: SBA-capable NFs are cloud-native microservices, orchestrated via container platforms with CI/CD pipelines.

- **SBA-capable NFs are cloud-native microservices**
- **SBA-capable NFs are orchestrated via container platforms**
- **Container platforms use CI/CD pipelines**

Factoid: Control-plane APIs (e.g., Nsmf, Namf) are secured via TLS, whereas user-plane protocols (e.g., GTP-U, PFCP) use UDP/SCTP.

- **Control-plane APIs include Nsmf**
- **Control-plane APIs include Namf**
- **Control-plane APIs are secured via TLS**
- **User-plane protocols include GTP-U**
- **User-plane protocols include PFCP**
- **User-plane protocols use UDP/SCTP**

Factoid: SBA's RESTful interfaces cannot directly carry UE or user-data-bound control messages (e.g., RRC or NAS).

- **SBA's RESTful interfaces cannot directly carry UE control messages**
- **SBA's RESTful interfaces cannot directly carry user-data-bound control messages**
- **UE control messages include RRC**
- **User-data-bound control messages include NAS**

Factoid: In CUPS, user-plane functions like UPF scale and deploy separately from their control-plane counterparts (SMF, PCF).

- **UPF scales separately from SMF**
- **UPF deploys separately from SMF**
- **UPF scales separately from PCF**
- **UPF deploys separately from PCF**

Factoid: Each Open5GS NF is configured via its own YAML file in `/etc/open5gs/`, specifying protocol ports, PLMN/TAC, and features.

- **Open5GS NF is configured via its own YAML file in `/etc/open5gs/`**
- **YAML file in `/etc/open5gs/` specifies protocol ports**
- **YAML file in `/etc/open5gs/` specifies PLMN/TAC**
- **YAML file in `/etc/open5gs/` specifies features**

Factoid: Changing bind addresses in NF YAML requires consistent updates across RAN configurations to establish NGAP/GTP-U links.

- **Changing bind addresses in NF YAML** *requires consistent updates across RAN configurations*
- **Consistent updates across RAN configurations** *establish NGAP/GTP-U links*

Factoid: AMF listens on NGAP SCTP port 38412; UPF listens for GTP-U on PFCP-assigned port, both reflected in NF logs.

- **AMF** *listens on NGAP SCTP port 38412*
- **UPF** *listens for GTP-U on PFCP-assigned port*
- **Port listening events** *are reflected in NF logs*

Factoid: Log entries such as `ngap_server()` and `gtp_server()` confirm the NF's protocol stack initialization and port bindings.

- `ngap_server()` *log entry confirms NF's protocol stack initialization*
- `ngap_server()` *log entry confirms port bindings*
- `gtp_server()` *log entry confirms NF's protocol stack initialization*
- `gtp_server()` *log entry confirms port bindings*

Factoid: SCTP INIT/ABORT messages in logs reveal SCTP handshake status between gNB and AMF.

- **SCTP INIT/ABORT messages** *occur in logs*
- **SCTP INIT/ABORT messages** *reveal SCTP handshake status*
- **SCTP handshake status** *is between gNB and AMF*

Factoid: Core setup creates `ogstun` TUN interface with IPv4/IPv6 subnets for UPF operations.

- **Core setup** *creates ogstun TUN interface*
- **ogstun TUN interface** *provides IPv4 subnets*
- **ogstun TUN interface** *provides IPv6 subnets*
- **ogstun TUN interface** *supports UPF operations*

Factoid: IP forwarding (`net.ipv4.ip_forward`) must be enabled to allow UE-originated packets to route to WAN.

- **IP forwarding** (`net.ipv4.ip_forward`) *must be enabled to allow UE-originated packets to route to WAN*
- **IP forwarding** (`net.ipv4.ip_forward`) *allows UE-originated packets to route to WAN*
- **UE-originated packets** *route to WAN*

Factoid: Firewall rules enforce subnet isolation (e.g., 10.45.0.0/16), blocking unauthorized access to NF services.

- **Firewall rules** *enforce subnet isolation*
- **Subnet isolation** *applies to 10.45.0.0/16*
- **Subnet isolation** *blocks unauthorized access to NF services*

Factoid: Open5GS NFs run as systemd services and can be stopped or disabled individually (e.g., `open5gs-amfd`).

- **Open5GS NFs** *run as systemd services*
- **systemd services** *can be stopped or disabled individually*
- **open5gs-amfd** *is an example of a systemd service for an Open5GS NF*

Factoid: Running only subsets of NFs is supported by stopping irrelevant services and editing YAML configs accordingly.

- **Running only subsets of NFs** *is supported by stopping irrelevant services*
- **Running only subsets of NFs** *is supported by editing YAML configs accordingly*

Factoid: In Docker/Kubernetes mode, NF bind addresses, Kubernetes service IPs, and port mappings must align with core and RAN network settings.

- **NF bind addresses** *must align with core network settings*
- **NF bind addresses** *must align with RAN network settings*
- **Kubernetes service IPs** *must align with core network settings*
- **Kubernetes service IPs** *must align with RAN network settings*
- **Port mappings** *must align with core network settings*
- **Port mappings** *must align with RAN network settings*

Factoid: Open5GS uses Docker manifests that configure `upfPublicIP` and `amfif.ip/port`, binding container network to host services.

- **Open5GS** *uses Docker manifests*
- **Docker manifests** *configure upfPublicIP*
- **Docker manifests** *configure amfif.ip/port*
- **Configuring upfPublicIP** *binds container network to host services*
- **Configuring amfif.ip/port** *binds container network to host services*

Factoid: Open5GS uses MongoDB (e.g., `mongodb://localhost/open5gs`) for stateful NF data storage like NRF, PCF, and UDR.

- **Open5GS** *uses MongoDB*
- **MongoDB** *stores stateful NF data*
- **Stateful NF data** *includes NRF*
- **Stateful NF data** *includes PCF*
- **Stateful NF data** *includes UDR*
- **MongoDB** *connects using mongodb://localhost/open5gs*

Factoid: To support service discovery and policy control, Open5GS PCF includes `dbi` configuration pointing to a MongoDB URI.

- **Open5GS PCF** *includes dbi configuration*
- **dbi configuration** *points to MongoDB URI*
- **dbi configuration** *supports service discovery*
- **dbi configuration** *supports policy control*

Factoid: Each NF (AMF, PCF, SMF, UPF) in Open5GS can expose Prometheus metrics via an HTTP server configured under **metrics:** in its YAML file.

- **AMF** *can expose Prometheus metrics via an HTTP server*
- **PCF** *can expose Prometheus metrics via an HTTP server*
- **SMF** *can expose Prometheus metrics via an HTTP server*
- **UPF** *can expose Prometheus metrics via an HTTP server*
- **Prometheus metrics** *are exposed via an HTTP server*
- **HTTP server** *is configured under metrics:*
- **metrics:** *is defined in the NF's YAML file*

Factoid: Metrics endpoints are individually defined per NF instance (e.g., **open5gs-amfd**, **open5gs-smfd1**, **open5gs-upfd2**) with distinct IP and port.

- **Metrics endpoints** *are individually defined per NF instance*
- **NF instances** *include open5gs-amfd, open5gs-smfd1, open5gs-upfd2*
- **Metrics endpoints** *have distinct IP and port*

Factoid: Prometheus scrapes **/metrics** endpoints every 10 seconds using job definitions matching NF names in **prometheus.yml**.

- **Prometheus** *scrapes /metrics endpoints every 10 seconds*
- **Prometheus** *uses job definitions in prometheus.yml*
- **Job definitions** *match NF names*

_____ BATCH 2 GEMINI _____

Factoid: Grafana connects to Prometheus as a data source to visualize NF-specific metrics like **ues_active** or **amf_session**.

- **Grafana** *connects to Prometheus as a data source*
- **Grafana** *visualizes NF-specific metrics*
- **NF-specific metrics** *include ues_active*
- **NF-specific metrics** *include amf_session*

Factoid: Multiple SMF and UPF instances can be monitored in parallel, supporting slice- or region-specific deployments.

- **Multiple SMF and UPF instances** *can be monitored in parallel*
- **Parallel monitoring** *supports slice-specific deployments*
- **Parallel monitoring** *supports region-specific deployments*

Factoid: Open5GS's build process includes `libogsmetrics` to compile `libprom` and `libmicrohttpd` support for metrics output.

- **Open5GS's build process *includes* `libogsmetrics`**
- **`libogsmetrics` *compiles* `libprom` support for metrics output**
- **`libogsmetrics` *compiles* `libmicrohttpd` support for metrics output**

Factoid: Example metrics include counters for NAS registration requests (`rm_reginitreq`) and memory/resource usage (`process_resident_memory_bytes`).

- **Example metrics *include* counters for NAS registration requests (`rm_reginitreq`)**
- **Example metrics *include* memory/resource usage (`process_resident_memory_bytes`)**

Factoid: Using distinct `job_name` entries in Prometheus enables selective scraping and dashboarding per NF type.

- **Distinct `job_name` entries in Prometheus *enable* selective scraping**
- **Distinct `job_name` entries in Prometheus *enable* dashboarding per NF type**

Factoid: The sample config binds metrics HTTP servers to host IPs, facilitating external observability for containerized NFs.

- **Sample config *binds* metrics HTTP servers to host IPs**
- **Binding metrics HTTP servers to host IPs *facilitates* external observability**
- **External observability *applies to* containerized NFs**

Factoid: Metrics integration demonstrates how Open5GS can operate as 'Prometheus-enabled' microservices in Kubernetes or Docker environments.

- **Metrics integration *demonstrates* Open5GS operating as 'Prometheus-enabled' microservices**
- **Open5GS *operates as* 'Prometheus-enabled' microservices**
- **'Prometheus-enabled' microservices *run in* Kubernetes environments**
- **'Prometheus-enabled' microservices *run in* Docker environments**

Factoid: Open5GS is installed via Ubuntu PPA and requires MongoDB for subscriber context storage.

- **Open5GS *is installed via* Ubuntu PPA**
- **Open5GS *requires* MongoDB**
- **MongoDB *stores* subscriber context**

Factoid: Hardware with Intel i5 is preferred for MongoDB compatibility over Celeron-based systems.

- **Hardware with Intel i5 *is preferred for* MongoDB compatibility**

- **Hardware with Intel i5 is preferred over Celeron-based systems**

Factoid: Default Open5GS configs use loopback IPs; to integrate with external RAN simulators, host LAN IPs must replace loopback addresses in NF YAML.

- **Default Open5GS configs use loopback IPs**
- **Integration with external RAN simulators requires replacing loopback addresses in NF YAML**
- **Replacing loopback addresses in NF YAML uses host LAN IPs**

Factoid: PLMN ID and TAC must match across core (Open5GS) and RAN (UERANSIM/gNB) configurations for successful connectivity.

- **PLMN ID must match Open5GS core configuration**
- **PLMN ID must match UERANSIM/gNB RAN configuration**
- **TAC must match Open5GS core configuration**
- **TAC must match UERANSIM/gNB RAN configuration**
- **Matching PLMN ID and TAC across configurations ensures successful connectivity**

Factoid: UERANSIM `open5gs-gnb.yaml` must specify `linkIp`, `ngapIp`, `gtpIp`, and core AMF address to enable N2 and N3 connectivity.

- **UERANSIM `open5gs-gnb.yaml` must specify `linkIp`**
- **UERANSIM `open5gs-gnb.yaml` must specify `ngapIp`**
- **UERANSIM `open5gs-gnb.yaml` must specify `gtpIp`**
- **UERANSIM `open5gs-gnb.yaml` must specify core AMF address**
- **Specifying `linkIp`, `ngapIp`, `gtpIp`, and core AMF address enables N2 connectivity**
- **Specifying `linkIp`, `ngapIp`, `gtpIp`, and core AMF address enables N3 connectivity**

Factoid: UE uses `gnbSearchList` in UERANSIM config to locate gNB IP for network attachment.

- **UE uses `gnbSearchList`**
- **`gnbSearchList` is defined in UERANSIM config**
- **`gnbSearchList` locates gNB IP**
- **gNB IP is used for network attachment**

Factoid: Upon PDU session setup, the UE forms a TUN interface (e.g., `uesimtun0`) to receive IP routes via the UPF.

- **UE forms TUN interface (e.g., `uesimtun0`)**
- **TUN interface (`uesimtun0`) receives IP routes**
- **IP routes are delivered via UPF**
- **PDU session setup triggers UE to form TUN interface**

Factoid: UE external Internet access can be validated using tools like curl or ping over the TUN interface.

- **UE external Internet access** *can be validated using curl*
- **UE external Internet access** *can be validated using ping*
- **curl** *operates over TUN interface*
- **ping** *operates over TUN interface*

Factoid: Test environments may employ TCP proxies on public IPs for multi-machine end-to-end network validation.

- **Test environments** *employ TCP proxies*
- **TCP proxies** *operate on public IPs*
- **TCP proxies on public IPs** *support multi-machine end-to-end network validation*

Factoid: Open5GS core services integrate with RAN simulators through manual IP alignment in both core and RAN YAML configurations.

- **Open5GS core services** *integrate with RAN simulators*
- **Integration** *occurs through manual IP alignment*
- **Manual IP alignment** *is done in core YAML configurations*
- **Manual IP alignment** *is done in RAN YAML configurations*

Factoid: Open5GS AMF by default binds NGAP/SCTP to loopback; must set `ngap.addr` in `amf.yaml` to LAN IP to enable RAN connectivity.

- **Open5GS AMF** *binds NGAP/SCTP to loopback by default*
- **Open5GS AMF** *must set ngap.addr in amf.yaml to LAN IP*
Setting ngap.addr in amf.yaml to LAN IP enables RAN connectivity

Factoid: N2 interface uses NGAP over SCTP to deliver UE NAS signaling and handover events from gNB to AMF.

- **N2 interface** *uses NGAP over SCTP*
- **N2 interface** *delivers UE NAS signaling from gNB to AMF*
- **N2 interface** *delivers handover events from gNB to AMF*

Factoid: For multi-host setups (Core and RAN on separate servers), explicit IP binding is required for N2 to function across hosts.

- **Multi-host setups** *involve Core on separate servers*
- **Multi-host setups** *involve RAN on separate servers*
- **Explicit IP binding** *is required for N2 to function across hosts*

Factoid: Running `sudo systemctl restart open5gs-amfd` applies the `ngap.addr` binding change to the AMF service.

- **sudo systemctl restart open5gs-amfd** *applies* **ngap.addr** binding change
- **ngap.addr** binding change *affects* **AMF** service

Factoid: UE RAN simulation with UERANSIM requires SCTP libraries (**libsctp-dev**) and CMake for building the gNB module.

- **UE RAN simulation with UERANSIM** *requires* **libsctp-dev**
- **UE RAN simulation with UERANSIM** *requires* **CMake**
- **libsctp-dev** *provides* **SCTP** libraries
- **CMake** *is used for* building the **gNB** module
- **gNB** module *is part of* **UERANSIM**

Factoid: Proper AMF binding allows remote UERANSIM-created gNB to successfully attach and exchange NGAP over SCTP.

- **Proper AMF binding** *allows* remote **UERANSIM-created gNB** *to successfully attach*
- **Remote UERANSIM-created gNB** *exchanges* **NGAP over SCTP**

Factoid: N2 SCTP handshake logs confirm successful gNB-AMF connectivity, validating multi-host network setup.

- **N2 SCTP handshake logs** *confirm* **successful gNB-AMF connectivity**
- **Successful gNB-AMF connectivity** *validates* multi-host network setup

Factoid: Open5GS metrics system was added to export performance counters and gauges for monitoring active PDP contexts and NF activity.

- **Open5GS metrics system** *exports* **performance counters**
- **Open5GS metrics system** *exports* **gauges**
- **Performance counters** *monitor* active **PDP** contexts
- **Gauges** *monitor* **NF** activity

Factoid: Metrics in Open5GS are implemented via **libprom** and **libpromhttp**, built on top of **libmicrohttpd**, enabling an embedded HTTP server.

- **Metrics in Open5GS** *are implemented via* **libprom**
- **Metrics in Open5GS** *are implemented via* **libpromhttp**
- **libprom** *is built on top of* **libmicrohttpd**
- **libpromhttp** *is built on top of* **libmicrohttpd**
- **Building on top of libmicrohttpd** *enables* embedded HTTP server

Factoid: The metrics subsystem uses a generic API in **lib/metrics/**, with a conditional Prometheus backend and a no-op fallback.

- **Metrics subsystem** *uses generic API in lib/metrics/*
Generic API in lib/metrics/ supports conditional Prometheus backend
- **Generic API in lib/metrics/ supports no-op fallback**

Factoid: Each NF (SMF, AMF, UPF, etc.) hosts its own `/metrics` endpoint, allowing individual scraping by Prometheus.

- **Each NF (SMF, AMF, UPF, etc.) hosts its own `/metrics` endpoint**
- **SMF hosts its own `/metrics` endpoint**
- **AMF hosts its own `/metrics` endpoint**
- **UPF hosts its own `/metrics` endpoint**
- **Hosting its own `/metrics` endpoint allows individual scraping by Prometheus**

Factoid: Metrics definitions include counters, gauges, and potentially histograms to capture NF performance and load.

- **Metrics definitions include counters**
- **Metrics definitions include gauges**
- **Metrics definitions include histograms**
- **Counters, gauges, and histograms capture NF performance and load**

Factoid: The HTTP metrics server is embedded in the NF process, so exposing metrics doesn't require external exporters.

- **HTTP metrics server is embedded in NF process**
- **Exposing metrics does not require external exporters**

Factoid: Prometheus scraping is enabled by building Open5GS with the Prometheus backend; otherwise, the metrics API is a stub.

- **Prometheus scraping is enabled by building Open5GS with the Prometheus backend**
- **Metrics API is a stub without the Prometheus backend**

Factoid: Example metrics include active session counts and internal NF resource telemetry.

- **Example metrics include active session counts**
- **Example metrics include internal NF resource telemetry**

Factoid: The conditional build approach allows operators to disable metrics support by omitting libprom-related dependencies.

- **Conditional build approach allows operators to disable metrics support**
- **Disabling metrics support occurs by omitting libprom-related dependencies**

Factoid: The architecture supports containerized deployments where each NF can be independently monitored.

- **The architecture *supports* containerized deployments**
- **Containerized deployments *allow* each NF to be independently monitored**

Factoid: CSPs transition from PNFs to CNFs using Kubernetes to gain vendor-agnostic, scalable infrastructure.

- **CSPs *transition from* PNFs to CNFs**
- **Transition to CNFs *uses* Kubernetes**
- **Transition using Kubernetes *enables* vendor-agnostic infrastructure**
- **Transition using Kubernetes *enables* scalable infrastructure**

Factoid: Disaggregated network functions (RU, DU, CU-UP, UPF) run at the edge under 5 ms RTT for optimal performance.

- **Disaggregated network functions (RU, DU, CU-UP, UPF) *run at* the edge**
- **Disaggregated network functions (RU, DU, CU-UP, UPF) *run under* 5 ms RTT**
- **Running under 5 ms RTT *provides* optimal performance**

Factoid: Control-plane NFs like AMF and SMF are deployed in centralized cloud regions, not latency-critical.

- **AMF *is deployed in* centralized cloud regions**
- **SMF *is deployed in* centralized cloud regions**
- **Control-plane NFs deployed in centralized cloud regions *are not* latency-critical**

Factoid: Anthos/GDC provides unified orchestration and policy across edge, private, and public 5G network deployments.

- **Anthos/GDC *provides* unified orchestration and policy**
- **Unified orchestration and policy *applies across* edge deployments**
- **Unified orchestration and policy *applies across* private deployments**
- **Unified orchestration and policy *applies across* public 5G network deployments**

Factoid: CNFs adhere to microservice design and CI/CD lifecycle models for fast feature deployment and upgrades.

- **CNFs *adhere to* microservice design**
- **CNFs *adhere to* CI/CD lifecycle models**
- **Microservice design and CI/CD lifecycle models *enable* fast feature deployment**
- **Microservice design and CI/CD lifecycle models *enable* upgrades**

Factoid: Edge deployment enables CSPs to host both 5G-CNFs and third-party edge applications (e.g., AR/VR) on shared infrastructure.

- **Edge deployment *enables* CSPs to host 5G-CNFs on shared infrastructure**
- **Edge deployment *enables* CSPs to host third-party edge applications on shared infrastructure**

- **Third-party edge applications** *include AR/VR*

Factoid: Use of hybrid clouds allows dynamic workload placement—edge for real-time tasks, cloud for batch or training.

- **Hybrid clouds** *allow dynamic workload placement*
- **Dynamic workload placement** *assigns edge to real-time tasks*
- **Dynamic workload placement** *assigns cloud to batch tasks*
- **Dynamic workload placement** *assigns cloud to training tasks*

Factoid: Telco workloads use infrastructure-as-code pipelines for security, scaling, and orchestration across distributed sites.

- **Telco workloads** *use Infrastructure-as-Code pipelines*
- **Infrastructure-as-Code pipelines** *enable security across distributed sites*
- **Infrastructure-as-Code pipelines** *enable scaling across distributed sites*
- **Infrastructure-as-Code pipelines** *enable orchestration across distributed sites*

Factoid: Latency-aware deployment ensures UPF/CU-UP services are co-located near user for ultra-low latency applications.

- **Latency-aware deployment** *ensures UPF/CU-UP services are co-located near users*
- **Co-location near users** *supports ultra-low latency applications*

Factoid: Container-based CNF approach enables feature-rich experiences while keeping costs and TCO under control.

- **Container-based CNF approach** *enables feature-rich experiences*
- **Container-based CNF approach** *keeps costs under control*
- **Container-based CNF approach** *keeps TCO under control*

Factoid: Telco workloads are migrating from VNFs in VMs to Kubernetes-managed CNFs deployed as Pods.

- **Telco workloads** *are migrating from VNFs in VMs*
- **Telco workloads** *are migrating to Kubernetes-managed CNFs deployed as Pods*

Factoid: Kubernetes Operators, via CRDs, support complex lifecycle tasks like upgrades and scaling for CNFs.

- **Kubernetes Operators** *use Custom Resource Definitions (CRDs)*
- **Kubernetes Operators** *support complex lifecycle tasks for CNFs*
- **Complex lifecycle tasks** *include upgrades*
- **Complex lifecycle tasks** *include scaling*
- **Custom Resource Definitions (CRDs)** *enable upgrades of CNFs*

- **Custom Resource Definitions (CRDs)** *enable* scaling of CNFs

Factoid: Microservice-based CNFs run each function in separate Pods, enabling independent scaling and failure isolation.

- **Microservice-based CNFs** *run* each function in separate Pods
- **Separate Pods** *enable* independent scaling
- **Separate Pods** *enable* failure isolation

Factoid: Monolithic CNFs simplify deployment but lack scalability granularity compared to microservice designs.

- **Monolithic CNFs** *simplify* deployment
- **Monolithic CNFs** *lack* scalability granularity
- **Scalability granularity** *is offered by* microservice designs

Factoid: KNI combines Kubernetes with DPDK, SR-IOV, and GPU networking to support telco-grade performance.

- **KNI** *combines* Kubernetes
- **KNI** *combines* DPDK
- **KNI** *combines* SR-IOV
- **KNI** *combines* GPU networking
- **Combining Kubernetes, DPDK, SR-IOV, and GPU networking** *supports* telco-grade performance

Factoid: Service mesh tools (Istio/Envoy) secure and manage APIs between control-plane CNFs in SBA.

- **Istio** *secures* APIs between control-plane CNFs in SBA
- **Envoy** *secures* APIs between control-plane CNFs in SBA
- **Service mesh tools (Istio/Envoy)** *manage* APIs between control-plane CNFs in SBA
- **APIs between control-plane CNFs in SBA** *operate under* Service-Based Architecture (SBA)

Factoid: Operators enable Site Reliability Engineering for CNFs by automating observability, self-healing, and dynamic scaling.

- **Operators** *enable* Site Reliability Engineering for CNFs
- **Site Reliability Engineering for CNFs** *involves* automating observability
- **Site Reliability Engineering for CNFs** *involves* self-healing
- **Site Reliability Engineering for CNFs** *involves* dynamic scaling

Factoid: KNI supports edge and core deployments using unified Kubernetes control plane for CNF orchestration.

- **KNI** *supports* edge deployments

- **KNI supports core deployments**
- **Edge and core deployments use unified Kubernetes control plane**
- **Unified Kubernetes control plane facilitates CNF orchestration**

Factoid: 5G core CNFs are deployed in multi-tier hub-and-spoke architecture with control-plane centralized and user-plane at edge.

- **5G core CNFs are deployed in multi-tier hub-and-spoke architecture**
- **Multi-tier hub-and-spoke architecture features centralized control-plane**
- **Multi-tier hub-and-spoke architecture features user-plane at edge**
- **Control-plane is centralized**
- **User-plane is deployed at edge**

Factoid: Edge CNFs (UPF/CU-UP) are co-located with RAN, while CNFs like AMF/SMF remain in central cloud.

- **UPF is co-located with RAN**
- **CU-UP is co-located with RAN**
- **AMF remains in central cloud**
- **SMF remains in central cloud**

Factoid: Istio-based service mesh provides mTLS, secure discovery, tracing and policy across distributed CNFs.

- **Istio-based service mesh provides mTLS**
- **Istio-based service mesh provides secure discovery**
- **Istio-based service mesh provides tracing**
- **Istio-based service mesh provides policy enforcement**
- **mTLS, secure discovery, tracing, and policy enforcement operate across distributed CNFs**

Factoid: Logs, metrics, and traces from edge sites are aggregated centrally (e.g., via Loki) for full-stack observability.

- **Logs are aggregated centrally from edge sites**
- **Metrics are aggregated centrally from edge sites**
- **Traces are aggregated centrally from edge sites**
- **Central aggregation uses Loki**
- **Central aggregation provides full-stack observability**

Factoid: Zero-touch provisioning and GitOps enable on-demand scaling and burst deployments of 5G CNF clusters.

- **Zero-touch provisioning enables on-demand scaling of 5G CNF clusters**
- **Zero-touch provisioning enables burst deployments of 5G CNF clusters**
- **GitOps enables on-demand scaling of 5G CNF clusters**
- **GitOps enables burst deployments of 5G CNF clusters**

Factoid: ACM placement rules ensure CNFs are deployed to appropriate clusters based on location and capacity policies.

- **ACM placement rules** *ensure* CNFs are deployed to appropriate clusters
- **CNFs deployed to appropriate clusters** *are based on* location policies
- **CNFs deployed to appropriate clusters** *are based on* capacity policies

Factoid: Network fabric layers include hub management, inter-cluster connectivity, and RAN access paths.

- **Network fabric layers** *include* hub management
- **Network fabric layers** *include* inter-cluster connectivity
- **Network fabric layers** *include* RAN access paths

Factoid: UPF selection is determined by SMF via NRF lookup using DNN, TAC, and cell_id metadata.

- **UPF selection** *is determined by* SMF
- **UPF selection** *is done via* NRF lookup
- **NRF lookup** *uses* DNN metadata
- **NRF lookup** *uses* TAC metadata
- **NRF lookup** *uses* cell_id metadata

Factoid: Partial edge deployment supports only UPF instances, while fully distributed bundles include SMF+UPF CNFs in remote sites.

- **Partial edge deployment** *supports only* UPF instances
- **Fully distributed bundles** *include* SMF CNFs in remote sites
- **Fully distributed bundles** *include* UPF CNFs in remote sites

Factoid: Service mesh federation allows secure cross-cluster communication and traffic splitting across hub and edge.

- **Service mesh federation** *allows* secure cross-cluster communication
- **Service mesh federation** *allows* traffic splitting across hub and edge

Factoid: Observability fabric leverages centralized policy enforcement across all distributed edge clusters.

- **Observability fabric** *leverages* centralized policy enforcement
- **Centralized policy enforcement** *applies across* distributed edge clusters

Factoid: Initial 5G deployments use NSA (RAN on 5G, control on 4G EPC); evolution to fully standalone (SA) core follows in later phases.

- **Initial 5G deployments** *use* NSA
- **NSA** *combines* RAN on 5G and control on 4G EPC
- **Evolution to fully standalone (SA) core** *follows in* later phases

BATCH 3 GEMINI

Factoid: AWS supports stateless microservices for 5G NFs via containers and managed databases like DynamoDB, Aurora, and ElastiCache.

- **AWS supports stateless microservices for 5G NFs**
- **Stateless microservices for 5G NFs are deployed via containers**
- **Stateless microservices for 5G NFs use managed databases like DynamoDB, Aurora, and ElastiCache**
- **Managed databases include DynamoDB**
- **Managed databases include Aurora**
- **Managed databases include ElastiCache**

Factoid: CUPS is implemented using AWS Outposts and Local Zones to deploy UPF at the edge, while control-plane NFs remain in central regions.

- **CUPS is implemented using AWS Outposts**
CUPS is implemented using Local Zones
- **AWS Outposts deploy UPF at the edge**
- **Local Zones deploy UPF at the edge**
- **Control-plane NFs remain in central regions**

Factoid: Kubernetes (EKS) orchestrates CNFs on AWS, with CI/CD pipelines enabling rapid lifecycle management and blue/green deployments.

- **Kubernetes (EKS) orchestrates CNFs on AWS**
- **CI/CD pipelines enable rapid lifecycle management**
- **CI/CD pipelines enable blue/green deployments**

Factoid: Stateless design externalizes NF state to UDSF-like stores, enabling resilience and container-based scalability.

- **Stateless design externalizes NF state to UDSF-like stores**
- **Externalizing NF state to UDSF-like stores enables resilience**
- **Externalizing NF state to UDSF-like stores enables container-based scalability**

Factoid: Network slicing is orchestrated through AWS service mesh (App Mesh) and programmable infrastructure via CDK and Step Functions.

- **Network slicing is orchestrated through AWS service mesh (App Mesh)**
- **Network slicing is orchestrated through programmable infrastructure**
- **Programmable infrastructure uses CDK**
- **Programmable infrastructure uses Step Functions**

Factoid: AWS offers enhanced instances with SR-IOV, DPDK, huge pages, and bare-metal to meet packet-processing demands of CNFs.

- **AWS offers enhanced instances**

- **Enhanced instances** *support SR-IOV*
- **Enhanced instances** *support DPDK*
Enhanced instances *support huge pages*
- **Enhanced instances** *support bare-metal*
- **Enhanced instances** *meet packet-processing demands of CNFs*

Factoid: Container-based NF design on AWS aligns with the ‘12-factor app’ model, ensuring process isolation, observability, and lifecycle automation.

- **Container-based NF design on AWS** *aligns with ‘12-factor app’ model*
- **‘12-factor app’ model** *ensures process isolation*
- **‘12-factor app’ model** *ensures observability*
- **‘12-factor app’ model** *ensures lifecycle automation*

Factoid: AWS Direct Connect and Global Accelerator deliver dedicated, high-performance network paths for 5G workload placement.

- **AWS Direct Connect** *delivers dedicated, high-performance network paths for 5G workload placement*
- **Global Accelerator** *delivers dedicated, high-performance network paths for 5G workload placement*
- **Dedicated, high-performance network paths** *support 5G workload placement*

Factoid: DevOps automation (CodePipeline, CloudFormation) allows NF operators to treat infrastructure as code, enabling reproducible, version-controlled deployments at scale.

- **DevOps automation (CodePipeline, CloudFormation)** *allows NF operators to treat infrastructure as code*
- **Treating infrastructure as code** *enables reproducible deployments at scale*
- **Treating infrastructure as code** *enables version-controlled deployments*

Factoid: UPF deployment may be hardware-based or virtualized; choice depends on SLA, performance needs, and vendor maturity.

- **UPF deployment** *may be hardware-based*
- **UPF deployment** *may be virtualized*
- **Choice of deployment** *depends on SLA*
- **Choice of deployment** *depends on performance needs*
- **Choice of deployment** *depends on vendor maturity*

Factoid: Hardware UPF suits high-performance needs; virtualized UPF aids flexibility and geo-distribution.

- **UPF deployment** *may be hardware-based*
- **UPF deployment** *may be virtualized*
- **Choice of deployment** *depends on SLA*
- **Choice of deployment** *depends on performance needs*

- **Choice of deployment** *depends on vendor maturity*
- **Hardware UPF** *suits high-performance needs*
- **Virtualized UPF** *aids flexibility*
- **Virtualized UPF** *aids geo-distribution*

Factoid: UPF steers traffic at edge using standardized APIs to route flows toward local or cloud-hosted applications.

- **UPF** *steers traffic at edge*
- **UPF** *uses standardized APIs*
- **Standardized APIs** *route flows toward local applications*
- **Standardized APIs** *route flows toward cloud-hosted applications*

Factoid: Edge IaaS platforms host VMs and containers managed by OpenStack and other platforms.

- **Edge IaaS platforms** *host VMs*
- **Edge IaaS platforms** *host containers*
- **VMs** *are managed by OpenStack*
- **Containers** *are managed by OpenStack*
- **VMs** *are managed by other platforms*
- **Containers** *are managed by other platforms*

Factoid: Edge PaaS layers provide networking, RNIS, location, user identity, firewall, DNS, and load balancing services.

- **Edge PaaS layers** *provide networking services*
- **Edge PaaS layers** *provide RNIS services*
- **Edge PaaS layers** *provide location services*
- **Edge PaaS layers** *provide user identity services*
- **Edge PaaS layers** *provide firewall services*
- **Edge PaaS layers** *provide DNS services*
- **Edge PaaS layers** *provide load balancing services*

Factoid: An orchestrator deploys NFs and applications based on service requirements, policies, and resource templates mapped to location.

- **Orchestrator** *deploys NFs and applications*
- **Deployment** *is based on service requirements*
- **Deployment** *is based on policies*
- **Deployment** *is based on resource templates mapped to location*

Factoid: Traffic steering APIs allow flow decisions to be made dynamically per application context (e.g., video, IoT).

- **Traffic steering APIs** *allow dynamic flow decisions per application context*
- **Dynamic flow decisions** *are made per application context*
- **Application contexts** *include video*

- **Application contexts** *include* IoT

Factoid: Cisco UPF implements 1:1 Active/Standby redundancy via SRP with ICSR-based state sync.

- **Cisco UPF** *implements* **1:1 Active/Standby redundancy**
- **1:1 Active/Standby redundancy** *uses* **SRP**
- **SRP** *performs* **ICSR-based state synchronization**

Factoid: Standby UPF assumes the same Sx/N4 address during switchover, making the transition transparent to the SMF.

- **Standby UPF** *assumes* **the same Sx/N4 address** during **switchover**
- **Assuming the same Sx/N4 address during switchover** *makes* **the transition transparent to the SMF**

Factoid: Sx/N4 control-plane heartbeat monitoring with BFD triggers fast UPF switchover in failure events.

- **Sx/N4 control-plane heartbeat monitoring with BFD** *triggers* **fast UPF switchover**
- **Fast UPF switchover** *occurs in* **failure events**

Factoid: Active UPF replicates IP-pool and session context to Standby during Sx/N4 association and checkpoint cycles.

- **Active UPF** *replicates* **IP-pool** to **Standby UPF**
- **Active UPF** *replicates* **session context** to **Standby UPF**
- **Replication** *occurs during* **Sx/N4 association**
- **Replication** *occurs during* **checkpoint cycles**

Factoid: Standby UPF starts in 'Pending-Active' until SRP elections and manual timeout configurations finalize switchover.

Factoid: VPP health and BGP monitoring are integrated into SRP for multi-layered UPF redundancy triggering.

- **Standby UPF** *starts in* **'Pending-Active'**
- **Standby UPF** *transitions out of* **'Pending-Active'** *after* **SRP elections**
- **Standby UPF** *transitions out of* **'Pending-Active'** *after* **manual timeout configurations**
- **SRP elections** and **manual timeout configurations** *finalize* **switchover**

Factoid: Manual CLI controls (e.g., `srp reset-sx-fail`) allow operators to override automatic failover conditions.

- **VPP health** *is integrated into SRP*
- **BGP monitoring** *is integrated into SRP*
- **Integration into SRP** *enables multi-layered UPF redundancy triggering*
- **Multi-layered UPF redundancy triggering** *relies on SRP*

Factoid: SMF is unaware of standby UPF and always interacts with the active endpoint via stable Sx/N4 address.

- **SMF** *is unaware of standby UPF*
- **SMF** *interacts with active UPF endpoint*
- **Interaction** *occurs via stable Sx/N4 address*

Factoid: SRP Active/Standby redundancy is supported without dual-active scenarios due to address takeover and heartbeat control.

- **SRP Active/Standby redundancy** *is supported without dual-active scenarios*
- **SRP Active/Standby redundancy** *relies on address takeover*
- **SRP Active/Standby redundancy** *relies on heartbeat control*

Factoid: Proper timing between SMF heartbeat and UPF SRP timeout is crucial to avoid false session drop detection.

- **SMF heartbeat** *must be timed appropriately relative to UPF SRP timeout*
- **Proper timing between SMF heartbeat and UPF SRP timeout** *is crucial to avoid false session drop detection*

Factoid: An NF Set groups interchangeable NF instances of the same type (e.g., AMF, SMF), enabling shared context and failover capability.

- **NF Set** *groups interchangeable NF instances of the same type*
- **Interchangeable NF instances of the same type** *include AMF*
- **Interchangeable NF instances of the same type** *include SMF*
- **Grouping by NF Set** *enables shared context*
- **Grouping by NF Set** *enables failover capability*

Factoid: NF Sets support geo-redundancy via distributed instances and externalized session state.

- **NF Sets** *support geo-redundancy*
- **Geo-redundancy** *is achieved via distributed instances*
- **Geo-redundancy** *is achieved via externalized session state*

Factoid: Stateless NFs in an NF Set externalize their state to UDSF, enabling context retrieval by any instance.

- **Stateless NFs in an NF Set** *externalize state to UDSF*
- **Externalizing state to UDSF** *enables context retrieval by any instance*

Factoid: NF instances in a set can have different software versions, enabling seamless rolling upgrades.

- **NF instances in an NF Set** *can have different software versions*
- **Different software versions** *enable seamless rolling upgrades*

Factoid: NF Set uses N+M redundancy, reducing overprovisioning while enabling fast failover.

- **NF Set** *uses N+M redundancy*
- **N+M redundancy** *reduces overprovisioning*
- **N+M redundancy** *enables fast failover*

Factoid: Context-sharing in NF Sets prevents signaling storms by avoiding re-registration during failover.

- **Context-sharing in NF Sets** *prevents signaling storms*
- **Context-sharing in NF Sets** *avoids re-registration during failover*
- **Avoiding re-registration during failover** *prevents signaling storms*

Factoid: An NF Set presents a single virtual interface (e.g., anycast IP) to clients despite multiple active instances.

- **NF Set** *presents single virtual interface (e.g., anycast IP)*
- **Single virtual interface (e.g., anycast IP)** *is presented to clients*

Factoid: Geographically distributed NF Set instances enable localized scaling and resilience per region.

- **Geographically distributed NF Set instances** *enable localized scaling*
- **Geographically distributed NF Set instances** *enable resilience per region*

Factoid: Disjoint UPF deployment with dual PDU sessions increases end-to-end availability by enabling parallel user-plane paths.

- **Disjoint UPF deployment with dual PDU sessions** *increases end-to-end availability*
- **End-to-end availability** *is increased by enabling parallel user-plane paths*

Factoid: Dual-connectivity through two separate gNBs allows simultaneous data flows to distinct UPFs, forming a redundant user-plane chain.

- **Dual-connectivity** *occurs through two separate gNBs*
- **Dual-connectivity** *allows simultaneous data flows*
- **Simultaneous data flows** *go to distinct UPFs*
- **Simultaneous data flows** *form a redundant user-plane chain*

Factoid: User-plane redundancy is achieved via dual UPFs; control-plane remains single-path to maintain management simplicity.

- **User-plane redundancy** *is achieved via dual UPFs*
- **Control-plane** *remains single-path*
- **Single-path control-plane** *maintains management simplicity*

Factoid: Industrial/private 5G deployments can reach availability levels above 99.9999% using dual-path UPF redundancy.

- **Industrial/private 5G deployments** *can reach availability levels above 99.9999%*
- **Availability levels above 99.9999%** *are achieved using dual-path UPF redundancy*
- **Dual-path UPF redundancy** *provides user-plane redundancy*

Factoid: E2E availability gains are calculated using parallel reliability models like Reliability Block Diagrams.

- **E2E availability gains** *are calculated using parallel reliability models*
- **Parallel reliability models** *include Reliability Block Diagrams*

Factoid: Mean Time To Repair (MTTR) and path redundancy are the two primary enablers of achieving telecom-grade availability targets.

- **Mean Time To Repair (MTTR)** *is a primary enabler of achieving telecom-grade availability targets*
- **Path redundancy** *is a primary enabler of achieving telecom-grade availability targets*

Factoid: Redundancy strategy leverages dual PDU sessions mapped across disjoint UPFs and RAN connections for ultra-reliable industrial use cases.

- **Redundancy strategy** *leverages dual PDU sessions*
- **Dual PDU sessions** *are mapped across disjoint UPFs and RAN connections*
- **Mapping across disjoint UPFs and RAN connections** *supports ultra-reliable industrial use cases*

Factoid: URLLC applications demand sub-millisecond latency and negligible packet loss, requiring stronger resilience than standard Fast ReRoute mechanisms.

- **URLLC applications** *demand sub-millisecond latency*
- **URLLC applications** *demand negligible packet loss*
- **URLLC applications** *require stronger resilience*
- **Stronger resilience** *exceeds standard Fast ReRoute mechanisms*

Factoid: PREOF enables 1+1 path protection by duplicating packets over two disjoint UPF paths and eliminating redundant packets at the receiver.

- **PREOF** *enables 1+1 path protection*
- **PREOF** *duplicates packets over two disjoint UPF paths*

- **Redundant packets are eliminated at the receiver**

Factoid: The PREOF mechanism can be deployed at the gNB or UE as a Protection Tunnel Ingress (PTI), replicating packets for redundancy.

- **PREOF mechanism can be deployed at gNB as Protection Tunnel Ingress (PTI)**
- **PREOF mechanism can be deployed at UE as Protection Tunnel Ingress (PTI)**
- **Protection Tunnel Ingress (PTI) replicates packets for redundancy**

Factoid: Replicated packets traverse two synchronized UPFs before reaching a Protection Tunnel Egress (PTE) node for elimination.

- **Replicated packets traverse two synchronized UPFs**
- **Replicated packets reach Protection Tunnel Egress (PTE) node**
- **Protection Tunnel Egress (PTE) node eliminates replicated packet**

Factoid: Ordering of packets is offloaded to an external PTE-O server or implemented via eBPF/DPDK, reducing complexity in programmable hardware switches.

- **Ordering of packets is offloaded to external PTE-O server**
- **Ordering of packets is implemented via eBPF/DPDK**
- **Offloading ordering of packets reduces complexity in programmable hardware switches**

Factoid: PREOF is fully compatible with existing GTP-U tunnels, encapsulating replicated packets transparently to UPFs.

- **PREOF is fully compatible with existing GTP-U tunnels**
- **PREOF encapsulates replicated packets**
- **Encapsulation of replicated packets is transparent to UPFs**

Factoid: 1+1 path protection increases system resilience at the cost of slightly higher latency and resource usage.

- **1+1 path protection increases system resilience**
- **1+1 path protection incurs slightly higher latency**
- **1+1 path protection incurs higher resource usage**

Factoid: PREOF aligns with 3GPP Release 18's support for packet duplication mechanisms to enhance URLLC resilience.

- **PREOF aligns with 3GPP Release 18's support for packet duplication mechanisms**
- **3GPP Release 18's support for packet duplication mechanisms enhances URLLC resilience**

Factoid: Disjoint UP paths with replicated UPFs ensure session survival even if one UPF or path experiences a failure.

- **Disjoint UP paths** *use replicated UPFs*
- **Disjoint UP paths with replicated UPFs** *ensure session survival*
- **Session survival** *occurs even if one UPF or path experiences a failure*

Factoid: Choosing between PTI placement at gNB versus UE reflects trade-offs in performance, complexity, and protection scope.

- **Choosing PTI placement at gNB** *reflects trade-offs in performance*
- **Choosing PTI placement at gNB** *reflects trade-offs in complexity*
- **Choosing PTI placement at gNB** *reflects trade-offs in protection scope*
- **Choosing PTI placement at UE** *reflects trade-offs in performance*
- **Choosing PTI placement at UE** *reflects trade-offs in complexity*
- **Choosing PTI placement at UE** *reflects trade-offs in protection scope*

Factoid: Static resource allocation fails in 5G; dynamic models like MCM are needed for real-time adaptability.

- **Static resource allocation** *fails in 5G*
- **Dynamic models** *include MCM*
- **Dynamic models** *are needed for real-time adaptability*

Factoid: MCM integrates bandwidth allocation, traffic prioritization, encryption, and network slicing to enforce QoS.

- **MCM** *integrates bandwidth allocation*
- **MCM** *integrates traffic prioritization*
- **MCM** *integrates encryption*
- **MCM** *integrates network slicing*
- **Integration of bandwidth allocation, traffic prioritization, encryption, and network slicing** *enforces QoS*

Factoid: Target QoS metrics include sub-ms latency and bounded jitter for URLLC, plus hundreds of Mbps throughput for eMBB.

- **Target QoS metrics** *include sub-ms latency for URLLC*
- **Target QoS metrics** *include bounded jitter for URLLC*
- **Target QoS metrics** *include hundreds of Mbps throughput for eMBB*
- **Sub-ms latency** *is for URLLC*
- **Bounded jitter** *is for URLLC*
- **Hundreds of Mbps throughput** *is for eMBB*

Factoid: ML-based traffic prediction is used in MCM to dynamically reallocate resources in anticipation of demand peaks.

- **ML-based traffic prediction** *is used in MCM*
- **ML-based traffic prediction** *dynamically reallocates resources*

- **Dynamically reallocating resources** *anticipates demand peaks*

Factoid: MDP/DRL-based RAN schedulers dynamically balance URLLC and eMBB performance on slot and mini-slot timescales.

- **MDP/DRL-based RAN schedulers** *dynamically balance URLLC performance*
MDP/DRL-based RAN schedulers *dynamically balance eMBB performance*
Dynamic balancing of URLLC and eMBB performance *occurs on slot timescales*
- **Dynamic balancing of URLLC and eMBB performance** *occurs on mini-slot timescales*

Factoid: EDQAS/LDI schedulers at MAC layer schedule resource blocks to minimize uRLLC latency and eMBB rate loss.

- **EDQAS/LDI schedulers** *schedule resource blocks*
- **EDQAS/LDI schedulers** *minimize uRLLC latency*
- **EDQAS/LDI schedulers** *minimize eMBB rate loss*
- **Scheduling resource blocks** *occurs at MAC layer*

Factoid: Network slices partition resources with per-slice SLA-driven enforcement across RAN and core domains.

- **Network slices** *partition resources*
- **Network slices** *provide per-slice SLA-driven enforcement*
- **Per-slice SLA-driven enforcement** *applies across RAN domains*
- **Per-slice SLA-driven enforcement** *applies across core domains*

Factoid: Encryption overhead is modeled in resource allocation decisions to meet QoS while ensuring security.

- **Encryption overhead** *is modeled in resource allocation decisions*
- **Resource allocation decisions** *aim to meet QoS*
- **Resource allocation decisions** *ensure security*

Factoid: QoS parameters considered include latency, jitter, packet loss, throughput, spectral efficiency, and energy efficiency.

- **QoS parameters** *include latency*
- **QoS parameters** *include jitter*
- **QoS parameters** *include packet loss*
- **QoS parameters** *include throughput*
- **QoS parameters** *include spectral efficiency*
- **QoS parameters** *include energy efficiency*

_____ BATCH 4 GEMINI _____

Factoid: Increasing URLLC performance often involves puncturing eMBB traffic in real-time, trading off throughput.

- **URLLC performance improvements** *involve puncturing eMBB traffic in real-time*
- **Puncturing eMBB traffic in real-time** *trades off throughput*

Factoid: QoS-aware slicing leverages vertical slice models to allocate bandwidth and priority between service types.

QoS-aware slicing *leverages vertical slice models*

Vertical slice models *allocate bandwidth between service types*

Vertical slice models *allocate priority between service types*

Factoid: Dynamic resource allocation models consider fairness, availability, and service resilience during congestion.

- **Dynamic resource allocation models** *consider fairness*
- **Dynamic resource allocation models** *consider availability*
- **Dynamic resource allocation models** *consider service resilience*
- **Consideration of fairness, availability, and service resilience** *occurs during congestion*

Factoid: Resource allocation frameworks now combine RAN-level scheduling with core-level queue management for SLA delivery.

- **Resource allocation frameworks** *combine RAN-level scheduling and core-level queue management*
- **RAN-level scheduling and core-level queue management** *enable SLA delivery*

Factoid: URLLC requires ≤ 1 ms one-way latency and $\geq 99.999\%$ reliability, posing stringent QoS targets.

- **URLLC** *requires ≤ 1 ms one-way latency*
- **URLLC** *requires $\geq 99.999\%$ reliability*
- **Stringent QoS targets** *are posed by URLLC*

Factoid: An inline DPDK-accelerated Snort IPS can meet median latency but shows tail latency spikes above 1 ms—problematic for URLLC.

- **Inline DPDK-accelerated Snort IPS** *can meet median latency*
- **Inline DPDK-accelerated Snort IPS** *shows tail latency spikes above 1 ms*
- **Tail latency spikes above 1 ms** *are problematic for URLLC*

Factoid: Worst-case latency with Snort IPS in a VM hit up to 2.5 ms at the 99.999th percentile even with no rule matching.

- **Snort IPS in a VM** *hit 2.5 ms worst-case latency*
- **2.5 ms worst-case latency** *occurred at 99.999th percentile*

- **2.5 ms worst-case latency at the 99.999th percentile occurred even with no rule matching**

Factoid: Packet processing stacks using Linux + DPDK + Snort reveal virtualization unpredictability due to interrupts and CPU scaling.

- **Packet processing stacks use Linux**
- **Packet processing stacks use DPDK**
- **Packet processing stacks use Snort**
- **Virtualization unpredictability is revealed due to interrupts**
- **Virtualization unpredictability is revealed due to CPU scaling**

Factoid: Mitigation for virtualization-induced latency includes dedicated cores, CPU isolation, interrupt pinning, and run-to-completion processing.

- **Packet processing stacks use Linux**
- **Packet processing stacks use DPDK**
- **Packet processing stacks use Snort**
- **Virtualization unpredictability is revealed due to interrupts**
- **Virtualization unpredictability is revealed due to CPU scaling**

Factoid: A predictive model estimates maximum sustainable IPS load under URLLC, enabling SLA-driven capacity planning.

- **Predictive model estimates maximum sustainable IPS load under URLLC**
- **Estimating maximum sustainable IPS load under URLLC enables SLA-driven capacity planning**

Factoid: Middleware security functions (like IPS) must be carefully optimized to avoid violating URLLC tail latency requirements.

- **Middleware security functions (like IPS) must be carefully optimized**
- **Careful optimization of middleware security functions (like IPS) avoids violating URLLC tail latency requirements**

Factoid: Hardware-assisted timestamping is essential for precise tail-latency measurement when evaluating IPS impact on URLLC.

- **Hardware-assisted timestamping is essential for precise tail-latency measurement**
- **Precise tail-latency measurement enables evaluating IPS impact on URLLC**

Factoid: Even with minimal packet inspection, packet processing latency remains unpredictable due to virtual environment effects.

- **Packet processing latency remains unpredictable due to virtual environment effects**

- **Minimal packet inspection** *does not prevent unpredictable packet processing latency*

Factoid: To support URLLC with inline security, packet processing pipelines must be architected for low jitter and deterministic behavior.

- **Packet processing pipelines** *must be architected for low jitter*
- **Packet processing pipelines** *must be architected for deterministic behavior*
- **Architecting packet processing pipelines for low jitter and deterministic behavior** *supports URLLC with inline security*

Factoid: 5G URLLC targets ≤ 1 ms user-plane latency—a $\sim 4\times$ reduction compared to LTE.

- **5G URLLC targets** *mandate ≤ 1 ms user-plane latency*
- **≤ 1 ms user-plane latency** *represents $\approx 4\times$ reduction compared to LTE*

Factoid: URLLC demands ultra-reliability ($\geq 99.999\%$), with block error rates down to 10^{-9} .

- **URLLC** *demands ultra-reliability ($\geq 99.999\%$)*
- **URLLC** *demands block error rates down to 10^{-9}*
- **Ultra-reliability ($\geq 99.999\%$)** *corresponds to block error rates down to 10^{-9}*

Factoid: QoS Flow Identifier (QFI) allows differentiated packet handling in the user plane.

- **QoS Flow Identifier (QFI)** *allows differentiated packet handling in the user plane*

Factoid: URLLC flows use resource reservation with preemption and mini-slots to meet sub-ms deadlines.

- **URLLC flows** *use resource reservation*
- **URLLC flows** *use preemption*
- **URLLC flows** *use mini-slots*
- **Use of resource reservation with preemption and mini-slots** *enables meeting sub-ms deadlines*

Factoid: Edge computing eliminates ~ 100 ms transport delay, enabling end-to-end latency ≤ 1 ms.

- **Edge computing** *eliminates ~ 100 ms transport delay*
- **Eliminating ~ 100 ms transport delay** *enables end-to-end latency ≤ 1 ms*

Factoid: Flexible TTIs and robust coding techniques are used to reduce transmission latency and BER.

- **Flexible TTIs** *are used to reduce transmission latency*
- **Robust coding techniques** *are used to reduce transmission latency*
- **Flexible TTIs** *are used to reduce BER*

- **Robust coding techniques** *are used to reduce BER*

Factoid: Multiple UPFs with local breakout support disjoint user-plane paths for low-latency and resilience.

- **Multiple UPFs with local breakout** *support disjoint user-plane paths*
- **Disjoint user-plane paths** *provide low-latency*
- **Disjoint user-plane paths** *provide resilience*

Factoid: URLLC scheduling uses prioritized access and TTI scaling to minimize delay and jitter.

- **URLLC scheduling** *uses prioritized access*
- **URLLC scheduling** *uses TTI scaling*
- **Prioritized access** *minimizes delay and jitter*
- **TTI scaling** *minimizes delay and jitter*

Factoid: Reliable URLLC transmission combines mini-slots, redundancy, and HARQ optimization.

- **Reliable URLLC transmission** *combines mini-slots*
- **Reliable URLLC transmission** *combines redundancy*
- **Reliable URLLC transmission** *combines HARQ optimization*

Factoid: Edge-deployed UPFs and MEC collaborate to enforce QoS and satisfy URLLC requirements.

- **Edge-deployed UPFs** *collaborate with MEC*
- **Collaboration between Edge-deployed UPFs and MEC** *enforces QoS*
- **Collaboration between Edge-deployed UPFs and MEC** *satisfies URLLC requirements*

Factoid: URLLC flows demand ≤ 1 ms latency, low jitter, and ultra-high reliability; eMBB flows prioritize high throughput with moderate latency; mMTC supports up to 1 M devices/km² with low rate requirements.

- **URLLC flows** *demand ≤ 1 ms latency*
- **URLLC flows** *demand low jitter*
- **URLLC flows** *demand ultra-high reliability*
- **eMBB flows** *prioritize high throughput*
- **eMBB flows** *prioritize moderate latency*
- **mMTC flows** *support up to 1 M devices/km²*
- **mMTC flows** *support low rate requirements*

Factoid: QoS metrics include latency, jitter, packet loss, throughput, reliability, fairness, and energy efficiency across 5G slices.

- **QoS metrics** *include latency*

- **QoS metrics** *include jitter*
- **QoS metrics** *include packet loss*
- **QoS metrics** *include throughput*
- **QoS metrics** *include reliability*
- **QoS metrics** *include fairness*
- **QoS metrics** *include energy efficiency*
- **QoS metrics** *apply across 5G slices*

Factoid: Slice-specific SLAs are enforced through dynamic bandwidth allocation and priority queueing per network slice.

- **Slice-specific SLAs** *are enforced through dynamic bandwidth allocation*
- **Slice-specific SLAs** *are enforced through priority queueing per network slice*

Factoid: Traffic shaping uses DiffServ: ingress classification, metering, and marking enable per-hop QoS enforcement.

- **Traffic shaping** *uses DiffServ*
- **Ingress classification** *enables per-hop QoS enforcement*
- **Metering** *enables per-hop QoS enforcement*
- **Marking** *enables per-hop QoS enforcement*

Factoid: URLLC scheduling uses preemption/puncturing or mini-slots at RAN to meet sub-ms latency requirements.

- **URLLC scheduling** *uses preemption/puncturing*
- **URLLC scheduling** *uses mini-slots at RAN*
- **Preemption/puncturing and mini-slots at RAN** *meet sub-ms latency requirements*

Factoid: eMBB scheduling applies proportional fairness or weighted round-robin to balance throughput and fairness.

Factoid: mMTC devices use dynamic channel access protocols to handle high device density and sporadic traffic.

- **eMBB scheduling** *applies proportional fairness*
- **eMBB scheduling** *applies weighted round-robin*
- **Proportional fairness** *balances throughput and fairness*
- **Weighted round-robin** *balances throughput and fairness*

Factoid: Transport-priority queuing ensures UL/DL URLLC packets are served ahead of eMBB and mMTC traffic.

- **Transport-priority queuing** *ensures UL/DL URLLC packets are served ahead of eMBB traffic*

- **Transport-priority queuing** *ensures UL/DL URLLC packets are served ahead of mMTC traffic*

Factoid: Hard slicing uses dedicated resources per slice; soft slicing shares resources with prioritization—supporting mixed traffic isolation.

- **Hard slicing** *uses dedicated resources per slice*
- **Soft slicing** *shares resources with prioritization*
- **Hard slicing and soft slicing** *support mixed traffic isolation*

Factoid: Reliability and latency goals for URLLC are modeled via resource block reservation and robust coding at physical and transport layers.

- **Reliability and latency goals for URLLC** *are modeled via resource block reservation*
Reliability and latency goals for URLLC *are modeled via robust coding*
- **Resource block reservation** *occurs at physical layer*
- **Resource block reservation** *occurs at transport layer*
- **Robust coding** *occurs at physical layer*
- **Robust coding** *occurs at transport layer*

Factoid: QoS enforcement spans RAN, transport, and core layers—ensuring consistent SLA adherence.

- **QoS enforcement** *spans RAN*
- **QoS enforcement** *spans transport*
- **QoS enforcement** *spans core layers*
- **Spanning RAN, transport, and core layers** *ensures consistent SLA adherence*

Factoid: Scheduler models use optimization and ML techniques to dynamically allocate resources based on real-time QoS demands.

- **Scheduler models** *use optimization techniques*
- **Scheduler models** *use ML techniques*
- **Scheduler models** *dynamically allocate resources*
- **Dynamic allocation of resources** *is based on real-time QoS demands*

Factoid: Each QoS Flow in 5G is tagged with a unique 4-bit QFI within its PDU session.

- **QoS Flow** *is tagged with unique 4-bit QFI*
- **Tagging with unique 4-bit QFI** *occurs within PDU session*

Factoid: QoS Flows are classified as GBR or Non-GBR, with specific QoS profiles defined.

- **QoS Flows** *are classified as GBR*
- **QoS Flows** *are classified as Non-GBR*
- **Specific QoS profiles** *are defined for GBR QoS Flows*

- **Specific QoS profiles** *are defined for Non-GBR QoS Flows*

Factoid: UPF applies Packet Detection Rules (PDRs) and UE applies QoS rules to map packets to QoS flows at NAS layer.

- **UPF** *applies Packet Detection Rules (PDRs)*
- **UE** *applies QoS rules*
- **Applying PDRs and QoS rules** *maps packets to QoS flows at the NAS layer*

Factoid: Packet filters use IP addresses, ports, protocols, and Ethernet tags to classify packets.

- **Packet filters** *use IP addresses*
- **Packet filters** *use ports*
- **Packet filters** *use protocols*
- **Packet filters** *use Ethernet tags*
- **Using IP addresses, ports, protocols, and Ethernet tags** *enables packet classification*

Factoid: SDAP sublayer maps QoS flows with given QFI to DRBs, supported by RRC-configured mapping rules.

- **Kubernetes-based control plane orchestration** *benefits from stateless NF design*
- **Stateless NF design** *enables auto-replacement of crashed pods*
- **Auto-replacement of crashed pods** *occurs without state loss*

Factoid: Multiple QoS flows can share a DRB if their service requirements align.

- **Multiple QoS flows** *can share a DRB*
- **Sharing a DRB** *requires aligned service requirements*

Factoid: Reflective QoS enables UE to derive uplink QoS-to-DRB mapping from downlink rules, using RQI/RDI flags.

- **Reflective QoS** *enables UE to derive uplink QoS-to-DRB mapping from downlink rules*
- **Deriving uplink QoS-to-DRB mapping from downlink rules** *uses RQI/RDI flags*

Factoid: N3 GTP-U headers carry QFI for QoS identification between UPF and gNB.

- **N3 GTP-U headers** *carry QFI*
- **QFI** *is used for QoS identification*
- **QoS identification** *occurs between UPF and gNB*

Factoid: SMF configures QoS flows by distributing QoS rules and PDRs to UE (N1), gNB (N2), and UPF (N4).

- **SMF** *configures QoS flows*
- **SMF** *distributes QoS rules to UE via N1*

- **SMF distributes PDRs to gNB via N2**
- **SMF distributes PDRs to UPF via N4**

Factoid: QoS control spans NAS classification, SDAP mapping, and GTP-U marking to enforce consistent end-to-end service quality.

- **QoS control spans NAS classification**
- **QoS control spans SDAP mapping**
- **QoS control spans GTP-U marking**
- **NAS classification enforces consistent end-to-end service quality**
- **SDAP mapping enforces consistent end-to-end service quality**
- **GTP-U marking enforces consistent end-to-end service quality**

Factoid: 5G mandates support for 5G-AKA and EAP-AKA' for unified, access-agnostic authentication.

- **5G mandates support for 5G-AKA**
- **5G mandates support for EAP-AKA'**
- **5G-AKA provides unified, access-agnostic authentication**
- **EAP-AKA' provides unified, access-agnostic authentication**

Factoid: EAP-TLS can be used as a secondary authentication method in private or enterprise deployments.

- **EAP-TLS can be used as secondary authentication method**
- **Secondary authentication method is used in private deployments**
- **Secondary authentication method is used in enterprise deployments**

Factoid: AUSF anchors the security key KSEAF during initial authentication and supplies authentication vectors.

- **AUSF anchors security key KSEAF**
- **Anchoring security key KSEAF occurs during initial authentication**
- **AUSF supplies authentication vectors**

Factoid: UDM/ARPF stores subscriber credentials and supports authentication via both AKA and EAP frameworks.

- **UDM/ARPF stores subscriber credentials**
- **UDM/ARPF supports authentication via AKA framework**
- **UDM/ARPF supports authentication via EAP framework**

Factoid: All SBA APIs must use mutual TLS with client and server certificates, often coupled with OAuth authorization.

- **All SBA APIs must use mutual TLS**
- **Mutual TLS requires client certificates**
- **Mutual TLS requires server certificates**

- **All SBA APIs are often coupled with OAuth authorization**

Factoid: IPsec tunnels (NWu/NWt) protect NAS signaling over non-3GPP access, with optional NULL encryption for trusted networks.

- **IPsec tunnels (NWu/NWt) protect NAS signaling**
- **IPsec tunnels (NWu/NWt) operate over non-3GPP access**
- **Optional NULL encryption is used for trusted networks**

Factoid: SEPP secures inter-PLMN signaling on N32 via TLS or PRINS, ensuring integrity and confidentiality.

- **SEPP secures inter-PLMN signaling on N32 via TLS**
- **SEPP secures inter-PLMN signaling on N32 via PRINS**
- **Securing inter-PLMN signaling via TLS or PRINS ensures integrity**
- **Securing inter-PLMN signaling via TLS or PRINS ensures confidentiality**

Factoid: IPUPS secures N9 GTP-U traffic between UPFs with IPsec-based filtering to avoid tunnel spoofing.

- **IPUPS secures N9 GTP-U traffic between UPFs**
- **IPUPS uses IPsec-based filtering**
- **IPsec-based filtering avoids tunnel spoofing**

Factoid: NSSAAF enables slice-level authentication using EAP-based credentials for each network slice.

- **NSSAAF enables slice-level authentication**
- **Slice-level authentication uses EAP-based credentials**
- **EAP-based credentials are used for each network slice**

Factoid: TLS 1.2/1.3 configurations require AEAD cipher suites and OCSP, removing legacy weak cipher support.

- **TLS 1.2/1.3 configurations require AEAD cipher suites**
- **TLS 1.2/1.3 configurations require OCSP**
- **TLS 1.2/1.3 configurations remove legacy weak cipher support**

Factoid: Zero-Trust security within 5G SBA demands explicit flow definitions, strict policies, and continuous monitoring.

- **Zero-Trust security demands explicit flow definitions**
- **Zero-Trust security demands strict policies**
- **Zero-Trust security demands continuous monitoring**

Factoid: Slice management interfaces leverage OAuth for access control, combined with mTLS for secure management messaging.

- **Slice management interfaces** *leverage OAuth*
Slice management interfaces *use mTLS*
- **OAuth** *provides access control*
mTLS *secures management messaging*

Factoid: UE identity is encrypted as SUCI using public-key encryption, protecting SUPI from over-the-air exposure.

- **UE identity** *is encrypted as SUCI*
- **Encryption as SUCI** *uses public-key encryption*
- **Encryption as SUCI** *protects SUPI from over-the-air exposure*

Factoid: New SBA functions SEAF, AUSF, UDM/ARPF, and SIDF coordinate to authenticate UE and manage key derivation.

- **SEAF, AUSF, UDM/ARPF, and SIDF** *coordinate to authenticate UE*
- **SEAF, AUSF, UDM/ARPF, and SIDF** *coordinate to manage key derivation*

Factoid: 5G key hierarchy includes KAUSF, KSEAF, and KAMF, offering deeper security separation than 4G.

- **5G key hierarchy** *includes KAUSF*
- **5G key hierarchy** *includes KSEAF*
- **5G key hierarchy** *includes KAMF*
- **5G key hierarchy** *offers deeper security separation than 4G*

Factoid: A single 5G authentication session can establish multiple security contexts across access types.

- **5G authentication session** *establishes multiple security contexts*
- **Security contexts** *span different access types*

Factoid: EAP-TLS leverages X.509 certificates for authentication without requiring USIM, suiting BYOD or enterprise devices.

- **EAP-TLS** *leverages X.509 certificates*
- **EAP-TLS** *authenticates without requiring USIM*
- **EAP-TLS** *suits BYOD devices*
- **EAP-TLS** *suits enterprise devices*

Factoid: EAP-AKA' is a symmetric key-based EAP method offering similar trust as 5G-AKA but via EAP exchange.

- **EAP-AKA'** *is symmetric key-based EAP method*
- **EAP-AKA'** *offers trust equivalent to 5G-AKA*
- **Trust equivalent to 5G-AKA** *is delivered via EAP exchange*

Factoid: SIDF decrypts SUCI to SUPI, enabling identifier confidentiality with public-key protection.

- **SIDF decrypts SUCI to SUPI**
- **Decrypting SUCI to SUPI enables identifier confidentiality**
- **Identifier confidentiality is protected by public-key protection**

Factoid: SEAF uses mutual TLS and PRINS to secure inter-PLMN communication and prevent downgrade attacks.

- **SEAF uses mutual TLS**
- **SEAF uses PRINS**
- **Using mutual TLS and PRINS secures inter-PLMN communication**
- **Using mutual TLS and PRINS prevents downgrade attacks**

Factoid: Mapping of keys: KAUSF → KSEAF → KAMF ensures layered trust and key separation in 5G.

- **KAUSF derives KSEAF**
- **KSEAF derives KAMF**
- **Mapping of keys (KAUSF → KSEAF → KAMF) ensures layered trust**
- **Mapping of keys (KAUSF → KSEAF → KAMF) ensures key separation in 5G**

Factoid: SUPI encryption, SBA authentication functions, and deeper key derivation collectively enhance privacy and home-network control in 5G.

- **SUPI encryption enhances privacy**
- **SBA authentication functions enhance privacy**
- **Deeper key derivation enhances privacy**
- **SUPI encryption enhances home-network control**
- **SBA authentication functions enhance home-network control**
- **Deeper key derivation enhances home-network control**
- **Enhancing privacy and home-network control occurs in 5G**

Factoid: EAP-TLS eliminates symmetric key dependency but introduces certificate lifecycle overhead—trading key management for lifecycle complexity.

- **EAP-TLS eliminates symmetric key dependency**
- **EAP-TLS introduces certificate lifecycle overhead**
- **Eliminating symmetric key dependency trades key management for lifecycle complexity**
- **Introducing certificate lifecycle overhead trades key management for lifecycle complexity**

Factoid: 5G-AKA and EAP-AKA' are mandatory primary authentication protocols across N1, N12, and N13 interfaces.

- **5G-AKA is a mandatory primary authentication protocol across N1 interface**
- **5G-AKA is a mandatory primary authentication protocol across N12 interface**

- **5G-AKA** is a mandatory primary authentication protocol across **N13 interface**
- **EAP-AKA'** is a mandatory primary authentication protocol across **N1 interface**
- **EAP-AKA'** is a mandatory primary authentication protocol across **N12 interface**
- **EAP-AKA'** is a mandatory primary authentication protocol across **N13 interface**

Factoid: EAP enables unified authentication regardless of whether a UE accesses via 3GPP or non-3GPP RAT.

- **EAP enables unified authentication**
- **Unified authentication applies regardless of UE access via 3GPP RAT**
- **Unified authentication applies regardless of UE access via non-3GPP RAT**

Factoid: Security threats like replay, MitM, and downgrade attacks target N1 and N12 if proper cryptographic protections are absent.

- **Security threats include replay attacks**
- **Security threats include MitM attacks**
- **Security threats include downgrade attacks**
- **Security threats target N1 interface if proper cryptographic protections are absent**
- **Security threats target N12 interface if proper cryptographic protections are absent**

Factoid: N1 and N2 interfaces must be protected using IPsec for non-3GPP access, and NAS integrity algorithms over 3GPP access.

- **N1 interface must be protected using IPsec over non-3GPP access**
- **N1 interface must be protected using NAS integrity algorithms over 3GPP access**
- **N2 interface must be protected using IPsec over non-3GPP access**
- **N2 interface must be protected using NAS integrity algorithms over 3GPP access**

Factoid: Control-plane interfaces to AUSF and UDM (N12/N13) require mutual TLS to enforce authentication and confidentiality.

- **N12 interface requires mutual TLS**
- **N13 interface requires mutual TLS**
- **Mutual TLS enforces authentication**
- **Mutual TLS enforces confidentiality**

Factoid: Inter-domain interfaces such as N32/N16 need TLS with PRINS or equivalent to ensure end-to-end signaling integrity.

- **Inter-domain interface N32 needs TLS with PRINS or equivalent**
- **Inter-domain interface N16 needs TLS with PRINS or equivalent**
- **TLS with PRINS or equivalent ensures end-to-end signaling integrity**

Factoid: Formal verification tools like Tamarin uncovered potential linkability vulnerabilities in 5G-AKA unless mitigations are implemented.

- **Formal verification tools like Tamarin** *uncovered potential linkability vulnerabilities*
- **Potential linkability vulnerabilities** *occur in 5G-AKA*
- **Potential linkability vulnerabilities** *remain unless mitigations are implemented*

Factoid: Even when authentication exchanges are secure, implementations must ensure serving network binding via proper key derivation to prevent impersonation.

- **Authentication exchanges** *can be secure*
Implementations *must ensure serving network binding via proper key derivation*
- **Proper key derivation for serving network binding** *prevents impersonation*

Factoid: Non-3GPP access (e.g., Wi-Fi) requires IPsec tunnels (e.g., NWu) to secure NAS and user-plane traffic.

- **Non-3GPP access (e.g., Wi-Fi)** *requires IPsec tunnels (e.g., NWu)*
IPsec tunnels (e.g., NWu) *secure NAS traffic*
- **IPsec tunnels (e.g., NWu)** *secure user-plane traffic*

Factoid: Upgrade or downgrade prevention is critical—NAS and interface-level protections must enforce version/context awareness to avoid downgrade exploits.

- **Upgrade or downgrade prevention** *is critical*
- **NAS protections** *must enforce version awareness*
- **Interface-level protections** *must enforce context awareness*
- **Enforcing version and context awareness** *avoids downgrade exploits*

Factoid: UE establishes an IPsec signaling SA (NWu) with the N3IWF over IKEv2 and EAP-5G when connected via untrusted WLAN.

- **UE** *establishes IPsec signaling SA (NWu)*
- **IPsec signaling SA (NWu)** *is established with N3IWF*
- **IPsec signaling SA (NWu)** *uses IKEv2*
- **IPsec signaling SA (NWu)** *uses EAP-5G*
- **UE** *is connected via untrusted WLAN*

Factoid: EAP-5G encapsulates NAS-based authentication (5G-AKA/EAP-AKA') within IKEv2 exchanges over non-3GPP access.

- **EAP-5G** *encapsulates NAS-based authentication (5G-AKA/EAP-AKA')*
- **NAS-based authentication** *includes 5G-AKA*
- **NAS-based authentication** *includes EAP-AKA'*
- **Encapsulation within IKEv2 exchanges** *occurs over non-3GPP access*

Factoid: User-plane over Wi-Fi uses separate IPsec child SA(s) after PDU session establishment, with encryption based on trust.

- **User-plane over Wi-Fi** *uses separate IPsec child SA(s)*
- **Separate IPsec child SA(s)** *are established after PDU session establishment*

Factoid: Trusted WLAN setups use layer-2 authentication (802.1x via TNAP) followed by IPsec signaling SA with NULL encryption to avoid double encryption.

- **Trusted WLAN setups** *use layer-2 authentication (802.1x via TNAP)*
- **Layer-2 authentication (802.1x via TNAP)** *is followed by IPsec signaling SA*
- **IPsec signaling SA** *employs NULL encryption*
- **NULL encryption** *avoids double encryption*

Factoid: Ta/Yw interfaces connect WLAN APs to TNGF/TWIF gateways, existing outside 3GPP scope but essential for trusted Wi-Fi integration.

- **Ta/Yw interfaces** *connect WLAN APs to TNGF/TWIF gateways*
- **Ta/Yw interfaces** *exist outside 3GPP scope*
- **Ta/Yw interfaces** *are essential for trusted Wi-Fi integration*

Factoid: Wi-Fi only devices without USIM rely on certificate-based EAP-TLS/EAP-TTLS, especially in private SNPN environments.

- **Wi-Fi only devices without USIM** *rely on certificate-based EAP-TLS*
- **Wi-Fi only devices without USIM** *rely on certificate-based EAP-TTLS*
- **Certificate-based EAP-TLS/EAP-TTLS** *are used in private SNPN environments*

Factoid: ATSSS enables steering or splitting of traffic between 3GPP and Wi-Fi based on central policy from SMF/PCF.

- **ATSSS** *enables steering or splitting of traffic between 3GPP and Wi-Fi*
- **Steering or splitting of traffic** *is based on central policy*
- **Central policy** *is provided by SMF and PCF*

Factoid: NAS messages over Wi-Fi rely on TCP/IP transport within IPsec tunnels to ensure message reliability and ordering.

- **NAS messages over Wi-Fi** *rely on TCP/IP transport*
- **TCP/IP transport** *occurs within IPsec tunnels*
- **IPsec tunnels** *ensure message reliability*
- **IPsec tunnels** *ensure message ordering*

Factoid: Building trust zones in Wi-Fi (untrusted vs trusted) determines whether IPsec encryption is NULL or standard to prevent redundancy.

- **Building trust zones in Wi-Fi** *determines IPsec encryption type*
IPsec encryption type *is NULL encryption in trusted zones*
- **IPsec encryption type** *is standard encryption in untrusted zones*
- **Using NULL encryption** *prevents redundant encryption*

Factoid: Integration of EAP-5G, IKEv2, and IPsec ensures secure control and data-plane transport across Wi-Fi links.

- **Integration of EAP-5G, IKEv2, and IPsec** *ensures secure control-plane transport*
- **Integration of EAP-5G, IKEv2, and IPsec** *ensures secure data-plane transport*
- **Secure control-plane transport** *occurs across Wi-Fi links*
- **Secure data-plane transport** *occurs across Wi-Fi links*

Factoid: Primary authentication in industrial 5G mandates use of 5G-AKA, EAP-AKA, or EAP-TLS to verify device identity.

- **Primary authentication in industrial 5G** *mandates use of 5G-AKA*
- **Primary authentication in industrial 5G** *mandates use of EAP-AKA*
- **Primary authentication in industrial 5G** *mandates use of EAP-TLS*
- **Use of 5G-AKA, EAP-AKA, or EAP-TLS** *verifies device identity*

Factoid: Secondary EAP authentication enables access to external services or slice-specific networks using alternate credentials.

- **Secondary EAP authentication** *enables access to external services*
- **Secondary EAP authentication** *enables access to slice-specific networks*
- **Access to external services** *uses alternate credentials*
- **Access to slice-specific networks** *uses alternate credentials*

Factoid: UDM/HSS securely stores USIM credentials and supports provisioning workflows for industrial devices.

- **UDM/HSS** *securely stores USIM credentials*
- **UDM/HSS** *supports provisioning workflows for industrial devices*

Factoid: SIM provisioning in industrial environments requires secure lifecycle procedures including revocation and updates.

- **SIM provisioning in industrial environments** *requires secure lifecycle procedures*
- **Secure lifecycle procedures** *include revocation*
- **Secure lifecycle procedures** *include updates*

Factoid: Industrial RAN requires IPsec tunnels (NWu/NWt) to protect NAS signaling and data-plane, with SPI protections.

- **Industrial RAN** *requires IPsec tunnels (NWu/NWt)*
- **IPsec tunnels (NWu/NWt)** *protect NAS signaling*
- **IPsec tunnels (NWu/NWt)** *protect data-plane*
- **IPsec tunnels (NWu/NWt)** *include SPI protections*

Factoid: Trusted industrial deployments may use NULL-encryption under IPsec tunnels to avoid double-layer encryption overhead.

- **Trusted industrial deployments** *may use NULL-encryption under IPsec tunnels*
- **NULL-encryption under IPsec tunnels** *avoids double-layer encryption overhead*

Factoid: Industrial 5G prioritizes availability and deterministic latency over confidentiality to maintain real-time control.

- **Industrial 5G** *prioritizes availability over confidentiality*
- **Industrial 5G** *prioritizes deterministic latency over confidentiality*
- **Prioritizing availability and deterministic latency over confidentiality** *supports maintaining real-time control*

Factoid: TSN traffic must be mapped into 5G slices with real-time scheduling and IPsec security for industrial applications.

- **TSN traffic** *must be mapped into 5G slices*
- **TSN traffic mapping into 5G slices** *requires real-time scheduling*
- **TSN traffic mapping into 5G slices** *requires IPsec security*
- **TSN traffic mapping into 5G slices** *is used for industrial applications*

Factoid: Industrial 5G must mitigate wireless-specific threats (e.g., jamming, spoofing) via hardened SIM provisioning and tunnel integrity.

- **Industrial 5G** *must mitigate wireless-specific threats*
- **Wireless-specific threats** *include jamming*
- **Wireless-specific threats** *include spoofing*
- **Mitigation of wireless-specific threats** *is achieved via hardened SIM provisioning*
- **Mitigation of wireless-specific threats** *is achieved via tunnel integrity*

Factoid: Slice-specific EAP authentication is critical where multiple industrial operations share the same 5G infrastructure.

- **Slice-specific EAP authentication** *is critical for multiple industrial operations sharing the same 5G infrastructure*
- **Multiple industrial operations** *share the same 5G infrastructure*

Factoid: Intent-Based Networking (IBN) enables high-level operator intent to be translated into network policy via NIT in 5G Pr-purposed NMA systems.

- **Intent-Based Networking (IBN)** *enables high-level operator intent to be translated into network policy*
- **Translation of high-level operator intent into network policy** *occurs via NIT*
- **Translation via NIT** *occurs in 5G Pr-purposed NMA systems*

Factoid: The architectural framework integrates NWDAF as IBN Analyzer to create a closed-loop system—monitoring, validation, and policy refinement.

- **Architectural framework** *integrates NWDAF as IBN Analyzer*
- **Integration of NWDAF as IBN Analyzer** *creates closed-loop system*
- **Closed-loop system** *includes monitoring*
- **Closed-loop system** *includes validation*
- **Closed-loop system** *includes policy refinement*

Factoid: NIT uses data-model mapping and optionally NLP to convert user intents (e.g., slice QoS, IoT data collection) into NF-specific configurations.

- **NIT** *uses data-model mapping*
- **NIT** *optionally uses NLP*
- **Data-model mapping and NLP** *convert user intents into NF-specific configurations*
- **User intents** *include slice QoS*
- **User intents** *include IoT data collection*

Factoid: IBN Controller dispatches translated policy rules to VNFs, CNFs, or PNFs using the NF-facing interface.

- **IBN Controller** *dispatches translated policy rules to VNFs*
- **IBN Controller** *dispatches translated policy rules to CNFs*
- **IBN Controller** *dispatches translated policy rules to PNFs*
- **Dispatching translated policy rules to VNFs, CNFs, or PNFs** *uses NF-facing interface*

Factoid: Network telemetry is collected via a monitoring interface from NFs, analyzed by NWDAF, and used to audit intent success.

- **Network telemetry** *is collected via monitoring interface from Network Functions (NFs)*
- **Network telemetry** *is analyzed by NWDAF*
- **Analysis of network telemetry by NWDAF** *is used to audit intent success*

Factoid: Consumer-facing interface accepts user intent defined via 3GPP TS-28.312 intent schemas.

- **Consumer-facing interface** *accepts user intent*
- **User intent** *is defined via 3GPP TS-28.312 intent schemas*

Factoid: Use cases like IoT aggregation and V2X QoS demonstrate how intent drives slice creation and SLA enforcement end-to-end.

- **Use cases** *include IoT aggregation*
- **Use cases** *include V2X QoS*
- **IoT aggregation use case** *demonstrates intent-driven slice creation*
- **V2X QoS use case** *demonstrates intent-driven slice creation*
- **IoT aggregation use case** *demonstrates intent-driven SLA enforcement end-to-end*
- **V2X QoS use case** *demonstrates intent-driven SLA enforcement end-to-end*

- **Intent** *drives slice creation*
- **Intent** *drives SLA enforcement end-to-end*

Factoid: Closed-loop automation ensures intent enforcement accuracy by verifying via telemetry and updating policies as needed.

- **Closed-loop automation** *ensures intent enforcement accuracy*
- **Closed-loop automation** *verifies intent enforcement accuracy via telemetry*
- **Closed-loop automation** *updates policies as needed*

Factoid: TS 28.312 defines an SBMA-based Intent-Driven Management Service (MnS) that allows intent producers to accept, fulfill, and report on network or service intents.

- **TS 28.312** *defines SBMA-based Intent-Driven Management Service (MnS)*
- **SBMA-based Intent-Driven Management Service (MnS)** *allows intent producers to accept network or service intents*
- **SBMA-based Intent-Driven Management Service (MnS)** *allows intent producers to fulfill network or service intents*
- **SBMA-based Intent-Driven Management Service (MnS)** *allows intent producers to report on network or service intents*

Factoid: An intent is decoupled from execution—it specifies **what** (intent expectations), not **how** (policies/actions), aligning with a model-driven SBMA approach.

- **Intent** *is decoupled from execution*
- **Intent** *specifies intent expectations*
- **Intent** *does not specify policies/actions*
- **Decoupling of intent from execution** *aligns with model-driven SBMA approach*

Factoid: Intent operations include create, modify, query, delete, activate, and deactivate, each managing intent lifecycle states.

- **Intent operations** *include create*
- **Intent operations** *include modify*
- **Intent operations** *include query*
- **Intent operations** *include delete*
- **Intent operations** *include activate*
- **Intent operations** *include deactivate*
- **Intent operations** *manage intent lifecycle states*
- **Managing intent lifecycle states** *applies to each intent operation*

Factoid: Intent content includes expectations (e.g., throughput targets, area coverage) linked to context objects such as PLMN and TAC.

- **Intent content** *includes expectations*
- **Expectations** *include throughput targets*
- **Expectations** *include area coverage*
- **Intent content** *is linked to context objects*

- **Context objects** *include* **PLMN**
- **Context objects** *include* **TAC**

Factoid: TR 28.912 enhances TS 28.312 with verification reports, conflict resolution, feasibility checks, AI/ML mapping, and intent-driven SON orchestration.

- **TR 28.912** *enhances* **TS 28.312**
- **TR 28.912** *provides* **verification reports**
- **TR 28.912** *provides* **conflict resolution**
- **TR 28.912** *provides* **feasibility checks**
- **TR 28.912** *provides* **AI/ML mapping**
- **TR 28.912** *provides* **intent-driven SON orchestration**

Factoid: Intent producers may implement actions via rule-based, closed-loop, or AI/ML-driven mechanisms.

- **Intent producers** *implement actions via* **rule-based mechanisms**
- **Intent producers** *implement actions via* **closed-loop mechanisms**
- **Intent producers** *implement actions via* **AI/ML-driven mechanisms**

Factoid: 3GPP intent services are interoperable with TM-Forum and O-RAN intent frameworks to enable cross-domain automation.

- **3GPP intent services** *are interoperable with* **TM-Forum intent frameworks**
- **3GPP intent services** *are interoperable with* **O-RAN intent frameworks**
- **Interoperability** *enables* **cross-domain automation**

Factoid: Intent expectations can include energy savings targets balanced against service performance metrics.

- **Intent expectations** *can include* **energy savings targets**
- **Intent expectations** *can include* **service performance metrics**
- **Energy savings targets** *are balanced against* **service performance metrics**

Factoid: Release 19 (TR 28.914) extends intent models, including RAN-level intent autopolicies for 6G readiness.

- **Release 19 (TR 28.914)** *extends* **intent models**
- **Extended intent models** *include* **RAN-level intent autopolicies**
- **RAN-level intent autopolicies** *support* **6G readiness**

Factoid: RadioNetworkExpectation in TS 28.312 includes attributes like coverageAreaPolygon, RAT types, target throughput, and latency thresholds.

- **RadioNetworkExpectation** *includes* **coverageAreaPolygon**
- **RadioNetworkExpectation** *includes* **RAT types**
- **RadioNetworkExpectation** *includes* **target throughput**
- **RadioNetworkExpectation** *includes* **latency thresholds**

Factoid: Network intents are high-level objectives defined as IntentTargets, IntentExpectations, and context constraints.

- **Network intents** *are high-level objectives*
- **Network intents** *are defined as IntentTargets*
- **Network intents** *are defined as IntentExpectations*
- **Network intents** *are defined as context constraints*

Factoid: End-to-end IDM architecture includes ingestion, translation (via ILUs/ILL), orchestration, assurance, and reporting components.

- **End-to-end IDM architecture** *includes ingestion components*
- **End-to-end IDM architecture** *includes translation components*
- **End-to-end IDM architecture** *includes orchestration components*
- **End-to-end IDM architecture** *includes assurance components*
- **End-to-end IDM architecture** *includes reporting components*
- **Translation components** *are implemented via ILUs/ILL*

Factoid: Intent lifecycle follows

Create → Refine → Validate → Fulfill → Monitor → Assure → Report → Drift → Correct.

Factoid: Inner loop automates fulfillment and assurance; outer loop involves user refinement and intent updates.

- **Inner loop** *automates fulfillment*
- **Inner loop** *automates assurance*
- **Outer loop** *involves user refinement*
- **Outer loop** *involves intent updates*

Factoid: Conflict detection mechanisms identify overlapping or contradictory intents, requiring resolution rules.

- **Conflict detection mechanisms** *identify overlapping intents*
- **Conflict detection mechanisms** *identify contradictory intents*
- **Identifying overlapping or contradictory intents** *requires resolution rules*

Factoid: Intent assurance semantically maps network telemetry to intent expectations to detect drift.

- **Intent assurance** *semantically maps network telemetry to intent expectations*
Semantic mapping of network telemetry to intent expectations *detects drift*

_____ BATCH 6 GEMINI _____

Factoid: Assurance systems may deploy LLMs to recommend policy changes when intent drift arises.

- **Assurance systems** *may deploy LLMs*
- **LLMs** *recommend policy changes*
- **Recommendation of policy changes** *occurs when intent drift arises*

Factoid: Intent frameworks align with TM Forum, 3GPP, and ETSI ZSM to enable multi-domain interoperability.

- **Intent frameworks** *align with TM Forum*
- **Intent frameworks** *align with 3GPP*
- **Intent frameworks** *align with ETSI ZSM*
- **Alignment with TM Forum, 3GPP, and ETSI ZSM** *enables multi-domain interoperability*

Factoid: Intent ingestion supports both user-interactive refinement and automated processing toward machine-actionable formats.

- **Intent ingestion** *supports user-interactive refinement*
- **Intent ingestion** *supports automated processing*
- **Automated processing** *targets machine-actionable formats*

Factoid: Intent translation uses reusable Intent Logic Units from an Intent Logic Library for mapping abstract intent to policies.

- **Intent translation** *uses reusable Intent Logic Units (ILUs)*
- **Reusable Intent Logic Units (ILUs)** *are sourced from Intent Logic Library*
- **Intent Logic Units (ILUs)** *map abstract intent to policies*

Factoid: Reporting abstracts low-level telemetry into intent-aligned summaries for operator decision-making.

- **Reporting** *abstracts low-level telemetry*
- **Reporting** *generates intent-aligned summaries*
- **Intent-aligned summaries** *support operator decision-making*

Factoid: Assurance loop may trigger corrective workflows or escalate to user when semantic compliance falls below thresholds.

- **Assurance loop** *may trigger corrective workflows*
- **Assurance loop** *may escalate to user when semantic compliance falls below thresholds*

Factoid: A natural-language intent interface parses English statements into formal policy definitions for private 5G management.

- **Natural-language intent interface** *parses English statements*
- **Natural-language intent interface** *translates English statements into formal policy definitions*

- **Formal policy definitions** *enable private 5G management*

Factoid: The 5G-CLARITY platform includes an Intent Engine plus MLFO and Data Lake for telemetry-driven intent mapping.

- **5G-CLARITY platform** *includes Intent Engine*
- **5G-CLARITY platform** *includes MLFO*
- **5G-CLARITY platform** *includes Data Lake*
- **Intent Engine** *is used for telemetry-driven intent mapping*
- **MLFO** *is used for telemetry-driven intent mapping*
- **Data Lake** *is used for telemetry-driven intent mapping*

Factoid: ML models convert high-level intents into API workflows targeting NFV or RAN orchestrators.

- **ML models** *convert high-level intents into API workflows*
- **API workflows** *target NFV orchestrators*
- **API workflows** *target RAN orchestrators*

Factoid: Use cases—slice provisioning, indoor positioning, and service deployment—were benchmarked for provisioning latency.

- **Use cases** *include slice provisioning*
- **Use cases** *include indoor positioning*
- **Use cases** *include service deployment*
- **Use cases** *were benchmarked for provisioning latency*

Factoid: Platform aligns with 3GPP TS 28.312, ETSI ZSM, and TM Forum intent frameworks for cross-SDO integration.

- **Platform** *aligns with 3GPP TS 28.312*
- **Platform** *aligns with ETSI ZSM*
- **Platform** *aligns with TM Forum intent frameworks*
- **Alignment with 3GPP TS 28.312, ETSI ZSM, and TM Forum intent frameworks** *enables cross-SDO integration*

Factoid: Future work includes embedding LLMs to automate intent-to-workflow translation and reduce manual configuration effort.

- **Future work** *includes embedding LLMs*
- **Embedding LLMs** *automates intent-to-workflow translation*
- **Automating intent-to-workflow translation** *reduces manual configuration effort*

Factoid: Intent systems govern RAN, core, and data-center domains, adjusting parameters in RUs, DUs, CUs to meet performance or power objectives.

- **Intent systems** *govern RAN domain*
- **Intent systems** *govern core domain*
- **Intent systems** *govern data-center domain*

- Intent systems *adjust* parameters in RUs
- Intent systems *adjust* parameters in DUs
- Intent systems *adjust* parameters in CUs
- Adjusting parameters in RUs, DUs, CUs *aims to meet* performance objectives
- Adjusting parameters in RUs, DUs, CUs *aims to meet* power objectives

Factoid: Intents are constrained to pre-defined capabilities; they cannot create new functionality beyond the data model.

- Intents *are constrained to* pre-defined capabilities
- Intents *cannot create* new functionality beyond the data model

Factoid: ETSI ENI, ZSM, TM Forum IG1253, and 3GPP TS 28.312 are the primary SDOs driving intent-driven network standards.

- ETSI ENI, ETSI ZSM, TM Forum IG1253, and 3GPP TS 28.312 *are* primary SDOs driving intent-driven network standards

Factoid: Intent architectures require a policy/action engine, monitoring/assurance subsystem, and optional AI/ML-enhanced mapping.

- Intent architectures *require* policy/action engine
- Intent architectures *require* monitoring/assurance subsystem
- Intent architectures *include* optional AI/ML-enhanced mapping

Factoid: Policy translation regenerates and updates when intents change to keep intent-data models and actions aligned.

- Policy translation *regenerates* intent-data models when intents change
- Policy translation *updates* actions when intents change
- Regeneration and updates *keep* intent-data models and actions aligned*

Factoid: Gen-2 orchestration must span PNFs, VNFs, CNFs, VMs, containers, and serve across private/public/edge clouds.

- Gen-2 orchestration *must span* PNFs
- Gen-2 orchestration *must span* VNFs
- Gen-2 orchestration *must span* CNFs
- Gen-2 orchestration *must span* VMs
- Gen-2 orchestration *must span* containers
- Gen-2 orchestration *must serve across* private clouds
- Gen-2 orchestration *must serve across* public clouds
- Gen-2 orchestration *must serve across* edge clouds

Factoid: Closed-loop telemetry is mandatory for dynamic reconfiguration and SLA compliance within intent-driven systems.

- Closed-loop telemetry *is mandatory for* dynamic reconfiguration

- **Closed-loop telemetry** *is mandatory for SLA compliance*
- **Dynamic reconfiguration and SLA compliance** *occur within intent-driven systems*

Factoid: IETF's intent model uses nested loops: an inner autonomous control loop and an outer user-in-the-loop refinement loop.

- **IETF's intent model** *uses nested loops*
- **Nested loops** *include inner autonomous control loop*
- **Nested loops** *include outer user-in-the-loop refinement loop*

Factoid: Vendor data models vary, so intent systems need abstraction bridges across polymorphic device APIs and configurations.

- **Vendor data models** *vary*
- **Intent systems** *need abstraction bridges*
- **Abstraction bridges** *span polymorphic device APIs*
- **Abstraction bridges** *span configurations*

Factoid: AI/ML can enhance intent mapping and assurance but remains bounded by intended expressiveness of policy and data models.

- **AI/ML** *can enhance intent mapping*
- **AI/ML** *can enhance assurance*
- **AI/ML** *remains bounded by intended expressiveness of policy and data models*

Factoid: A slice is a customer SLA-backed end-to-end logical network running on shared infrastructure with performance bounds.

- **Slice** *is a customer SLA-backed end-to-end logical network*
- **End-to-end logical network** *runs on shared infrastructure*
- **End-to-end logical network** *has performance bounds*

Factoid: Slices comprise both dedicated and shared resources and must be logically isolated from one another.

- **Slices** *comprise dedicated resources*
- **Slices** *comprise shared resources*
- **Dedicated and shared resources** *must be logically isolated from one another*

Factoid: Business bundles allow operators to combine multiple slice types (e.g., URLLC + eMBB) under a unified SLA.

- **Business bundles** *allow operators to combine multiple slice types under a unified SLA*
- **Operators** *combine URLLC and eMBB slice types*
- **Combination of multiple slice types** *occurs under unified SLA*

Factoid: Generic Slice Templates define attribute domains (e.g., latency range, reliability targets) for network slices.

- **Generic Slice Templates** *define attribute domains*
- **Attribute domains** *include latency range*
- **Attribute domains** *include reliability targets*
- **Attribute domains** *apply to network slices*

Factoid: Network Slice Type instantiates a GST with specific values tailored to a vertical use case's requirements.

- **Network Slice Type** *instantiates Generic Slice Template (GST)*
- **Network Slice Type** *applies specific values tailored to a vertical use case's requirements*

Factoid: Slices can span RAN, Core, and Transport domains and be shared across multiple operators for roaming support.

- **Slices** *can span RAN, Core, and Transport domains*
- **Slices** *can be shared across multiple operators*
- **Sharing slices across multiple operators** *supports roaming support*

Factoid: Operators publish standardized slice blueprints to visited PLMNs to preserve SLAs and enable slice continuity.

- **Operators** *publish standardized slice blueprints*
- **Standardized slice blueprints** *are published to visited PLMNs*
- **Publishing standardized slice blueprints to visited PLMNs** *preserves SLAs*
- **Publishing standardized slice blueprints to visited PLMNs** *enables slice continuity*

Factoid: Slice customization is capped by GST/NEST—customers cannot exceed operator-defined attribute ranges.

- **Slice customization** *is capped by GST/NEST*
- **Customers** *cannot exceed operator-defined attribute ranges*

Factoid: Use cases like V2X and industrial IoT require elasticity and strict traffic isolation from slices.

- **Use cases like V2X and industrial IoT** *require elasticity*
- **Use cases like V2X and industrial IoT** *require strict traffic isolation from slices*

Factoid: Regulatory models must handle neutrality, data sovereignty, and cross-border slicing SLAs.

- **Regulatory models** *must handle neutrality*
- **Regulatory models** *must handle data sovereignty*

- **Regulatory models** *must handle cross-border slicing SLAs*

Factoid: Slice descriptors and SLAs are central to operator interoperability and roaming in sliced 5G environments.

- **Slice descriptors** *are central to operator interoperability*
- **Slice descriptors** *are central to roaming in sliced 5G environments*
- **SLAs** *are central to operator interoperability*
- **SLAs** *are central to roaming in sliced 5G environments*
- **Operator interoperability** *occurs in sliced 5G environments*
- **Roaming** *occurs in sliced 5G environments*

Factoid: 5G slicing enables dynamic per-slice traffic treatment, creating potential information asymmetries vis-à-vis regulators.

- **5G slicing** *enables dynamic per-slice traffic treatment*
- **Dynamic per-slice traffic treatment** *creates potential information asymmetries*
- **Potential information asymmetries** *occur vis-à-vis regulators*

Factoid: NWDAF can expose per-slice KPIs (packet loss, delay, QoS metrics) to external regulators via NEF.

- **NWDAF** *exposes per-slice KPIs*
- **Per-slice KPIs** *include packet loss*
- **Per-slice KPIs** *include delay*
- **Per-slice KPIs** *include QoS metrics*
- **NWDAF** *exposes per-slice KPIs to external regulators*
- **Exposure to external regulators** *occurs via NEF*

Factoid: NWDAF employs a consumer–producer model with standardized interfaces for periodic KPI and analytics distribution.

- **NWDAF** *employs consumer–producer model*
- **Consumer–producer model** *uses standardized interfaces*
- **Standardized interfaces** *support periodic KPI distribution*
- **Standardized interfaces** *support analytics distribution*

Factoid: Analytics functions include historical trend analysis, predictive modeling, anomaly detection, and QoS sustainability reports.

- **Analytics functions** *include historical trend analysis*
- **Analytics functions** *include predictive modeling*
- **Analytics functions** *include anomaly detection*
- **Analytics functions** *include QoS sustainability reports*

Factoid: A proof-of-concept showed eMBB slice analytics can be used by regulators to detect traffic differentiation.

- **Proof-of-concept** *showed eMBB slice analytics can be used by regulators*
- **eMBB slice analytics** *enables detection of traffic differentiation*
- **Regulators** *detect traffic differentiation*

Factoid: NWDAF enables real-time regulatory monitoring of net neutrality compliance based on slice-level transparency.

- **NWDAF** *enables real-time regulatory monitoring of net neutrality compliance*
- **Real-time regulatory monitoring of net neutrality compliance** *is based on slice-level transparency*

Factoid: Deployment options include SNPN, PNI-NPN, and hybrid models with shared RAN and TSN integration.

- **Deployment options** *include SNPN*
Deployment options *include PNI-NPN*
- **Deployment options** *include hybrid models*
- **Hybrid models** *use shared RAN*
- **Hybrid models** *use TSN integration*

Factoid: Rel-16 slicing lacks native support for industrial Ethernet integration; enhancements planned in Rel-17.

- **Rel-16 slicing** *lacks native support for industrial Ethernet integration*
- **Enhancements for industrial Ethernet integration** *are planned in Rel-17*

Factoid: Security zones and VLAN-based segmentation enable dynamic isolation of traffic within slices.

- **Security zones** *enable dynamic isolation of traffic within slices*
- **VLAN-based segmentation** *enables dynamic isolation of traffic within slices*

Factoid: Operator policy defines tenant-customizable parameters like coverage area and TSN support levels.

- **Operator policy** *defines tenant-customizable parameters*
- **Tenant-customizable parameters** *include coverage area*
- **Tenant-customizable parameters** *include TSN support levels*

Factoid: Qualitative mapping of use-case requirements (latency, reliability) to NPN architecture types informs deployment choice.

- **Use-case requirements** *include latency*
- **Use-case requirements** *include reliability*
- **Use-case requirements** *are qualitatively mapped to NPN architecture types*
- **Qualitative mapping of use-case requirements to NPN architecture types** *informs deployment choice*

Factoid: Slice access can be geographically constrained using TAC or other location identifiers.

- **Slice access** *can be geographically constrained using TAC*
- **Slice access** *can be geographically constrained using other location identifiers*

Factoid: Edge cloud and TSN functions are co-located based on geo-proximity to meet tight latency targets.

- **Edge cloud and TSN functions** *are co-located based on geo-proximity*
- **Co-location based on geo-proximity** *enables meeting tight latency targets*

Factoid: Dynamic traffic-driven forwarding rules support adaptive slice isolation and performance optimization.

- **Dynamic traffic-driven forwarding rules** *support adaptive slice isolation*
- **Dynamic traffic-driven forwarding rules** *support performance optimization*

Factoid: A Maximum Slice Data Rate (UL/DL) can be configured by the operator per S-NSSAI to cap slice-wide throughput.

- **Maximum Slice Data Rate (UL/DL)** *can be configured by operator*
- **Maximum Slice Data Rate (UL/DL)** *is configured per S-NSSAI*
- **Configuring Maximum Slice Data Rate (UL/DL)** *caps slice-wide throughput*

Factoid: PCF enforces Maximum Slice Data Rate by rejecting SM Policy or PDU session establishment if slice budget is exceeded.

- **PCF enforces Maximum Slice Data Rate**
- **PCF rejects SM Policy establishment**
- **PCF rejects PDU session establishment**
- **Rejection of SM Policy or PDU session establishment** *occurs if slice budget is exceeded*

Factoid: With NWDAF integration, PCF uses slice usage data (volume and duration) to apply or relax constraints dynamically.

- **PCF uses slice usage data (volume and duration)**
- **Slice usage data** *is provided by NWDAF integration*
- **PCF applies or relaxes constraints dynamically**

Factoid: Without NWDAF, the PCF deducts used capacity from UDR via Session-AMBR and MBR during flow provisioning.

- **PCF deducts used capacity** from **UDR** via **Session-AMBR**
- **PCF deducts used capacity** from **UDR** via **MBR**

- **Deducting used capacity via Session-AMBR and MBR occurs during flow provisioning**

Factoid: If PCF rejects due to exceeded data rate, SMF returns HTTP 403 'EXCEEDED_SLICE_DATA_RATE' to UE.

- **PCF rejects flow setup due to exceeded data rate**
- **SMF returns HTTP 403 'EXCEEDED_SLICE_DATA_RATE' to UE**

Factoid: Slice capacity is restored to UDR when sessions or GBR rules terminate or are modified downward.

- **Slice capacity is restored to UDR**
- **Restoration of slice capacity to UDR occurs when sessions terminate**
- **Restoration of slice capacity to UDR occurs when GBR rules are modified downward**

Factoid: Multiple PCFs can synchronously enforce slice caps using conditional UDR updates with etags.

- **Multiple PCFs can synchronously enforce slice caps**
- **Slice cap enforcement uses conditional UDR updates with etags**

Factoid: Maximum Group Data Rate control is extended to VN groups similarly to slice-based rate limiting.

- **Maximum Group Data Rate control is extended to VN groups**
- **Extension of Maximum Group Data Rate control to VN groups is similar to slice-based rate limiting**

Factoid: Emergency or prioritized services can bypass slice data rate limits based on operator policy or regulation.

- **Emergency services can bypass slice data rate limits**
- **Prioritized services can bypass slice data rate limits**
- **Bypassing slice data rate limits is based on operator policy**
- **Bypassing slice data rate limits is based on regulation**

Factoid: Slice caps apply across both GBR and non-GBR flows—Non-GBR via Session-AMBR, GBR via MBR in PCC rules.

- **Slice caps apply across non-GBR flows**
- **Slice caps apply across GBR flows**
- **Non-GBR flows are enforced via Session-AMBR**
- **GBR flows are enforced via MBR in PCC rules**

Factoid: ENISA identifies core control-plane NFs (AUSF, SEAF, SEPP, SIDP, AMF, UDM, NSSAAF) and NFV/MANO components as critical assets requiring trust zone isolation.

- **ENISA identifies AUSF as critical asset requiring trust zone isolation**
- **ENISA identifies SEAF as critical asset requiring trust zone isolation**
- **ENISA identifies SEPP as critical asset requiring trust zone isolation**
- **ENISA identifies SIDP as critical asset requiring trust zone isolation**
- **ENISA identifies AMF as critical asset requiring trust zone isolation**
- **ENISA identifies UDM as critical asset requiring trust zone isolation**
- **ENISA identifies NSSAAF as critical asset requiring trust zone isolation**
- **ENISA identifies NFV/MANO components as critical assets requiring trust zone isolation**

Factoid: ENISA's threat landscape uses nine architectural 'zoom-ins' to map vulnerabilities across CP–UP, NFV, SDN, MEC, slicing, and orchestration domains.

- **ENISA's threat landscape uses nine architectural 'zoom-ins'**
- **Nine architectural 'zoom-ins' map vulnerabilities**
- **Mapping vulnerabilities occurs across control-plane and user-plane (CP–UP)**
- **Mapping vulnerabilities occurs across NFV domain**
- **Mapping vulnerabilities occurs across SDN domain**
- **Mapping vulnerabilities occurs across MEC domain**
- **Mapping vulnerabilities occurs across slicing domain**
- **Mapping vulnerabilities occurs across orchestration domain**

Factoid: EU regulations (NISD, Telecom Security Act, Article 13a) require strict operational controls: IAM, configuration management, logging, continuity planning.

- **EU regulations (NISD, Telecom Security Act, Article 13a) require IAM**
- **EU regulations (NISD, Telecom Security Act, Article 13a) require configuration management**
- **EU regulations (NISD, Telecom Security Act, Article 13a) require logging**
- **EU regulations (NISD, Telecom Security Act, Article 13a) require continuity planning**

Factoid: Control-plane functions are architecturally isolated from user-plane, mitigating CP-targeted attacks without affecting UP traffic.

- **Control-plane functions are architecturally isolated from user-plane**
- **Architectural isolation mitigates CP-targeted attacks**
- **Mitigating CP-targeted attacks occurs without affecting UP traffic**

_____ BATCH 7 GEMINI _____

Factoid: MANO, NFVO, and VNF/VIM orchestrators are critical CP assets vulnerable to supply chain and virtualization-layer attacks.

- **MANO is critical CP asset**
- **NFVO is critical CP asset**
- **VNF orchestrators are critical CP assets**

- **VIM orchestrators** *are critical CP assets*
- **Critical CP assets** *are vulnerable to supply chain attacks*
- **Critical CP assets** *are vulnerable to virtualization-layer attacks*

Factoid: Attacks on CP APIs and slice orchestration functions can compromise multiple domains—ENISA recommends hardened authentication and policy enforcement.

- **Attacks on CP APIs and slice orchestration functions** *can compromise multiple domains*
- **ENISA** *recommends hardened authentication*
- **ENISA** *recommends policy enforcement*

Factoid: Operational resilience requirements include incident reporting, network audits, disaster recovery, and service continuity under EU telecom law.

- **Operational resilience requirements** *include incident reporting*
- **Operational resilience requirements** *include network audits*
- **Operational resilience requirements** *include disaster recovery*
- **Operational resilience requirements** *include service continuity*
- **Operational resilience requirements** *are mandated under EU telecom law*

Factoid: NFV and SDN systems centralize both CP and UP control, making them high-value attack vectors for DoS or configuration tampering.

- **NFV systems** *centralize CP control*
- **NFV systems** *centralize UP control*
- **SDN systems** *centralize CP control*
- **SDN systems** *centralize UP control*
- **Centralization of CP and UP control** *makes NFV and SDN systems high-value attack vectors*
- **High-value attack vectors** *include DoS*
- **High-value attack vectors** *include configuration tampering*

Factoid: Trust boundaries are enforced using certificate-based isolation, role-based access, and dedicated API profiles per NF in control plane.

- **Trust boundaries** *are enforced using certificate-based isolation*
- **Trust boundaries** *are enforced using role-based access*
- **Trust boundaries** *are enforced using dedicated API profiles*
- **Dedicated API profiles** *apply per NF in control plane*

Factoid: Threat actors may exploit cross-domain trust misuse in orchestration and roaming as part of multi-stage CP attacks.

- **Threat actors** *may exploit cross-domain trust misuse in orchestration*
- **Threat actors** *may exploit cross-domain trust misuse in roaming*
- **Exploitation of cross-domain trust misuse** *is part of multi-stage control-plane attacks*

Factoid: UPF serves as the primary User-Plane NF in 5GC, handling packet forwarding, QoS, buffering, usage reporting, and mobility anchoring.

- **UPF serves as primary User-Plane NF in 5GC**
- **UPF handles packet forwarding**
- **UPF handles QoS**
- **UPF handles buffering**
- **UPF handles usage reporting**
- **UPF handles mobility anchoring**

Factoid: The UPF interconnects RAN and Data Networks, and performs reflective QoS marking and application-specific flow detection.

- **UPF interconnects RAN and Data Networks**
- **UPF performs reflective QoS marking**
- **UPF performs application-specific flow detection**

Factoid: UPF uses four reference points: N3 (gNB), N4 (SMF via PCF), N6 (Data Network), and N9 (UPF-to-UPF chaining).

- **UPF uses N3 reference point**
- **N3 reference point connects UPF and gNB**
- **UPF uses N4 reference point**
- **N4 reference point connects UPF and SMF via PCF**
- **UPF uses N6 reference point**
- **N6 reference point connects UPF and Data Network**
- **UPF uses N9 reference point**
- **N9 reference point supports UPF-to-UPF chaining**

Factoid: PCF over N4 installs forwarding (FAR), buffering (BAR), QoS (QER), usage (URR), and detection (PDR) rules in UPF.

- **PCF over N4 installs forwarding rules (FAR) in UPF**
- **PCF over N4 installs buffering rules (BAR) in UPF**
- **PCF over N4 installs QoS rules (QER) in UPF**
- **PCF over N4 installs usage rules (URR) in UPF**
- **PCF over N4 installs detection rules (PDR) in UPF**

Factoid: GTP-U encapsulated traffic traverses N3/N9 for user-plane data exchange between RAN and/or UPF instances.

- **GTP-U encapsulated traffic traverses N3 interface**
- **GTP-U encapsulated traffic traverses N9 interface**
- **Traversal of N3 and N9 facilitates user-plane data exchange**
- **User-plane data exchange occurs between RAN and UPF instances**

Factoid: UPF fulfills the CUPS model by decoupling data-plane processing from control-plane SMF logic and enabling edge deployment.

- **UPF fulfills Control and User Plane Separation (CUPS) model**
- **Control and User Plane Separation (CUPS) model decouples data-plane processing from control-plane SMF logic**
- **Decoupling of data-plane processing from control-plane SMF logic enables edge deployment**

Factoid: Cloud-native UPF architectures using microservices and Kubernetes enable elastic scaling for diverse throughput demands.

- **Cloud-native UPF architectures use microservices**
- **Cloud-native UPF architectures use Kubernetes**
- **Use of microservices and Kubernetes enables elastic scaling**
- **Elastic scaling supports diverse throughput demands**

Factoid: SmartNIC offload (Napatech) enables UPFs to process dual 100 Gbps data-plane loads with zero CPU usage.

- **SmartNIC offload (Napatech) enables UPFs to process dual 100 Gbps data-plane loads**
- **Processing dual 100 Gbps data-plane loads occurs with zero CPU usage**

Factoid: Branching via N9 or uplink classification allows UPF to intelligently split or redirect traffic flows.

- **Branching via N9 allows UPF to split or redirect traffic flows intelligently**
- **Uplink classification allows UPF to split or redirect traffic flows intelligently**

Factoid: UPF supports IPv4/IPv6, NAT at N6, multi-tenancy, session anchoring, and reflective QoS for end-to-end service compliance.

- **UPF supports IPv4/IPv6**
- **UPF supports NAT at N6**
- **UPF supports multi-tenancy**
- **UPF supports session anchoring**
- **UPF supports reflective QoS**
- **Reflective QoS enables end-to-end service compliance**

Factoid: UPF is the CUPS-based user-plane NF in 5GC, enabling separation from SMF and independent scaling/deployment.

- **UPF is the CUPS-based user-plane NF in 5GC**
- **UPF enables separation from SMF**
- **UPF enables independent scaling**
- **UPF enables independent deployment**

Factoid: CUPS allows UPF to evolve and scale independently, enabling edge deployment for low-latency use cases.

- **CUPS *allows* UPF to evolve independently**
- **CUPS *allows* UPF to scale independently**
- **Independent evolution and scaling of UPF *enables* edge deployment**
- **Edge deployment *supports* low-latency use cases**

Factoid: UPF handles mobility anchoring, IP allocation, routing, classification, branching, buffering, and PFD-based application detection.

- **UPF *handles* mobility anchoring**
- **UPF *handles* IP allocation**
- **UPF *handles* routing**
- **UPF *handles* classification**
- **UPF *handles* branching**
- **UPF *handles* buffering**
- **UPF *handles* PFD-based application detection**

Factoid: Packet-level policy enforcement in UPF includes gating, redirection, traffic steering, reflective QoS marking, and rate control.

- **Packet-level policy enforcement in UPF *includes* gating**
- **Packet-level policy enforcement in UPF *includes* redirection**
- **Packet-level policy enforcement in UPF *includes* traffic steering**
- **Packet-level policy enforcement in UPF *includes* reflective QoS marking**
- **Packet-level policy enforcement in UPF *includes* rate control**

Factoid: Scientific interfaces include N3 for GTP-U, N4 for PFCP rule installation, N6 for Data Network connectivity, and N9 for UPF-to-UPF chaining.

- **N3 interface *uses* GTP-U**
- **N4 interface *is used for* PFCP rule installation**
- **N6 interface *connects* Data Network**
- **N9 interface *enables* UPF-to-UPF chaining**

Factoid: PFCP enables UPF to install PDR, FAR, BAR, QER, and URR via N4 under SMF control.

- **PFCP *enables* UPF to install PDR via N4 under SMF control**
- **PFCP *enables* UPF to install FAR via N4 under SMF control**
- **PFCP *enables* UPF to install BAR via N4 under SMF control**
- **PFCP *enables* UPF to install QER via N4 under SMF control**
- **PFCP *enables* UPF to install URR via N4 under SMF control**

Factoid: N9 chaining supports branching, multi-anchoring, and breakout of traffic flows in complex deployment scenarios.

- **N9 chaining** *supports branching*
- **N9 chaining** *supports multi-anchoring*
- **N9 chaining** *supports breakout of traffic flows*
- **Branching, multi-anchoring, and breakout** *apply in complex deployment scenarios*

Factoid: Edge-deployed UPFs and cloud-native architecture support elastic scaling managed via orchestration platforms.

- **Edge-deployed UPFs** *support elastic scaling*
- **Cloud-native architecture** *supports elastic scaling*
- **Elastic scaling** *is managed via orchestration platforms*

Factoid: UPF supports GTP-U-level packet duplication and elimination for handover and redundancy.

- **UPF** *supports GTP-U-level packet duplication*
- **UPF** *supports GTP-U-level packet elimination*
- **Packet duplication and elimination** *are used for handover*
- **Packet duplication and elimination** *are used for redundancy*

Factoid: Intent-driven orchestration of UPF includes lifecycle automated deployment, test validation, and performance assurance using DevOps aligned tools.

- **Intent-driven orchestration of UPF** *includes automated lifecycle deployment*
- **Intent-driven orchestration of UPF** *includes test validation*
- **Intent-driven orchestration of UPF** *includes performance assurance*
- **Automated deployment, validation, and assurance** *use DevOps-aligned tools*

Factoid: GTP octet counters on N3 (GTP.In/OutDataOctetsN3) monitor transport bandwidth usage per QoS class.

- **GTP octet counters on N3** *include GTP.InDataOctetsN3 and GTP.OutDataOctetsN3*
- **GTP.In/OutDataOctetsN3** *monitor transport bandwidth usage*
- **Transport bandwidth usage** *is tracked per QoS class*

Factoid: Packet loss on N3 is tracked via `GTP.InDataPktLossN3QoS`, enabling detection of per-class service degradation.

- **Packet loss on N3** *is tracked via GTP.InDataPktLossN3QoS*
- **GTP.InDataPktLossN3QoS** *enables detection of per-class service degradation*

Factoid: Average and distribution metrics of N3 round-trip delay per DSCP (e.g., `GTP.RttDelayN3DIPsaUpfMean.DSCP`) support QoS-level latency modelling.

- **Average and distribution metrics of N3 round-trip delay per DSCP** *are measured using GTP.RttDelayN3DIPsaUpfMean.DSCP*

- **GTP.RttDelayN3DIPsaUpfMean.DSCP** *supports QoS-level latency modelling*

Factoid: N4 PFCP session metrics measure **N4SessionEstabReq/Fail** and **N4SessionReport/ReportSucc** to assess control-plane signaling performance.

- **N4 PFCP session metrics** *measure N4SessionEstabReq*
- **N4 PFCP session metrics** *measure N4SessionEstabFail*
- **N4 PFCP session metrics** *measure N4SessionReport*
- **N4 PFCP session metrics** *measure N4SessionReportSucc*
- **These metrics** *assess control-plane signaling performance*

Factoid: N6 interface metrics (**IP.N6IncLinkUsage**, **IP.N6OutLinkUsage**) record IP-layer data volume, consistent with RFC 5136 counters.

- **N6 interface metrics** *include IP.N6IncLinkUsage*
- **N6 interface metrics** *include IP.N6OutLinkUsage*
- **IP.N6IncLinkUsage** and **IP.N6OutLinkUsage** *record IP-layer data volume*
- **These metrics** *are consistent with RFC 5136 counters*

Factoid: Per-QoS-class and per-slice measurements enable slice-specific SLA validation across UPF interfaces.

- **Per-QoS-class measurements** *enable slice-specific SLA validation*
- **Per-slice measurements** *enable slice-specific SLA validation*
- **Slice-specific SLA validation** *occurs across UPF interfaces*

Factoid: N3, N4, and N6 measurement data provide key inputs to dynamic QoS and scaling actions in intent-driven orchestration.

- **N3 measurement data** *provides key inputs to dynamic QoS and scaling actions*
- **N4 measurement data** *provides key inputs to dynamic QoS and scaling actions*
- **N6 measurement data** *provides key inputs to dynamic QoS and scaling actions*
- **Dynamic QoS and scaling actions** *are used in intent-driven orchestration*

Factoid: Monitoring of GTP RTT and packet loss informs real-time traffic steering or re-orchestration decisions.

- **Monitoring of GTP RTT** *informs real-time traffic steering*
- **Monitoring of GTP RTT** *informs re-orchestration decisions*
- **Monitoring of packet loss** *informs real-time traffic steering*
- **Monitoring of packet loss** *informs re-orchestration decisions*

Factoid: Control-plane PFCP session metrics support monitoring of SMF-UPF coordination health.

- **Control-plane PFCP session metrics** *support monitoring of SMF-UPF coordination health*

Factoid: UESPLIT per-interface metrics feed operators' QoS dashboards and NWDAF analytics for performance assurance.

- **UESPLIT per-interface metrics** *feed operators' QoS dashboards*
- **UESPLIT per-interface metrics** *feed NWDAF analytics*
- **Operators' QoS dashboards and NWDAF analytics** *support performance assurance*

Factoid: UPCR (UPF placement and chaining reconfiguration) addresses UPF placement in MEC environments to handle user mobility.

- **UPCR (UPF placement and chaining reconfiguration)** *addresses UPF placement in MEC environments*
- **UPCR** *handles user mobility through UPF placement and chaining*

Factoid: An ILP-based model optimizes UPF placement to minimize deployment and migration costs under QoS constraints.

- **ILP-based model** *optimizes UPF placement*
- **Optimization of UPF placement** *minimizes deployment costs*
- **Optimization of UPF placement** *minimizes migration costs*
- **Optimization of UPF placement** *operates under QoS constraints*

Factoid: The DPC-UPCR heuristic achieves within 15% of optimal placement outcomes with significantly reduced computation time.

- **DPC-UPCR heuristic** *achieves within 15% of optimal placement outcomes*
- **DPC-UPCR heuristic** *provides significantly reduced computation time*

Factoid: An optimal stopping theory–based scheduler triggers UPF reconfiguration when latency thresholds are breached.

- **Optimal stopping theory–based scheduler** *triggers UPF reconfiguration*
- **UPF reconfiguration** *is triggered when latency thresholds are breached*

Factoid: DPC-UPCR outperforms baseline schedulers by reducing reconfig event count and QoS violations.

- **DPC-UPCR** *outperforms baseline schedulers*
- **DPC-UPCR** *reduces reconfiguration event count*
- **DPC-UPCR** *reduces QoS violations*

Factoid: UPF repositioning supports proximity anchoring to minimize latency during access node handovers.

- **UPF repositioning** *supports proximity anchoring*
- **Proximity anchoring** *minimizes latency during access node handovers*

Factoid: Cost-QoS trade-offs are balanced via dynamic UPF chaining instead of static placement clusters.

- **Cost-QoS trade-offs** *are balanced via dynamic UPF chaining*
- **Dynamic UPF chaining** *is used instead of static placement clusters*

Factoid: User mobility patterns drive dynamic reconfiguration decisions rather than purely periodic adjustments.

- **User mobility patterns** *drive dynamic reconfiguration decisions*
- **Dynamic reconfiguration decisions** *are preferred over purely periodic adjustments*

Factoid: Scheduling parameters—e.g., latency violation rate and QoS thresholds—govern when to deploy UPF migrations.

- **Scheduling parameters** *include latency violation rate*
- **Scheduling parameters** *include QoS thresholds*
- **Latency violation rate and QoS thresholds** *govern when to deploy UPF migrations*

Factoid: The heuristic scheduler aims to minimize cost of migration while preserving slice-level performance during reconfigurations.

1. **Heuristic scheduler** *aims to minimize cost of migration*
2. **Heuristic scheduler** *preserves slice-level performance*
3. **Slice-level performance** *is preserved during reconfigurations*

Factoid: The 5G-PPP architecture integrates multi-domain orchestration through ETSI NFV-MANO and SDN agents across RAN, transport, core, and edge.

- **5G-PPP architecture** *integrates multi-domain orchestration*
- **Multi-domain orchestration** *is enabled through ETSI NFV-MANO*
- **Multi-domain orchestration** *is enabled through SDN agents*
- **SDN agents** *operate across RAN, transport, core, and edge domains*

Factoid: RAN functions are split: DU/PHY/MAC run at edge for performance, while PDCP/RRC run centrally as VNFs.

- **RAN functions** *are split*
- **DU/PHY/MAC** *run at edge for performance**
- **PDCP/RRC** *run centrally as VNFs*

Factoid: Core network topology includes NFVI in core and edge, fronthaul/backhaul, and WLAN/small-cell components.

- **Core network topology** *includes NFVI in core*
- **Core network topology** *includes NFVI in edge*
- **Core network topology** *includes fronthaul*
- **Core network topology** *includes backhaul*
- **Core network topology** *includes WLAN components*
- **Core network topology** *includes small-cell components*

Factoid: CUPS architecture enables independent scaling by separating CU-CP from CU-UP and using SMF–UPF control-plane splits.

- **CUPS architecture** *enables independent scaling*
- **CUPS architecture** *separates CU-CP from CU-UP*
- **CUPS architecture** *uses SMF–UPF control-plane splits*

Factoid: Edge deployment of DU, CU-UP, and UPF ensures low-latency, bandwidth-intensive data handling close to users.

- **Edge deployment of DU, CU-UP, and UPF** *ensures low-latency data handling*
- **Edge deployment of DU, CU-UP, and UPF** *ensures bandwidth-intensive data handling*
- **Low-latency and bandwidth-intensive data handling** *occurs close to users*

Factoid: SDN agents co-locate with domain NFs and are managed via NFVO/VNFM, enabling programmable domain control.

- **SDN agents** *co-locate with domain NFs*
- **SDN agents** *are managed via NFVO/VNFM*
- **Management via NFVO/VNFM** *enables programmable domain control*

Factoid: Intent ingestion, telemetry, and AI/ML closed-loop orchestration are best-practice control methods for slice and service SLAs.

- **Intent ingestion** *is a best-practice control method for slice and service SLAs*
- **Telemetry** *is a best-practice control method for slice and service SLAs*
- **AI/ML closed-loop orchestration** *is a best-practice control method for slice and service SLAs*

Factoid: PHY and MAC layers benefit from hardware acceleration at the DU, while PDCP/RRC benefit from centralized programmability.

- **PHY and MAC layers** *benefit from hardware acceleration at the DU*
- **PDCP/RRC layers** *benefit from centralized programmability*

Factoid: Transport-level programmability leveraged via T-API and intent-based NBIs integrated with ONF/SDN frameworks.

- **Transport-level programmability** *is leveraged via T-API*

- **Transport-level programmability** *is leveraged via intent-based NBIs*
- **T-API and intent-based NBIs** *are integrated with ONF/SDN frameworks*

Factoid: Application-aware orchestration includes plug-in service and function managers integrated into MANO for vertical-specific services.

- **Application-aware orchestration** *includes plug-in service managers*
- **Application-aware orchestration** *includes plug-in function managers*
- **Plug-in service and function managers** *are integrated into MANO*
- **Application-aware orchestration** *supports vertical-specific services*

Factoid: Kubernetes is used for edge orchestration to support sub-20 ms latency use cases in IoT and real-time applications.

- **Kubernetes** *is used for edge orchestration*
- **Edge orchestration with Kubernetes** *supports sub-20 ms latency use cases*
- **Sub-20 ms latency use cases** *include IoT applications*
- **Sub-20 ms latency use cases** *include real-time applications*

Factoid: Native Kubernetes lacks network-aware scheduling and topology-based placement essential for edge scenarios.

- **Native Kubernetes** *lacks network-aware scheduling*
- **Native Kubernetes** *lacks topology-based placement*
- **Network-aware scheduling and topology-based placement** *are essential for edge scenarios*

Factoid: Custom edge Kubernetes deployments use virtual kubelets, multi-cluster federation, and custom schedulers.

- **Custom edge Kubernetes deployments** *use virtual kubelets*
- **Custom edge Kubernetes deployments** *use multi-cluster federation*
- **Custom edge Kubernetes deployments** *use custom schedulers*

Factoid: Edge K8s implementations struggle with real-time metric processing, fault-tolerance, and container registry placement.

- **Edge Kubernetes implementations** *struggle with real-time metric processing*
- **Edge Kubernetes implementations** *struggle with fault-tolerance*
- **Edge Kubernetes implementations** *struggle with container registry placement*

_____ BATCH 8 GEMINI _____

Factoid: Edge-focused improvements include installing local metrics servers, topology-aware custom schedulers, and edge-located container registries.

- **Edge-focused improvements** *include installing local metrics servers*
- **Edge-focused improvements** *include topology-aware custom schedulers*

- **Edge-focused improvements** *include* **edge-located container registries**

Factoid: Edge orchestration requires orchestration decisions based on realtime network measurements and locality knowledge.

- **Edge orchestration** *requires* **orchestration decisions based on real-time network measurements**
- **Edge orchestration** *requires* **orchestration decisions based on locality knowledge**

Factoid: Fault tolerance at the edge demands resilient control planes and backup mechanisms for remote node failures.

- **Fault tolerance at the edge** *demands* **resilient control planes**
- **Fault tolerance at the edge** *demands* **backup mechanisms for remote node failures**

Factoid: Placing container registries on edge nodes reduces network usage and speeds up workload deployment.

- **Placing container registries on edge nodes** *reduces* **network usage**
- **Placing container registries on edge nodes** *speeds up* **workload deployment**

Factoid: Virtual kubelets enable edge nodes to publish as Kubernetes nodes without running full control plane components.

- **Virtual kubelets** *enable* **edge nodes to publish as Kubernetes nodes**
- **Virtual kubelets** *operate without* **running full control plane components**

Factoid: Topology-aware scheduling uses knowledge of node network connectivity to prioritize workload placement.

- **Topology-aware scheduling** *uses* **knowledge of node network connectivity**
- **Knowledge of node network connectivity** *is used to* **prioritize workload placement**

Factoid: Container placement spans cloud, fog, and edge domains, optimizing resource use, energy consumption, and fault tolerance.

- **Container placement** *spans* **cloud domain**
- **Container placement** *spans* **fog domain**
- **Container placement** *spans* **edge domain**
- **Container placement** *optimizes* **resource use**
- **Container placement** *optimizes* **energy consumption**
- **Container placement** *optimizes* **fault tolerance**

Factoid: Placement algorithms include optimization-based (e.g. LP, bin-packing), meta-heuristics (ACO, genetic), and reinforcement learning (MDP).

- Placement algorithms *include* optimization-based methods
- Optimization-based methods *include* Linear Programming (LP) and bin-packing
- Placement algorithms *include* meta-heuristics
- Meta-heuristics *include* Ant Colony Optimization (ACO) and genetic algorithms
- Placement algorithms *include* reinforcement learning methods
- Reinforcement learning methods *include* Markov Decision Process (MDP)

Factoid: Migration techniques include cold, pre-copy, post-copy, and hybrid methods, each with different downtime/resource overhead profiles.

- Migration techniques *include* cold migration
- Migration techniques *include* pre-copy migration
- Migration techniques *include* post-copy migration
- Migration techniques *include* hybrid migration
- Each migration method *has different* downtime and resource overhead profiles

Factoid: Edge/fog placements consider latency, energy, and resource constraints in trade-off-aware optimization.

- Edge/fog placements *consider* latency constraints
- Edge/fog placements *consider* energy constraints
- Edge/fog placements *consider* resource constraints
- Latency, energy, and resource constraints *are balanced in* trade-off-aware optimization

Factoid: Pre-copy migration reduces downtime via iterative state replication; post-copy focuses on rapid restart with on-demand state fetch.

- Pre-copy migration *reduces downtime via* iterative state replication
- Post-copy migration *focuses on rapid restart with* on-demand state fetch

Factoid: Meta-heuristic strategies like ACO and Grey Wolf optimize container placement for load balancing and energy efficiency.

- Meta-heuristic strategies *include* Ant Colony Optimization (ACO)
- Meta-heuristic strategies *include* Grey Wolf Optimization
- ACO and Grey Wolf Optimization *optimize* container placement
- Container placement optimization *targets* load balancing
- Container placement optimization *targets* energy efficiency

Factoid: Fault tolerance in edge setups is supported via proactive migrations and container replication to avoid single-node failures.

- Fault tolerance in edge setups *is supported via* proactive migrations
- Fault tolerance in edge setups *is supported via* container replication
- Proactive migrations and container replication *help avoid* single-node failures

Factoid: Energy-aware placement utilizes bin-packing to minimize the number of active fog nodes, saving power.

- **Energy-aware placement** *utilizes bin-packing*
- **Bin-packing** *minimizes number of active fog nodes*
- **Minimizing active fog nodes** *saves power*

Factoid: Placement models often include QoS constraints, formulated using MILP or knapsack to enforce latency and reliability SLAs.

- **Placement models** *often include QoS constraints*
- **QoS constraints** *are formulated using Mixed Integer Linear Programming (MILP)*
- **QoS constraints** *are formulated using knapsack problem models*
- **Formulated constraints** *enforce latency SLAs*
- **Formulated constraints** *enforce reliability SLAs*

Factoid: Reinforcement learning enables adaptive container placement sensitive to real-time metrics and evolving constraints.

- **Reinforcement learning** *enables adaptive container placement*
- **Adaptive container placement** *is sensitive to real-time metrics*
- **Adaptive container placement** *responds to evolving constraints*

Factoid: VNF and container placement across cloud, fog, and edge are NP-hard and directly impact latency, cost, energy, and scalability.

- **VNF and container placement** *across cloud, fog, and edge are NP-hard*
- **VNF and container placement** *directly impact latency*
- **VNF and container placement** *directly impact cost*
- **VNF and container placement** *directly impact energy*
- **VNF and container placement** *directly impact scalability*

Factoid: Placement techniques are classified as optimization-based (e.g., LP, MIQCP), meta-heuristics (ACO, GWO), or ML-based (reinforcement learning, deep RL).

- **Placement techniques** *are classified as optimization-based*
- **Optimization-based techniques** *include Linear Programming (LP) and Mixed-Integer Quadratically Constrained Programming (MIQCP)*
- **Placement techniques** *are classified as meta-heuristics*
- **Meta-heuristics** *include Ant Colony Optimization (ACO) and Grey Wolf Optimization (GWO)*
- **Placement techniques** *are classified as ML-based*
- **ML-based techniques** *include reinforcement learning and deep reinforcement learning (deep RL)*

Factoid: Edge placements prioritize low latency, energy efficiency, resource constraints, and intermittent connectivity.

- **Edge placements** *prioritize low latency*
- **Edge placements** *prioritize energy efficiency*
- **Edge placements** *prioritize resource constraints*
- **Edge placements** *prioritize intermittent connectivity*

Factoid: Multi-objective placement models balance latency, energy, cost, resource use, fault tolerance, and QoS among VNFs/containers.

- **Multi-objective placement models** *balance latency*
- **Multi-objective placement models** *balance energy*
- **Multi-objective placement models** *balance cost*
- **Multi-objective placement models** *balance resource use*
- **Multi-objective placement models** *balance fault tolerance*
- **Multi-objective placement models** *balance QoS*
- **Balanced objectives** *apply to VNFs and containers*

Factoid: Reinforcement learning enables dynamic adaptation of placement strategies in response to traffic and resource fluctuations.

- **Reinforcement learning** *enables dynamic adaptation of placement strategies*
- **Dynamic adaptation of placement strategies** *responds to traffic fluctuations*
- **Dynamic adaptation of placement strategies** *responds to resource fluctuations*

Factoid: Meta-heuristic algorithms like ant-colony and grey wolf are widely used for balancing load, energy, and QoS objectives.

- **Meta-heuristic algorithms** *include ant-colony optimization*
- **Meta-heuristic algorithms** *include grey wolf optimization*
- **Ant-colony and grey wolf algorithms** *are widely used for balancing load*
- **Ant-colony and grey wolf algorithms** *are widely used for balancing energy*
- **Ant-colony and grey wolf algorithms** *are widely used for balancing QoS objectives*

Factoid: Optimization models like MIQCP or bin-packing address latency and energy constraints, especially in uRLLC scenarios.

- **Optimization models** *include MIQCP (Mixed-Integer Quadratically Constrained Programming)*
- **Optimization models** *include bin-packing*
- **MIQCP and bin-packing** *address latency constraints*
- **MIQCP and bin-packing** *address energy constraints*
- **Latency and energy constraints** *are especially relevant in uRLLC scenarios*

Factoid: Dynamic placement systems migrate VNFs/containers proactively to maintain SLAs in changing network states.

- **Dynamic placement systems** *migrate VNFs/containers proactively*
- **Proactive migration of VNFs/containers** *maintains SLAs*

- **Proactive migration** *responds to changing network states*

Factoid: Container placement extends beyond VNFs to include microservices in MEC and private 5G environments.

- **Container placement** *extends beyond VNFs*
- **Container placement** *includes microservices*
- **Microservices** *are placed in MEC environments*
- **Microservices** *are placed in private 5G environments*

Factoid: Taxonomy connects placement techniques to objectives—single vs multi-objective, static vs dynamic, method class used.

- **Taxonomy** *connects placement techniques to objectives*
- **Objectives** *include single-objective and multi-objective*
- **Placement techniques** *are categorized as static or dynamic*
- **Placement techniques** *are grouped by method class used* (e.g., optimization-based, meta-heuristic, ML-based)

Factoid: C-RAN base station functions (BBU/RRH) can be containerized using Docker and orchestrated via Kubernetes.

- **C-RAN base station functions** *include BBU (Baseband Unit) and RRH (Remote Radio Head)*
- **BBU/RRH functions** *can be containerized using Docker*
- **Containerized BBU/RRH functions** *can be orchestrated via Kubernetes*

Factoid: Kubernetes StatefulSets provide pod identity and ordering guarantees for RAN component scaling.

- **Kubernetes StatefulSets** *provide pod identity guarantees*
- **Kubernetes StatefulSets** *provide pod ordering guarantees*
- **Pod identity and ordering guarantees** *support RAN component scaling*

Factoid: Calico CNI enables layer-3 networking among RAN containers in orchestration deployments.

- **Calico CNI** *enables layer-3 networking*
- **Layer-3 networking** *operates among RAN containers*
- **Calico CNI** *is used in orchestration deployments*

Factoid: etcd acts as a distributed key-value store for dynamic runtime configuration between containerized RRH and BBU modules.

- **etcd** *acts as a distributed key-value store*
- **etcd** *supports dynamic runtime configuration*
- **Dynamic runtime configuration** *occurs between containerized RRH and BBU modules*

Factoid: Doubling the number of BBU–RRH pods linearly increases fronthaul data throughput in the containerized RAN testbed.

- **Doubling the number of BBU–RRH pods** *linearly increases fronthaul data throughput*
- **Linear increase in fronthaul data throughput** *is observed in containerized RAN testbed*

Factoid: OAISIM-based RRH emulation consumes over 60% CPU, which is useful for creating autoscaling policies.

- **OAISIM-based RRH emulation** *consumes over 60% CPU*
- **High CPU consumption** *is useful for creating autoscaling policies*

Factoid: Container orchestration enables dynamic resource-driven scaling of RAN baseband functions for performance and efficiency.

- **Container orchestration** *enables dynamic resource-driven scaling*
- **Dynamic scaling** *applies to RAN baseband functions*
- **Dynamic scaling of RAN baseband functions** *improves performance*
- **Dynamic scaling of RAN baseband functions** *improves efficiency*

Factoid: Resource metrics (CPU, memory) from Kubernetes can trigger scaling based on observed load patterns.

- **Resource metrics (CPU, memory)** *from Kubernetes*
- **Kubernetes resource metrics** *can trigger scaling actions*
- **Scaling actions** *are based on observed load patterns*

Factoid: Container-based virtualization (Docker/LXC) in real-time industrial systems incurs latency jitter ($\approx 37\text{--}102$ ms), posing challenges for hard real-time constraints

- **Container-based virtualization (Docker/LXC)** *in real-time industrial systems incurs latency jitter ($\approx 37\text{--}102$ ms)*
- **Latency jitter** *poses challenges for hard real-time constraints*

Factoid: Soft real-time deadlines (50–100 ms) are feasible with containers when the host is RT-kernel tuned and orchestration supports deterministic latency.

- **Soft real-time deadlines (50–100 ms)** *are feasible with containers*
- **Feasibility of soft real-time deadlines** *requires RT-kernel tuning on the host*
- **Feasibility of soft real-time deadlines** *requires orchestration support for deterministic latency*

Factoid: Lack of built-in real-time scheduling in container runtimes requires use of RT-patched kernels or RT-specific orchestration frameworks.

- **Lack of built-in real-time scheduling in container runtimes** *requires use of RT-patched kernels*
- **Lack of built-in real-time scheduling in container runtimes** *requires RT-specific orchestration frameworks*

Factoid: Experimental task latency varies by virtualization type: 37–102 ms for virtual PLCs; 47–54 ms for kernel-level containers.

- **Experimental task latency** *varies by virtualization type*
- **Virtual PLCs** *experience latency of 37–102 ms*
- **Kernel-level containers** *experience latency of 47–54 ms*

Factoid: Current orchestration platforms (e.g., Kubernetes) lack real-time awareness—missing scheduling priorities and time-bound placement logic.

- **Current orchestration platforms (e.g., Kubernetes)** *lack real-time awareness*
- **Lack of real-time awareness** *includes missing scheduling priorities*
- **Lack of real-time awareness** *includes missing time-bound placement logic*

Factoid: Industrial-grade container platforms need integration of real-time metrics in orchestration controllers to meet strict latency SLAs.

- **Industrial-grade container platforms** *need integration of real-time metrics in orchestration controllers*
- **Integration of real-time metrics** *is required to meet strict latency SLAs*

Factoid: Real-time container maturity is still limited—jitter and unpredictable scheduling up to 100 ms impede hard real-time use.

- **Real-time container maturity** *is still limited*
- **Jitter and unpredictable scheduling up to 100 ms**
- **Jitter and scheduling unpredictability** *impede hard real-time use*

Factoid: Future solution directions include real-time scheduling extensions, metrics-driven orchestration, and explicit task-level prioritization.

- **Future solution directions** *include real-time scheduling extensions*
- **Future solution directions** *include metrics-driven orchestration*
- **Future solution directions** *include explicit task-level prioritization*

Factoid: Time-sensitive container requirements demand enhancements at both runtime (cgroup, RT kernel) and orchestration (placement, scaling) layers.

- **Time-sensitive container requirements** *demand enhancements at runtime layer*
- **Runtime layer enhancements** *include cgroup and RT kernel*
- **Time-sensitive container requirements** *demand enhancements at orchestration layer*
- **Orchestration layer enhancements** *include placement and scaling*

Factoid: Systematic surveys identify container overhead as a key bottleneck for industrial control loops, guiding research into latency-aware virtualization.

- **Systematic surveys** *identify container overhead* as a **key bottleneck for industrial control loops**
- **Identification of container overhead** *guides research into latency-aware virtualization*

Factoid: Control-plane (N2) and user-plane (N3) traffic require separate IP subnets, unless using the same VLAN.

- **Control-plane (N2) traffic** *requires a separate IP subnet*
- **User-plane (N3) traffic** *requires a separate IP subnet*
- **Separate IP subnets** *are not required if N2 and N3 use the same VLAN*

Factoid: Edge deployments use dedicated subnets or UE IP pools, allocated via N6 interfaces for dynamic and static addresses.

- **Edge deployments** *use dedicated subnets*
- **Edge deployments** *use UE IP pools*
- **Dedicated subnets and UE IP pools** *are allocated via N6 interfaces*
- **N6 interfaces** *support allocation of dynamic addresses*
- **N6 interfaces** *support allocation of static addresses*

Factoid: NAPT must be disabled on N6 when external servers need to initiate connections to UEs.

- **NAPT** *must be disabled on N6*
- **Disabling NAPT on N6** *is required when external servers need to initiate connections to UEs*

Factoid: Firewalls must allow routes and port access between corporate networks and UE subnets connected via N6.

- **Firewalls** *must allow routes between corporate networks and UE subnets*
- **Firewalls** *must allow port access between corporate networks and UE subnets*
- **UE subnets** *are connected via N6 interface*

Factoid: HA setups use active-standby ASE pairs with VRRP and BFD to maintain control- and user-plane continuity within 2.5s failover.

- **HA setups** *use active-standby ASE pairs*
- **Active-standby ASE pairs** *utilize VRRP and BFD*
- **VRRP and BFD** *maintain control- and user-plane continuity*
- **Failover continuity** *is maintained within 2.5 seconds*

Factoid: Gateway routers require single static routes per network, with virtual IPs managed across redundant ASE devices.

- **Gateway routers** *require single static routes per network*
- **Virtual IPs** *are managed across redundant ASE devices*

Factoid: Separate IP pools are needed for access, data, management, cluster nodes, ACS/NFS, and virtual IP services in ASE deployments.

- **Separate IP pools** *are needed for access traffic*
- **Separate IP pools** *are needed for data traffic*
- **Separate IP pools** *are needed for management interfaces*
- **Separate IP pools** *are needed for cluster nodes*
- **Separate IP pools** *are needed for ACS/NFS services*
- **Separate IP pools** *are needed for virtual IP services*
- **These IP pools** *are required in ASE deployments*

Factoid: Azure reserves five private IPs in each subnet, limiting available resources for NF instances.

- **Azure** *reserves five private IPs in each subnet*
- **Reserved private IPs** *limit available resources for NF instances*

Factoid: UE IP address pool is specified in CIDR block and passed to AP5GC during site deployment via ARM template.

- **UE IP address pool** *is specified in CIDR block*
- **CIDR block** *is passed to AP5GC*
- **CIDR block** *is passed during site deployment*
- **Site deployment** *uses ARM template*

Factoid: Proper segmentation of N2/N3/N6 subnets supports network isolation, QoS, and firewall rule enforcement.

- **Proper segmentation of N2/N3/N6 subnets** *supports network isolation*
- **Proper segmentation of N2/N3/N6 subnets** *supports QoS enforcement*
- **Proper segmentation of N2/N3/N6 subnets** *supports firewall rule enforcement*

Factoid: IPv6 is preferred in private 5G deployments for its address space, built-in IPsec, and autoconfiguration (SLAAC).

- **IPv6** *is preferred in private 5G deployments*
- **Preference for IPv6** *is due to larger address space*
- **Preference for IPv6** *is due to built-in IPsec support*
- **Preference for IPv6** *is due to autoconfiguration (SLAAC)*

Factoid: IPv4 address pools are limited and require NAT, which introduces latency and complexity, especially in URLLC contexts.

- **IPv4 address pools** *are limited*
- **Limited IPv4 pools** *require NAT*
- **NAT** *introduces latency*
- **NAT** *introduces complexity*
- **Latency and complexity from NAT** *are problematic in URLLC contexts*

Factoid: InterLIR recommends a dual-stack approach during migration from IPv4 to IPv6 in private 5G networks.

- **InterLIR** *recommends a dual-stack approach*
- **Dual-stack approach** *is used during migration from IPv4 to IPv6*
- **Migration from IPv4 to IPv6** *applies to private 5G networks*

Factoid: RFC 1918 private IPv4 ranges are used to isolate network slices and device classes in private 5G.

- **RFC 1918 private IPv4 ranges** *are used to isolate network slices*
- **RFC 1918 private IPv4 ranges** *are used to isolate device classes*
- **Isolation of slices and devices** *occurs in private 5G networks*

Factoid: DHCP is used for dynamic IP assignment; static pools may serve fixed-function devices.

- **DHCP** *is used for dynamic IP assignment*
- **Static IP pools** *may serve fixed-function devices*

Factoid: NAT segmentation across slices enhances isolation and security in multi-slice deployments.

- **NAT segmentation across slices** *enhances isolation*
- **NAT segmentation across slices** *enhances security*
- **Isolation and security improvements** *apply to multi-slice deployments*

Factoid: Automated IPAM systems with monitoring capabilities are critical for conflict detection and utilization visibility.

- **Automated IPAM systems** *with monitoring capabilities are critical for conflict detection*
- **Automated IPAM systems** *are critical for utilization visibility*

_____ BATCH 9 GEMINI _____

Factoid: Real-time alerts from IPAM aid in rogue device detection and prevent IP conflicts in private 5G environments.

- **Real-time alerts from IPAM** *aid in rogue device detection*
- **Real-time alerts from IPAM** *prevent IP conflicts*
- **Rogue device detection and IP conflict prevention** *are essential in private 5G environments*

Factoid: IPv6 eliminates NAT overhead, improving latency performance in low-latency applications.

- **IPv6** *eliminates NAT overhead*
- **Elimination of NAT overhead** *improves latency performance*
- **Improved latency performance** *benefits low-latency applications*

Factoid: Dual-stack deployment supports gradual IPv6 adoption while maintaining IPv4 compatibility with legacy devices.

- **Dual-stack deployment** *supports gradual IPv6 adoption*
- **Dual-stack deployment** *maintains IPv4 compatibility with legacy devices*

Factoid: In CUPS, SMF dynamically assigns IP chunks to UPFs; UPFs allocate individual IPs per UE and report usage back to SMF.

- **In CUPS, SMF** *dynamically assigns IP chunks to UPFs*
- **UPFs** *allocate individual IPs per UE*
- **UPFs** *report IP usage back to SMF*

Factoid: Cisco's Cloud-Native IPAM includes a central IPAM Server and distributed IPAM Caches for multi-cluster consistency.

- **Cisco's Cloud-Native IPAM** *includes a central IPAM Server*
- **Cisco's Cloud-Native IPAM** *includes distributed IPAM Caches*
- **Central IPAM Server and distributed IPAM Caches** *ensure multi-cluster consistency*

Factoid: SMF triggers new IP chunk allocations when UPF usage exceeds 70%, and reclaims underutilized chunks.

- **SMF** *triggers new IP chunk allocations* when **UPF usage exceeds 70%**
- **SMF** *reclaims underutilized IP chunks*

Factoid: Optimal chunk size balances even load distribution against IP exhaustion and provisioning overhead.

- **Optimal chunk size** *balances even load distribution*
- **Optimal chunk size** *balances IP exhaustion*
- **Optimal chunk size** *balances provisioning overhead*

Factoid: Maximum UPFs per chunk-limited pool are determined by pool size divided by chunk size (e.g., 1M/65k → ~16 UPFs).

- **Maximum UPFs per chunk-limited pool** *are determined by pool size divided by chunk size*
- **Example calculation:** 1M IPs / 65k per chunk → ~16 UPFs

Factoid: UP Groups tied to APN define chunk associations; dynamic pool updates work without Sx reassociation.

- **UP Groups tied to APN** *define chunk associations*
- **Dynamic pool updates** *operate without requiring Sx reassociation*

Factoid: DNS-based UPF selection uses TAC/RAC to bind UPF and chunk assignment in geo-aware deployments.

- **DNS-based UPF selection** *uses TAC (Tracking Area Code) and RAC (Routing Area Code)*
- **TAC/RAC** *bind UPF and chunk assignment*
- **Binding UPF and chunk assignment** *enables geo-aware deployments*

Factoid: SMF enforces chunk throttling based on UPF's advertised max session capacity to prevent over-allocation.

- **SMF enforces chunk throttling**
- **Chunk throttling** *is based on UPF's advertised max session capacity*
- **Throttling** *prevents over-allocation*

Factoid: Static IP pools are split and distributed to multiple UPFs; SMF rejects requests outside defined static blocks.

- **Static IP pools** *are split and distributed to multiple UPFs*
- **SMF rejects requests outside defined static blocks**

Factoid: DHCP-based IP allocation via UPF requires PFCP N4 NF integration with VLAN ID tagging and DHCP DORA exchange.

- **DHCP-based IP allocation via UPF** *requires PFCP N4 NF integration*
- **DHCP-based allocation** *requires VLAN ID tagging*
- **DHCP-based allocation** *requires DHCP Dora exchange*

Factoid: Cisco best practice suggests hierarchical subnet design with Access, Distribution, and Core layers to improve scalability and resilience.

- **Cisco best practice** *suggests hierarchical subnet design*
- **Hierarchical subnet design** *includes Access layer*
- **Hierarchical subnet design** *includes Distribution layer*
- **Hierarchical subnet design** *includes Core layer*
- **Access, Distribution, and Core layers** *improve scalability and resilience*

Factoid: Subnetting divides an IP space into network and host portions, ensuring uniqueness and route summarization.

- **Subnetting** *divides an IP space into network and host portions*
- **Subnetting** *ensures address uniqueness*
- **Subnetting** *enables route summarization*

Factoid: VLSM allows subnets of varying sizes within a block—key for allocating optimal address ranges per NF/service.

- **VLSM (Variable Length Subnet Masking)** *allows subnets of varying sizes within a block*
- **VLSM** *is key for allocating optimal address ranges per NF/service*

Factoid: IPv6's 16-bit Subnet ID supports up to 65,535 subnets under a /64 prefix, enabling zonal network segmentation.

- **IPv6's 16-bit Subnet ID** *supports up to 65,535 subnets under a /64 prefix*
- **65,535 subnets under /64** *enable zonal network segmentation*

Factoid: Cisco recommends using modified EUI-64 interface IDs for stable, unique host addressing in IPv6 subnets.

- **Cisco** *recommends using modified EUI-64 interface IDs*
- **Modified EUI-64 interface IDs** *provide stable, unique host addressing*
- **Stable, unique host addressing** *is applied in IPv6 subnets*

Factoid: Hierarchical design localizes network changes, improves fault isolation, and reduces routing table size.

- **Hierarchical design** *localizes network changes*
- **Hierarchical design** *improves fault isolation*
- **Hierarchical design** *reduces routing table size*

Factoid: Typical subnet sizes include /28 for NF instances and /24 for core uplinks, providing headroom for growth.

- **Typical subnet size /28** *is used for NF instances*
- **Typical subnet size /24** *is used for core uplinks*
- **/28 and /24 subnet sizing** *provides headroom for growth*

Factoid: Hierarchical and CIDR-based subnetting enhances scalability and ease of management in telecom infrastructures.

- **Hierarchical subnetting** *enhances scalability*
- **Hierarchical subnetting** *enhances ease of management*
- **CIDR-based subnetting** *enhances scalability*
- **CIDR-based subnetting** *enhances ease of management*
- **Scalability and manageability improvements** *apply to telecom infrastructures*

Factoid: Separated subnets per network plane (control, user, management) aid in traffic isolation and QoS policy enforcement.

- **Separated subnets per network plane (control, user, management) aid in traffic isolation**
- **Separated subnets per network plane (control, user, management) aid in QoS policy enforcement**

Factoid: Efficient address planning prevents waste from fixed-length subnetting and traps address exhaustion before it occurs.

- **Efficient address planning prevents waste from fixed-length subnetting**
- **Efficient address planning prevents address exhaustion before it occurs**

Factoid: ZTE classifies UPFs into Central, Regional, Edge, and Campus deployments, each with defined throughput and latency metrics.

- **ZTE classifies UPFs into Central, Regional, Edge, and Campus deployments**
- **Each UPF deployment type has defined throughput metrics**
- **Each UPF deployment type has defined latency metrics**

Factoid: Campus UPFs offer ~50 Gbps throughput with sub-15 ms latency and include TSN, URLLC, and 5G-LAN enhancements.

- **Campus UPFs offer ~50 Gbps throughput**
- **Campus UPFs deliver sub-15 ms latency**
- **Campus UPFs include TSN enhancements**
- **Campus UPFs include URLLC enhancements**
- **Campus UPFs include 5G-LAN enhancements**

Factoid: Subnet planning for UPFs defines separate IP ranges per interface: N3, N4, N6, ensuring flow isolation and QoS.

- **Subnet planning for UPFs defines separate IP ranges per interface**
- **Separate IP ranges are assigned for N3 interface**
- **Separate IP ranges are assigned for N4 interface**
- **Separate IP ranges are assigned for N6 interface**
- **Separate IP ranges per interface ensure flow isolation**
- **Separate IP ranges per interface ensure QoS enforcement**

Factoid: VLAN or subnet segmentation at edge/campus UPFs enables localized firewall and NAT policies.

- **VLAN or subnet segmentation at edge/campus UPFs enables localized firewall policies**
- **VLAN or subnet segmentation at edge/campus UPFs enables localized NAT policies**

Factoid: Local firewall/NAT in campus UPFs ensures enterprise-grade security and traffic control at the site level.

- **Local firewall/NAT in campus UPFs ensures enterprise-grade security**
- **Local firewall/NAT in campus UPFs ensures traffic control at the site level**

Factoid: UPFs apply DNN or slice-based policy steering, allowing flexible access control across user data flows.

- **UPFs apply DNN-based policy steering**
- **UPFs apply slice-based policy steering**
- **DNN or slice-based policy steering allows flexible access control across user data flows**

Factoid: Separate subnets for user, control, and data traffic per plane support QoS isolation and policy enforcement.

- **Separate subnets for user, control, and data traffic per plane support QoS isolation**
- **Separate subnets for user, control, and data traffic per plane support policy enforcement**

Factoid: Edge and campus UPFs are positioned at county/data-center proximity to offload traffic and reduce backhaul latency (~10–30 ms).

- **Edge and campus UPFs are positioned at county/data-center proximity**
- **Positioning of edge and campus UPFs offloads traffic**
- **Positioning of edge and campus UPFs reduces backhaul latency (~10–30 ms)**

Factoid: Campus UPFs integrate simplified O&M and local data storage to satisfy local processing and security compliance.

- **Campus UPFs integrate simplified O&M**
- **Campus UPFs integrate local data storage**
- **Simplified O&M and local data storage satisfy local processing requirements**
- **Simplified O&M and local data storage satisfy security compliance**

Factoid: Central and regional UPFs focus on high throughput and broader connectivity; edge/campus UPFs target performance-sensitive enterprise usage.

- **Central and regional UPFs focus on high throughput**
- **Central and regional UPFs focus on broader connectivity**
- **Edge and campus UPFs target performance-sensitive enterprise usage**

Factoid: N3 octet counters (**GTP.In/OutDataOctetsN3UPF**) measure data volume between gNB and UPF, optionally per QoS or slice.

- **N3 octet counters (GTP.In/OutDataOctetsN3UPF)** *measure data volume between gNB and UPF*
- **Measurement** *can be done per QoS*
- **Measurement** *can be done per slice*

Factoid: Packet loss on N3 is captured as `GTP.In/OutDataPktLossN3UPF`, enabling per-QoS reliability tracking.

- **Packet loss on N3** *is captured as GTP.In/OutDataPktLossN3UPF*
- **GTP.In/OutDataPktLossN3UPF** *enables per-QoS reliability tracking*

Factoid: Average RTT per DSCP on N3 (`GTP.RttDelayN3DIPsaUpfMean.DSCP`) provides microsecond-level latency metrics for QoS classes.

- **Average RTT per DSCP on N3** *is measured using GTP.RttDelayN3DIPsaUpfMean.DSCP*
- **GTP.RttDelayN3DIPsaUpfMean.DSCP** *provides microsecond-level latency metrics*
- **Microsecond-level latency metrics** *are provided for QoS classes*

Factoid: Out-of-order GTP packet counts (`GTP.InDataPktOutOfOrderN3UPF`) help detect sequence and jitter issues on N3.

- **Out-of-order GTP packet counts (GTP.InDataPktOutOfOrderN3UPF)** *help detect sequence issues on N3*
- **Out-of-order GTP packet counts (GTP.InDataPktOutOfOrderN3UPF)** *help detect jitter issues on N3*

Factoid: N4 interface PFCP session metrics (`SessionEstab`, `ReportSucc`) reflect SMF-UPF control-plane signaling health.

- **N4 interface PFCP session metrics** *include SessionEstab*
- **N4 interface PFCP session metrics** *include ReportSucc*
- **SessionEstab and ReportSucc metrics** *reflect SMF-UPF control-plane signaling health*

Factoid: N6 link usage counters (`IP.N6IncLinkUsage`, `IP.N6OutLinkUsage`) aggregate user-plane bandwidth data toward external networks.

- **N6 link usage counters** *include IP.N6IncLinkUsage*
- **N6 link usage counters** *include IP.N6OutLinkUsage*
- **IP.N6IncLinkUsage and IP.N6OutLinkUsage** *aggregate user-plane bandwidth data*
- **User-plane bandwidth data** *is toward external networks*

Factoid: N9 RTT metrics (`GTP.RttDelayN9*`) quantify latency across chained UPFs, supporting multi-hop performance assessment.

- *N9 RTT metrics (GTP.RttDelayN9)* quantify latency across chained UPFs*
- **Latency quantification via N9 RTT metrics supports multi-hop performance assessment**

Factoid: N9 GTP packet/byte counters support slice-level throughput and control-plane scaling insights.

- **N9 GTP packet counters support slice-level throughput insights**
- **N9 GTP byte counters support slice-level throughput insights**
- **N9 GTP packet and byte counters support control-plane scaling insights**

Factoid: Measurement split by DSCP/QoS/S-NSSAI enables slice-specific SLA monitoring and resource orchestration.

- **Measurement split by DSCP/QoS/S-NSSAI enables slice-specific SLA monitoring**
- **Measurement split by DSCP/QoS/S-NSSAI enables resource orchestration**

Factoid: These interface-level metrics feed QoS enforcement, auto-scaling, and anomaly detection modules in intent-driven orchestration.

- **Interface-level metrics feed QoS enforcement modules**
- **Interface-level metrics feed auto-scaling modules**
- **Interface-level metrics feed anomaly detection modules**
- **These modules operate within intent-driven orchestration**

Factoid: Intel + SK Telecom UPF uses hardware NIC classification and software steering to deliver URLLC-grade priority flows.

Factoid: High-priority UPF traffic achieved consistent RTT of ~0.07–0.09 ms under ~87% CPU load.

Factoid: Priority jitter was reduced to ~±0.014 ms, an ~88% improvement over best-effort traffic.

Factoid: Best-effort traffic latency remained at ~0.3 ms, enabling URLLC without sacrificing throughput.

Factoid: Across varied traffic mixes, high-priority batches consistently saw 32–45 µs latency and 12–14 µs jitter.

Factoid: Deterministic low-latency performance was achieved on COTS Xeon + Intel Ethernet 800 hardware.

Factoid: Priority-based packet handling yields ~78% latency and ~88% jitter improvements relative to normal traffic.

- **Priority-based packet handling** *yields ~78% latency improvement*
- **Priority-based packet handling** *yields ~88% jitter improvement*
- **Latency and jitter improvements** *are relative to normal traffic*

Factoid: Performance remains stable across CPU load variations, ensuring URLLC resilience in production environments.

- **Performance** *remains stable across CPU load variations*
- **Stable performance** *ensures URLLC resilience*
- **URLLC resilience** *is critical in production environments*

Factoid: Hardware-enabled packet steering ensures flow steering to designated cores, optimizing cache and order.

- **Hardware-enabled packet steering** *ensures flow steering to designated cores*
- **Flow steering to designated cores** *optimizes cache utilization*
- **Flow steering to designated cores** *optimizes packet order*

Factoid: Throughput and latency separation in UPF enable multi-tier traffic handling (eMBB and URLLC) on same infrastructure.

- **Throughput and latency separation in UPF** *enable multi-tier traffic handling*
- **Multi-tier traffic handling** *supports eMBB traffic*
- **Multi-tier traffic handling** *supports URLLC traffic*
- **eMBB and URLLC traffic** *share the same infrastructure*

Factoid: Throughput, latency, packet loss, and jitter are the four key QoS metrics defining 5G service quality.

- **Throughput** *is a key QoS metric*
- **Latency** *is a key QoS metric*
- **Packet loss** *is a key QoS metric*
- **Jitter** *is a key QoS metric*
- **These four metrics** *define 5G service quality*

Factoid: Voice, autonomous vehicles, and remote medical services rely critically on low-latency and minimal packet loss.

- **Voice services** *rely critically on low latency*
- **Voice services** *rely critically on minimal packet loss*
- **Autonomous vehicles** *rely critically on low latency*
- **Autonomous vehicles** *rely critically on minimal packet loss*
- **Remote medical services** *rely critically on low latency*
- **Remote medical services** *rely critically on minimal packet loss*

Factoid: UPF, AMF, SMF, PCF, and AUSF form a cohesive 5GC suite enabling scalable, reliable, and flexible network operation.

- **UPF, AMF, SMF, PCF, and AUSF** *form a cohesive 5GC suite*
- **5GC suite** *enables scalable network operation*
- **5GC suite** *enables reliable network operation*
- **5GC suite** *enables flexible network operation*

Factoid: Scalability supports high traffic load, reliability ensures mission-critical service continuity, and flexibility enables smooth 3G/4G to 5G transition.

- **Scalability** *supports high traffic load*
- **Reliability** *ensures mission-critical service continuity*
- **Flexibility** *enables smooth 3G/4G to 5G transition*

Factoid: Future-proof architecture ensures readiness for emerging 5G services while balancing performance objectives and cost.

- **Future-proof architecture** *ensures readiness for emerging 5G services*
- **Future-proof architecture** *balances performance objectives and cost*

Factoid: A logical slicing architecture partitions 5G infrastructure into QoS-aligned logical networks for tailored service delivery.

- **Logical slicing architecture** *partitions 5G infrastructure*
- **Partitioning of 5G infrastructure** *creates QoS-aligned logical networks*
- **QoS-aligned logical networks** *enable tailored service delivery*

Factoid: Slice-aware mobility mechanisms are required to support seamless handover at high speeds (up to 500 km/h).

- **Slice-aware mobility mechanisms** *are required to support seamless handover*
- **Seamless handover** *occurs at high speeds (up to 500 km/h)*

Factoid: Joint optimization of power and subchannel allocation using ILP mitigates co-tier and cross-tier interference in shared spectrum.

Joint optimization of power and subchannel allocation uses ILP (Integer Linear Programming)

- **ILP-based optimization** *mitigates co-tier interference*
- **ILP-based optimization** *mitigates cross-tier interference*
- **Interference mitigation** *occurs in shared spectrum*

Factoid: Dynamic resource allocation among slices supports flexible, demand-driven QoS fulfillment under interference variability.

- **Dynamic resource allocation among slices** *supports flexible QoS fulfillment*
- **Dynamic resource allocation among slices** *supports demand-driven QoS fulfillment*

- **QoS fulfillment** *occurs under interference variability*

Factoid: Network slicing introduces open challenges: slice orchestration, mobility coordination, SDN/NFV integration, and infrastructure reconfiguration.

- **Network slicing** *introduces open challenges*
- **Open challenges** *include slice orchestration*
- **Open challenges** *include mobility coordination*
- **Open challenges** *include SDN/NFV integration*
- **Open challenges** *include infrastructure reconfiguration*

Factoid: 5G target is ≤ 1 ms end-to-end latency with 99.99% reliability for use cases like tactile internet.

- **5G target** *is ≤ 1 ms end-to-end latency*
- **5G target** *is 99.99% reliability*
- **Latency and reliability targets** *support use cases like tactile internet*

Factoid: RAN delays comprise ttx, tbsp, tmpt; LTE's 1 ms TTI must be reduced to ≤ 0.25 ms.

- **RAN delays** *comprise ttx, tbsp, tmpt*
- **LTE's 1 ms TTI** *must be reduced to ≤ 0.25 ms*

Factoid: Sub-ms TTI (0.25 ms) lowers latency at the expense of higher control overhead.

- **Sub-ms TTI (0.25 ms)** *lowers latency*
- **Sub-ms TTI (0.25 ms)** *increases control overhead*

Factoid: Advanced waveforms (GFDM, SC-FDM), and symbol-detection (MMSE, ZF) methods reduce RAN processing delays.

- **Advanced waveforms** *include GFDM and SC-FDM*
- **Symbol-detection methods** *include MMSE and ZF*
- **Advanced waveforms and symbol-detection methods** *reduce RAN processing delays*

Factoid: SDN/NFV and MEC in core/backhaul reduce per-packet latency by eliminating centralized processing hops.

- **SDN/NFV and MEC in core/backhaul** *reduce per-packet latency*
- **Latency reduction** *occurs by eliminating centralized processing hops*

_____ BATCH 10 GEMINI _____

Factoid: Dynamic GTP termination and ultra-dense WDM backhaul can contribute < 0.1 ms transport latency improvements.

- **Dynamic GTP termination** *contributes to transport latency improvements (< 0.1 ms)*
- **Ultra-dense WDM backhaul** *contributes to transport latency improvements (< 0.1 ms)*

Factoid: Edge caching of popular content shortens backhaul dependencies and content retrieval latency.

- **Edge caching of popular content** *shortens backhaul dependencies*
- **Edge caching of popular content** *shortens content retrieval latency*

Factoid: NF resource requirements for low latency include fast CPU, low-latency I/O, and real-time orchestration responsiveness.

- **NF resource requirements for low latency** *include fast CPU*
- **NF resource requirements for low latency** *include low-latency I/O*
- **NF resource requirements for low latency** *include real-time orchestration responsiveness*

Factoid: RAN optimization combined with transport and caching enables sub-ms latency but demands trade-off in overhead and complexity.

- **RAN optimization** combined with **transport and caching** *enables sub-ms latency*
- **Achieving sub-ms latency** *demands trade-offs in overhead and complexity*

Factoid: Reliable ultra-low latency requires harmonised RAN, core, transport, and caching co-design with real-time resource adaptation.

- **Reliable ultra-low latency** *requires harmonised RAN, core, transport, and caching co-design*
- **Reliable ultra-low latency** *requires real-time resource adaptation*

Factoid: An intent like '7 Mbps uplink / ≤ 50 ms latency' can specify AR/VR service requirements without detailing implementation.

- **Intent '7 Mbps uplink / ≤ 50 ms latency'** *specifies AR/VR service requirements*
- **Intent specification** *does not detail implementation specifics*

Factoid: Impacted layers parse intents using RAN identifiers like 5QI and S-NSSAI to scope service intents.

- **Impacted layers** *parse intents* using **RAN identifiers**
- **RAN identifiers** *include 5QI and S-NSSAI*
- **5QI and S-NSSAI** *are used to scope service intents*

Factoid: The Intent Management Function (IMF) checks feasibility, decomposes intent, and orchestrates domain resources.

- **Intent Management Function (IMF)** *checks intent feasibility*
- **IMF** *decomposes intent into actionable components*
- **IMF** *orchestrates domain resources*

Factoid: IMF uses TMF921 and 3GPP APIs for intent lifecycle management, negotiation, and compliance reporting.

- **Intent Management Function (IMF)** *uses TMF921 APIs*
- **IMF** *uses 3GPP APIs*
- **TMF921 and 3GPP APIs** *support intent lifecycle management*
- **TMF921 and 3GPP APIs** *support intent negotiation*
- **TMF921 and 3GPP APIs** *support compliance reporting*

Factoid: Utility functions embedded in intents guide decision-making trade-offs, e.g., latency vs. energy.

- **Utility functions embedded in intents** *guide decision-making trade-offs*
- **Decision-making trade-offs** *include latency vs. energy*

Factoid: E2E intent latency targets (e.g., 100 ms) are decomposed to domain-specific contributions (e.g., 60 ms in RAN).

- **End-to-end (E2E) intent latency targets (e.g., 100 ms)** *are decomposed into domain-specific contributions*
- **Domain-specific contributions** *include 60 ms in RAN*

Factoid: Closed-loop cognitive operations continuously monitor intent compliance, evaluate metrics, and trigger corrective actions.

- **SDAP sublayer** *maps QoS flows to DRBs*
- **QoS flows** *are identified by given QFI*
- **Mapping** *is supported by RRC-configured mapping rules*

Factoid: Multi-layer autonomous domains resolve policy conflicts (e.g. energy-saving vs. performance) through intent exchange.

- **Multi-layer autonomous domains** *resolve policy conflicts*
- **Policy conflicts** *include energy-saving vs. performance*
- **Resolution of policy conflicts** *is achieved through intent exchange*

Factoid: Intent adoption is phased—from static SLOs to dynamic assurance, to full autonomy integrated into business-service workflows.

- **Intent adoption** *is phased*
- **Phases of intent adoption** *include static SLOs*
- **Phases of intent adoption** *include dynamic assurance*

- **Phases of intent adoption** *include full autonomy integrated into business-service workflows*

Factoid: Intent-driven autonomy is expected to materialize in commercial networks between 2025 and 2027.

Factoid: MEF uses controlled natural-language DSLs (Allegro, Cantata) to express business-level intents like ‘Skype for Business → mission-critical SLA.’

- **MEF** *uses controlled natural-language DSLs (Allegro, Cantata)*
- **Controlled natural-language DSLs (Allegro, Cantata)** *express business-level intents*
- **Business-level intents** *include ‘Skype for Business → mission-critical SLA’*

Factoid: Intent expressions from diverse stakeholders are harmonized using MEF models before mapping into enforceable policy rules.

- **Intent expressions from diverse stakeholders** *are harmonized using MEF models*
- **Harmonized intent expressions** *are mapped into enforceable policy rules*

Factoid: Intent-to-policy mapping translates DSL intent into LSO API calls (Legato, Presto, Adagio) for network enforcement.

- **Intent-to-policy mapping** *translates DSL intent into LSO API calls*
- **LSO API calls** *include Legato, Presto, and Adagio*
- **LSO API calls** *are used for network enforcement*

Factoid: MEF enforces declared intents—including performance/security objectives—continuously until manually removed.

- **MEF** *enforces declared intents*
- **Declared intents** *include performance objectives*
- **Declared intents** *include security objectives*
- **Declared intents** *are enforced continuously until manually removed*

Factoid: The intent processing pipeline roles span Business, System, Admin, and Device layers to ensure end-to-end policy coherence.

- **Intent processing pipeline roles** *span Business layer*
- **Intent processing pipeline roles** *span System layer*
- **Intent processing pipeline roles** *span Admin layer*
- **Intent processing pipeline roles** *span Device layer*
- **Spanning all layers** *ensures end-to-end policy coherence*

Factoid: AI/ML modules assist in interpreting intent, synthesizing policies, and orchestrating LSO-controlled network resources.

- **AI/ML modules** *assist in interpreting intent*

- **AI/ML modules** *assist in synthesizing policies*
- **AI/ML modules** *assist in orchestrating LSO-controlled network resources*

Factoid: LSO APIs serve as enforcement endpoints for intent-derived policy rules within automated service lifecycles.

- **LSO APIs** *serve as enforcement endpoints*
- **Enforcement endpoints** *apply to intent-derived policy rules*
- **Intent-derived policy rules** *operate within automated service lifecycles*

Factoid: Intent-based automation addresses scaling limitations and policy complexity by abstracting intent from implementation details.

- **Intent-based automation** *addresses scaling limitations*
- **Intent-based automation** *addresses policy complexity*
- **Intent-based automation** *abstracts intent from implementation details*

Factoid: Continuous intent validation ensures that performance/security objectives remain enforced over time.

- **Continuous intent validation** *ensures performance objectives remain enforced over time*
- **Continuous intent validation** *ensures security objectives remain enforced over time*

Factoid: MEF's IBN framework forms a foundation for autonomous networking by linking declarative intent to operational APIs.

- **MEF's IBN framework** *forms a foundation for autonomous networking*
- **MEF's IBN framework** *links declarative intent to operational APIs*

Factoid: Large-scale programmable 5G/B5G networks require meta-schedulers to coordinate domain-specific schedulers and fulfill user intents.

- **Large-scale programmable 5G/B5G networks** *require meta-schedulers*
- **Meta-schedulers** *coordinate domain-specific schedulers*
- **Meta-schedulers** *fulfill user intents*

Factoid: Meta-schedulers issue high-level intent directives, while local schedulers handle domain-specific resource allocations autonomously.

- **Meta-schedulers** *issue high-level intent directives*
- **Local schedulers** *handle domain-specific resource allocations*
- **Local schedulers** *operate autonomously*

Factoid: Active inference—using causal models—is proposed to predict local scheduler behavior for better meta-scheduling decision-making.

Factoid: Hierarchical and federated learning allow shared intent-driven optimization while preserving domain autonomy.

Factoid: Sub-millisecond scheduler coordination is essential for future 5G/6G latency targets.

Factoid: Meta-scheduler architecture integrates multiple local schedulers via intent-based control loops.

Factoid: Open research areas include formal meta-scheduler modeling, coordination protocols, active inference application, and real-world validation.

Factoid: Intent representation at the meta-scheduler level must be both human-readable and machine-operational across network domains.

Factoid: Programmable network orchestration complexity increases with scale and necessitates intent + meta-scheduling layers.

Factoid: Federated meta-scheduling enables scalability and conflict resolution among hierarchical domain schedulers.

Factoid: User intents (e.g., motion commands in volumetric streaming) are captured online and linked to QoE goals in the SPIRIT framework.

Factoid: A Multi-Armed Bandit path-selection algorithm dynamically chooses SDN paths using application delay and congestion feedback.

Factoid: The system issues SDN flow-rule updates to steer traffic along the best path matching user intent and current metrics.

- **The system *issues* SDN flow-rule updates**
- **SDN flow-rule updates *steer* traffic along the best path**
- **Best path *matches* user intent and current metrics**

Factoid: Intent mapping rules define how each recognized user intent maps to network adaptation policies (delay thresholds, path priority).

- **Intent mapping rules *define how* recognized user intent *maps to* network adaptation policies**
- **Network adaptation policies *include* delay thresholds**
- **Network adaptation policies *include* path priority**

Factoid: QoE is maintained through network reconfiguration even as user motion or network load changes.

- **QoE *is maintained through* network reconfiguration**
- **Network reconfiguration *occurs even as* user motion changes**
- **Network reconfiguration *occurs even as* network load changes**

Factoid: SPIRIT integrates offline intent registration with online capture and dynamic rule enforcement.

Factoid: SDN-controlled path adaptation acts as the execution layer for intent translation in real-time streaming scenarios.

- **SDN-controlled path adaptation** *acts as execution layer for intent translation*
- **SDN-controlled path adaptation** *is used in real-time streaming scenarios*

Factoid: MAB-based adaptation balances exploitation vs exploration to handle environmental variability.

Factoid: Intent-driven adaptation closes the loop: intent capture → path selection → flow steering → QoE validation.

Factoid: Volumetric streaming under dynamic intent benefits from adaptive pathing to satisfy stringent latency/QoE constraints.

Factoid: LLM-centric intent life-cycle architecture spans decomposition, translation, negotiation, activation, and assurance.

Factoid: Plain-language intents obviate the need for manual JSON/YAML structuring by experts.

Factoid: EURECOM implementation uses Code Llama on A100 GPU to map NL intents into Infrastructure-Level Intents.

Factoid: Few-shot learning with human-in-loop feedback supports continuous improvement in intent translation.

Factoid: System demonstrates lifecycle support: NL intent → ILI → activation via NMS → compliance assurance.

Factoid: Key challenges include multi-domain orchestration, LLM interpretability, scalability, and real-time performance.

Factoid: NL intent examples include deploying XR apps requiring vCPUs, memory, throughput, and latency constraints.

Factoid: LLM handles cross-domain decomposition (Cloud, Edge, RAN) in a single unified lifecycle pipeline.

Factoid: Human feedback loop enables refinement of intent translation accuracy over system lifetime.

- **Human feedback loop** *enables refinement of intent translation accuracy*

- **Refinement of intent translation accuracy *occurs over system lifetime***

Factoid: Effectiveness demonstrated in real-world deployment within the EURECOM 5G facility.

Factoid: 5G networks can support IPv6-only user-plane using 464XLAT with CLAT at UE and PLAT in the network.

- **5G networks *can support IPv6-only user-plane***
- **IPv6-only user-plane support *uses 464XLAT***
- **464XLAT *includes CLAT at UE***
- **464XLAT *includes PLAT in the network***

Factoid: IPv4-only UEs receive their address via CLAT while IPv6-only UEs are fully native on the user-plane.

- **IPv4-only UEs *receive their address via CLAT***
- **IPv6-only UEs *are fully native on the user-plane***

Factoid: In roaming scenarios, the location of PDU anchor (home vs visited 5GC) dictates CLAT/PLAT placement.

- **In roaming scenarios, the location of PDU anchor *(home vs visited 5GC) dictates CLAT/PLAT placement***
- **CLAT/PLAT placement *is determined by whether the PDU anchor is in the home or visited 5GC***

Factoid: 464XLAT is deployed in mobile networks; DS-Lite is used in wireline, but end-to-end IPv6 backbone deployment is not yet complete.

Factoid: The draft uses 'MUST', 'SHOULD' per RFC2119 to define policy-level IPv6-only deployment rules.

Factoid: Gradual IPv6-only migration must accommodate mixed UE support: IPv4-only, IPv6-only, and dual-stack.

Factoid: Policy configurations include static IP assignment and network translation behavior per UE capability.

- **Policy configurations *include static IP assignment***
- **Policy configurations *include network translation behavior per UE capability***

Factoid: Backbone IPv6-only deployment lags behind access deployments; draft calls for multi-domain planning.

Factoid: Network translation points should be stateless and per RFC6877-compliant (464XLAT).

Factoid: IPv6-only user-plane offers scalability and IPv4 address exhaustion mitigation in 5G networks.

Factoid: A phased deployment begins with dual-stack, pilots IPv6, and transitions to IPv6-only where feasible.

Factoid: IPv6-only branch office deployment can reduce IPv4 addresses by ~99% using selective dual-stack endpoints.

Factoid: SLAAC is recommended for host auto-configuration; DHCPv6 is used for server/static assignment.

- **SLAAC** *is recommended for host auto-configuration*
- **DHCPv6** *is used for server/static assignment*

Factoid: IPv6 firewalls should use stateful implicit-deny policies rather than NAT for security.

- **IPv6 firewalls** *should use stateful implicit-deny policies*
- **IPv6 firewalls** *should not use NAT for security*

Factoid: Static NAT-PT is deprecated; NAT64/DNS64 is the preferred IPv4-compatibility mechanism.

- **Static NAT-PT** *is deprecated*
- **NAT64/DNS64** *is the preferred IPv4-compatibility mechanism*

Factoid: DNS64 synthesizes AAAA records when only A records exist, enabling IPv6-only host access via NAT64.

- **DNS64** *synthesizes AAAA records when only A records exist*
- **Synthesized AAAA records** *enable IPv6-only host access via NAT64*

Factoid: SLAAC-only clients should use DNS64/NAT64 to interact with IPv4-only services.

- **SLAAC-only clients** *should use DNS64/NAT64*
- **DNS64/NAT64** *enables interaction with IPv4-only services*

Factoid: IPSec is recommended to secure IPv6 communication channels requiring confidentiality.

- **IPSec** *is recommended to secure IPv6 communication channels*
- **Securing IPv6 communication channels** *is required for confidentiality*

Factoid: Firewalls must limit IPv6 multicast scope and monitor IPv6 transition tunnels (e.g., Teredo).

- **Firewalls** *must limit IPv6 multicast scope*

- **Firewalls** *must monitor IPv6 transition tunnels (e.g., Teredo)*

Factoid: IPv6 compatibility requires testing for IPv6 literals, DNS64 behavior, and supporting dual-stack protocols.

Factoid: IPv6 allows mapping of each IP address to a single subscriber, improving accountability and reducing mass surveillance needs.

Factoid: IPv4 shortage in LAC region (~400M users, ~190M addresses) results in ~2 users sharing each IPv4—IPv6 is required to serve the unconnected ~300M people.

Factoid: IPv6 deployment supports exponential growth of connected devices, including IoT, smart cities, and Industry 4.0.

Factoid: Regulators view IPv6 as strategic infrastructure, supporting traceability, security, and resource sustainability.

Factoid: IPv6 enables transparent subscriber tracking, replacing opaque NAT-based sharing with explicit assignment.

Factoid: NCCoE's project demonstrates secure IPv6-only deployment using dual-stack transition in enterprise environments.

Factoid: Micro-segmentation and software-defined perimeters implement zero-trust control by limiting lateral traffic flow.

- **Micro-segmentation** *implements zero-trust control*
- **Software-defined perimeters** *implement zero-trust control*
- **Zero-trust control** *limits lateral traffic flow*

Factoid: Zero-trust architecture on IPv6 emphasizes identity-based access, continuous monitoring, and strict firewall enforcement.

Factoid: IPv6 migration is staged—starting with dual-stack and ending with IPv6-only clients or services.

Factoid: Policies include DHCPv6 shielding, multicast scope enforcement, RADIUS/AAA, and transition protocol monitoring.

Factoid: Use cases range from management of IPv6-only clients to fully IPv6-only enterprise infrastructure.

Factoid: The implementation integrates RFC 7610, 7404, 7381, and NIST SP 800-207 for standards-based IPv6 policy enforcement.

Factoid: Enterprise components like firewalls, MDM, SIEM, and IPS support dual-stack to IPv6-only migration securely.

Factoid: The project includes representative lab environments showing the secure deployment of IPv6 across enterprise scenarios.

Factoid: Zero-trust enforcement in IPv6-only networks hinges on micro-segmentation and strict identity/policy frameworks.

Factoid: 5G private network IPAM systems require dynamic automation and conflict resolution to support millions of connected UE and IoT devices.

- **5G private network IPAM systems** *require dynamic automation*
- **5G private network IPAM systems** *require conflict resolution*
- **Dynamic automation and conflict resolution** *support millions of connected UE and IoT devices*

Factoid: Operators enforce IPv6-only control-plane policies—no public IPv4 allowed on CP links.

- **Operators** *enforce IPv6-only control-plane policies*
- **IPv6-only control-plane policies** *do not allow public IPv4 on CP links*

Factoid: IPv4 addressing is restricted to external-facing slices or middleboxes, with strict NAT at the UPF.

- **IPv4 addressing** *is restricted to external-facing slices*
- **IPv4 addressing** *is restricted to middleboxes*
- **Strict NAT** *is enforced at the UPF*

Factoid: IPAM subnets are isolated per slice (e.g., URLLC, mMTC) and NF zone (e.g., AMF/SMF vs UPF) to enforce QoS and isolation.

- **IPAM subnets** *are isolated per slice (e.g., URLLC, mMTC)*
- **IPAM subnets** *are isolated per NF zone (e.g., AMF/SMF vs UPF)*
- **Isolation per slice and NF zone** *enforces QoS*
- **Isolation per slice and NF zone** *enforces network isolation*

Factoid: Real-time IPAM monitoring should alert network managers to conflicts or rogue addresses in dense deployments.

- **Real-time IPAM monitoring** *should alert network managers*
- **Alerts** *indicate conflicts*
- **Alerts** *indicate rogue addresses*
- **Conflicts and rogue addresses** *are common in dense deployments*

Factoid: IPAM–orchestrator integration enables automated allocation of subnets/IPs during on-the-fly NF instantiation.

- **IPAM–orchestrator integration** *enables automated allocation of subnets/IPs*
- **Automated allocation of subnets/IPs** *occurs during on-the-fly NF instantiation*

Factoid: IPAM must support IPv6-only, dual-stack, and IPv4-translation modes for flexible slice support.

- **IPAM** *must support IPv6-only mode*
- **IPAM** *must support dual-stack mode*
- **IPAM** *must support IPv4-translation mode*
- **Support for these modes enables flexible slice support**

Factoid: IPv6 control-plane networks simplify policy and isolation strategies, improving compliance with operator frameworks.

- **IPv6 control-plane networks** *simplify policy strategies*
- **IPv6 control-plane networks** *simplify isolation strategies*
- **Simplified policy and isolation strategies** *improve compliance with operator frameworks*

Factoid: Operator policies mandate no public IPv4 on CP while allowing IPv4 for UPF-based NAT translation.

- **Operator policies** *mandate no public IPv4 on control plane (CP)*
- **Operator policies** *allow IPv4 for UPF-based NAT translation*

Factoid: Slice-specific addressing enhances enforcement of slice-level policies and traffic segregation.

- **Slice-specific addressing** *enhances enforcement of slice-level policies*
- **Slice-specific addressing** *enhances traffic segregation*

Factoid: NIST blueprint supports both aggregated and disaggregated O-RAN testbeds using srsRAN, FlexRIC, OSC, and Open5GS stacks.

- **NIST blueprint** *supports aggregated O-RAN testbeds*
- **NIST blueprint** *supports disaggregated O-RAN testbeds*
- **NIST blueprint** *uses srsRAN, FlexRIC, OSC, and Open5GS stacks*

Factoid: Time synchronization in O-RAN requires PPS and 10 MHz frequency distribution to USRPs.

- **Time synchronization in O-RAN** *requires PPS (Pulse Per Second)*
- **Time synchronization in O-RAN** *requires 10 MHz frequency distribution*
- **PPS and 10 MHz signals** *are distributed to USRPs*

Factoid: Testbed servers use Ubuntu (20.04/22.04), low-latency kernels, and BIOS tuning (disable C-states, HT, secure boot).

- **Testbed servers** *use Ubuntu (20.04/22.04)*
- **Testbed servers** *use low-latency kernels*

- **BIOS tuning** *includes disabling C-states*
- **BIOS tuning** *includes disabling Hyper-Threading (HT)*
- **BIOS tuning** *includes disabling secure boot*

Factoid: The NIST Testbed Automation Tool automates deployment of gNB, UE, RIC, xApps, and 5G Core on bare-metal or virtual hosts.

- **NIST Testbed Automation Tool** *automates deployment of gNB*
- **NIST Testbed Automation Tool** *automates deployment of UE*
- **NIST Testbed Automation Tool** *automates deployment of RIC*
- **NIST Testbed Automation Tool** *automates deployment of xApps*
- **NIST Testbed Automation Tool** *automates deployment of 5G Core*
- **Deployment** *can occur on bare-metal hosts*
- **Deployment** *can occur on virtual hosts*

Factoid: Automation tool supports configuration of ZMQ/E2 messaging and default Docker IP like 10.53.1.2.

- **Automation tool** *supports configuration of ZMQ/E2 messaging*
- **Automation tool** *supports configuration of default Docker IP (e.g., 10.53.1.2)*

Factoid: Modular orchestration with support for multiple RIC and RAN stacks reduces setup complexity and ensures repeatability.

- **Modular orchestration** *supports multiple RIC and RAN stacks*
- **Support for multiple RIC and RAN stacks** *reduces setup complexity*
- **Support for multiple RIC and RAN stacks** *ensures repeatability*

Factoid: Testbed requires integration of DNS and clock services to support real-time control and function discovery.

- **Testbed** *requires integration of DNS services*
- **Testbed** *requires integration of clock services*
- **DNS and clock services** *support real-time control*
- **DNS and clock services** *support function discovery*

Factoid: xApps (e.g., KPI monitor) are deployed on near-RT RIC, subscribing to E2KP metrics like RSRP via ZMQ or container orchestration.

- **xApps (e.g., KPI monitor)** *are deployed on near-RT RIC*
- **xApps** *subscribe to E2KP metrics*
- **E2KP metrics** *include RSRP*
- **Subscription to metrics** *uses ZMQ*
- **Subscription to metrics** *can use container orchestration*

Factoid: Hardware tuning ensures low-latency and deterministic scheduling suitable for RF-based O-RAN operations.

- **Hardware tuning** *ensures low-latency*
- **Hardware tuning** *ensures deterministic scheduling*
- **Low-latency and deterministic scheduling** *are suitable for RF-based O-RAN operations*

Factoid: Testbed uses srsRAN and Open5GS on commodity hardware and USRP NI-2944R via PCIe for 5G SA operation.

- **Testbed** *uses srsRAN*
- **Testbed** *uses Open5GS*
- **srsRAN and Open5GS** *run on commodity hardware*
- **Testbed** *uses USRP NI-2944R via PCIe*
- **srsRAN, Open5GS, and USRP NI-2944R** *support 5G SA operation*

Factoid: Backhaul uses gigabit Ethernet; upgrade to 10 Gbps is recommended for high-throughput experiments.

- **Backhaul** *uses gigabit Ethernet*
- **Upgrade to 10 Gbps** *is recommended for high-throughput experiments*

Factoid: Time sync via NTP and DNS are essential for control-plane protocols like PFCP, NGAP, HTTP2.

- **Time sync via NTP** *is essential for control-plane protocols*
- **DNS** *is essential for control-plane protocols*
- **Control-plane protocols** *include PFCP, NGAP, and HTTP2*

Factoid: Tuning involves disabling HyperThreading/VT, enabling real-time scheduling, and optimizing buffer sizes and MTU.

- **Tuning** *involves disabling HyperThreading/VT*
- **Tuning** *involves enabling real-time scheduling*
- **Tuning** *involves optimizing buffer sizes*
- **Tuning** *involves optimizing MTU (Maximum Transmission Unit)*

Factoid: RF planning requires configuring USRP gains, duplex mode, SCS, and matching UE APN and compatible bands.

- **RF planning** *requires configuring USRP gains*
- **RF planning** *requires configuring duplex mode*
- **RF planning** *requires configuring subcarrier spacing (SCS)*
- **RF planning** *requires matching UE APN*
- **RF planning** *requires matching compatible bands*

Factoid: Some consumer UEs must be rooted or APN-modified to function reliably in 5G SA testbeds.

- **Some consumer UEs** *must be rooted*

- **Some consumer UEs *must be* APN-modified**
- **Rooting and APN modification *ensure* reliable function in 5G SA testbeds**

Factoid: Open5GS logs (/var/log/open5gs) and srsRAN/ZMQ traces are critical for NGAP, GTP-U, and slice debugging.

- **Open5GS logs (/var/log/open5gs) *are critical for* NGAP debugging**
- **Open5GS logs (/var/log/open5gs) *are critical for* GTP-U debugging**
- **Open5GS logs (/var/log/open5gs) *are critical for* slice debugging**
- **srsRAN/ZMQ traces *are critical for* NGAP, GTP-U, and slice debugging**

Factoid: Basic network slicing using NSSAI (SST/SD) is validated in both RAN and Core domains with packet trace inspection.

- **Basic network slicing *uses* NSSAI (SST/SD)**
- **NSSAI (SST/SD) *is validated in* RAN domain**
- **NSSAI (SST/SD) *is validated in* Core domain**
- **Validation *is performed via* packet trace inspection**

Factoid: Consistent slice enforcement is observable in PFCP session setups and NGAP signaling for UE-slice mapping.

- **Consistent slice enforcement *is observable in* PFCP session setups**
- **Consistent slice enforcement *is observable in* NGAP signaling**
- **PFCP session setups and NGAP signaling *support* UE-slice mapping**

Factoid: Consumer-grade UE variability necessitates RF tuning, root access, and APN configurations to ensure SA connectivity.

- **Consumer-grade UE variability *necessitates* RF tuning**
- **Consumer-grade UE variability *necessitates* root access**
- **Consumer-grade UE variability *necessitates* APN configurations**
- **RF tuning, root access, and APN configurations *ensure* SA connectivity**

Factoid: open5gs-k8s provides both microservice and all-in-one Kubernetes manifests for Open5GS.

- **open5gs-k8s *provides* microservice Kubernetes manifests for Open5GS**
- **open5gs-k8s *provides* all-in-one Kubernetes manifests for Open5GS**

Factoid: Multus and OVS-CNI enable distinct network attachments for N2, N3, N4, preserving plane separation.

- **Multus *enables* distinct network attachments for N2, N3, N4**
- **OVS-CNI *enables* distinct network attachments for N2, N3, N4**
- **Distinct network attachments *preserve* control and user plane separation**

Factoid: MongoDB statefulsets with PVCs store subscriber and NF profile data externally.

- **MongoDB statefulsets with PVCs** *store subscriber data externally*
- **MongoDB statefulsets with PVCs** *store NF profile data externally*

Factoid: CLI and Python scripts automate slice and subscriber provisioning via MongoDB for multi-slice deployments.

- **CLI scripts** *automate slice and subscriber provisioning via MongoDB*
- **Python scripts** *automate slice and subscriber provisioning via MongoDB*
- **Slice and subscriber provisioning** *supports multi-slice deployments*

Factoid: UERANSIM manifests facilitate gNB and UE emulation, including automated ping tests for AMF connectivity.

- **UERANSIM manifests** *facilitate gNB emulation*
- **UERANSIM manifests** *facilitate UE emulation*
- **UERANSIM manifests** *include automated ping tests*
- **Automated ping tests** *verify AMF connectivity*

Factoid: Monarch integration supports real-time slice KPI monitoring when enabled in manifest.

- **Monarch integration** *supports real-time slice KPI monitoring*
- **Real-time slice KPI monitoring** *is enabled in manifest*

Factoid: Supported configurations include Kubernetes v1.28, Ubuntu 22.04, containerd 1.6 – documented via release tags.

- **Supported configurations** *include Kubernetes v1.28*
- **Supported configurations** *include Ubuntu 22.04*
- **Supported configurations** *include containerd 1.6*
- **These configurations** *are documented via release tags*

Factoid: Init containers orchestrate proper startup ordering, ensuring dependency readiness before NF launch.

- **Init containers** *orchestrate proper startup ordering*
- **Proper startup ordering** *ensures dependency readiness before NF launch*

Factoid: HPA-based autoscaling is used for 5G CNFs, while VPA is avoided due to potential service disruptions.

- **HPA-based autoscaling** *is used for 5G CNFs*
- **VPA** *is avoided due to potential service disruptions*

Factoid: Cluster Autoscaler dynamically adds/removes nodes based on HPA pod requirements.

- **Cluster Autoscaler** *dynamically adds/removes nodes*
- **Node scaling decisions** *are based on HPA pod requirements*

Factoid: GitOps (via ArgoCD) ensures consistent CNF configuration and prevents drift across clusters.

- **GitOps (via ArgoCD)** *ensures consistent CNF configuration*
- **GitOps (via ArgoCD)** *prevents configuration drift across clusters*

Factoid: Node pools must remain homogeneous in capacity and configuration, enforced by PodDisruptionBudgets and resource requests.

- **Node pools** *must remain homogeneous in capacity and configuration*
- **Homogeneity** *is enforced by PodDisruptionBudgets*
- **Homogeneity** *is enforced by resource requests*

Factoid: A scale-down fuse protects critical CNFs (e.g., AMF, SMF) from being evicted during resource scaling events.

- **Scale-down fuse** *protects critical CNFs (e.g., AMF, SMF)*
- **Protection from eviction** *occurs during resource scaling events*

Factoid: SCTP protocol (132) must be permitted at network and OS level to support 5G CNF control-plane communication.

- **SCTP protocol (132)** *must be permitted at network level*
- **SCTP protocol (132)** *must be permitted at OS level*
- **Permitting SCTP protocol (132)** *supports 5G CNF control-plane communication*

Factoid: Use of external monitoring (metrics API, Istio) provides data-driven trigger inputs for autoscaling decisions.

- **Use of external monitoring** *provides data-driven trigger inputs for autoscaling decisions*
- **External monitoring tools** *include metrics API*
- **External monitoring tools** *include Istio*

Factoid: Avoid multiple autoscalers in the same node group and check cloud quotas when configuring scaling policies.

- **Avoiding multiple autoscalers in the same node group** *prevents scaling conflicts*
- **Checking cloud quotas** *is important when configuring scaling policies*

Factoid: GitOps-based management hub enables central control of Day-0/Day-2 lifecycle operations for distributed 5G environments.

- **GitOps-based management hub** *enables* central control of Day-0/Day-2 lifecycle operations
- **Day-0/Day-2 lifecycle operations** *apply to* distributed 5G environments

Factoid: Cluster and workload configuration consistency is essential for reliable auto-scaling in telecommunication-grade CNFs.

- **Cluster and workload configuration consistency** *is essential for* reliable auto-scaling
- **Reliable auto-scaling** *applies to* telecommunication-grade CNFs

Factoid: 5G testbeds include remote-area (5G-RANGE), neutral-host city deployment (5GCity), and UAV-assisted RAN with NFV flexibility.

- **5G testbeds** *include* remote-area deployment (5G-RANGE)
- **5G testbeds** *include* neutral-host city deployment (5GCity)
- **5G testbeds** *include* UAV-assisted RAN with NFV flexibility

Factoid: External DNS, internet connectivity, and time sync (PPS/NTP) are mandatory for realistic end-to-end testbed operation.

- **External DNS** *is mandatory for* realistic end-to-end testbed operation
- **Internet connectivity** *is mandatory for* realistic end-to-end testbed operation
- **Time sync (PPS/NTP)** *is mandatory for* realistic end-to-end testbed operation

Factoid: Automated deployment uses open-source toolchains with Docker/K8s for NF instantiation, telemetry, and orchestration.

- **Automated deployment** *uses* open-source toolchains
- **Open-source toolchains** *include* Docker and Kubernetes (K8s)
- **Docker/K8s** *enable* NF instantiation
- **Docker/K8s** *enable* telemetry
- **Docker/K8s** *enable* orchestration

Factoid: KPI measurement includes latency, throughput, reliability, cross-domain sync, collected via telemetry and ML analytics.

- **KPI measurement** *includes* latency
- **KPI measurement** *includes* throughput
- **KPI measurement** *includes* reliability
- **KPI measurement** *includes* cross-domain sync
- **KPI data** *is collected via* telemetry
- **KPI data** *is analyzed with* ML analytics

Factoid: Testbeds bridge lab and field environments through consistent resource pipelines, enabling repeatable validation cycles.

- **Testbeds** *bridge* lab and field environments

- **Testbeds** *use consistent resource pipelines*
- **Consistent resource pipelines** *enable repeatable validation cycles*

Factoid: Neutral-host frameworks (5GCity) allow operators to share infrastructure dynamically using slicing and multitenancy.

- **Neutral-host frameworks (e.g., 5GCity)** *allow operators to share infrastructure dynamically*
- **Infrastructure sharing** *is enabled by slicing*
- **Infrastructure sharing** *is enabled by multitenancy*

Factoid: 5G-RANGE combines fixed RAN with UAVs supported by NFV for sporadic deployments in remote regions.

- **5G-RANGE** *combines fixed RAN with UAVs*
- **5G-RANGE** *is supported by NFV*
- **5G-RANGE** *enables sporadic deployments in remote regions*

Factoid: End-to-end KPI validation must consider NFV stack maturity, telemetry quality, and orchestration tool chain alignment.

- **End-to-end KPI validation** *must consider NFV stack maturity*
- **End-to-end KPI validation** *must consider telemetry quality*
- **End-to-end KPI validation** *must consider orchestration tool chain alignment*

Factoid: Edge and core synchronization in testbeds hinges on coordinated time sync, network path management, and telemetry integration.

- **Edge and core synchronization in testbeds** *hinges on coordinated time sync*
- **Edge and core synchronization in testbeds** *hinges on network path management*
- **Edge and core synchronization in testbeds** *hinges on telemetry integration*

Factoid: Lab-to-production deployment is facilitated by standardized virtualization patterns across SDR, core, and orchestration environments.

- **Lab-to-production deployment** *is facilitated by standardized virtualization patterns*
- **Standardized virtualization patterns** *apply across SDR environments*
- **Standardized virtualization patterns** *apply across core environments*
- **Standardized virtualization patterns** *apply across orchestration environments*